



# OpenSplice|DDS

Delivering Performance, Openness, and Freedom

**Angelo Corsaro, Ph.D.**

Chief Technology Officer

OMG DDS SIG Co-Chair

[angelo.corsaro@prismtech.com](mailto:angelo.corsaro@prismtech.com)



# The DDS Tutorial

::Part II

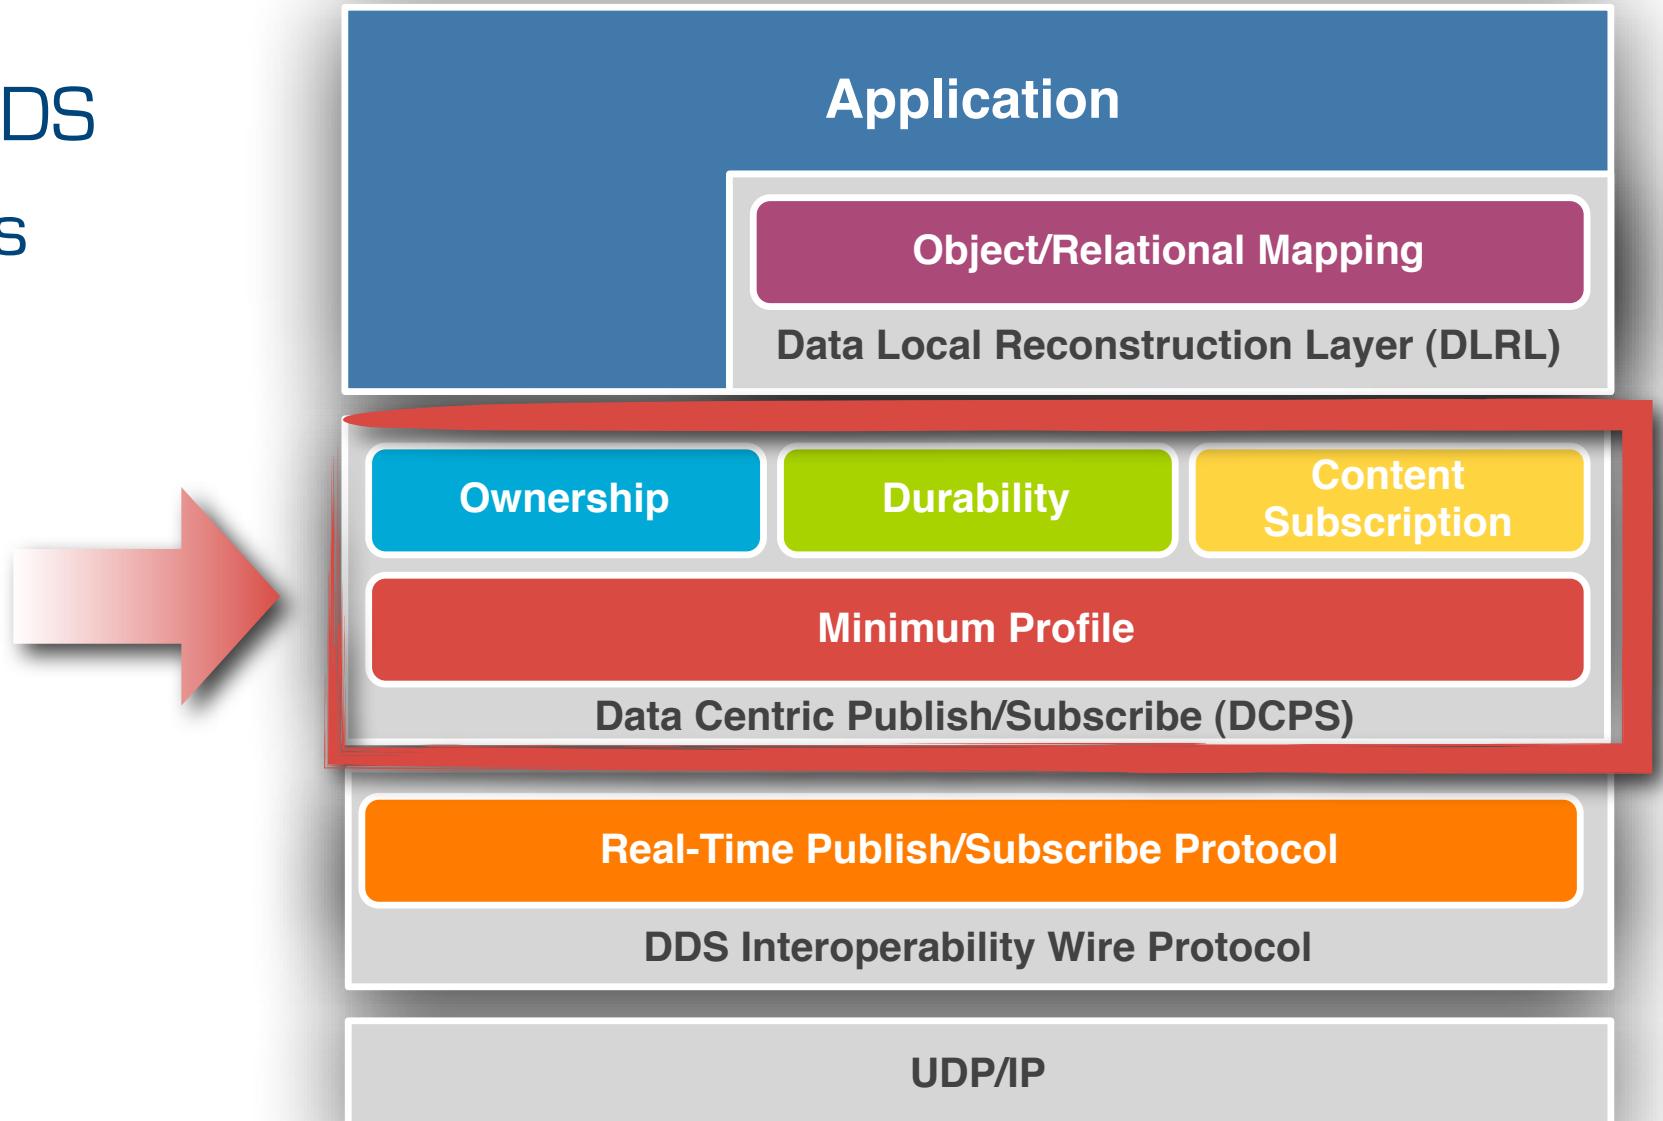
# Tutorial Scope

## Scope & Goals

- ▶ The Tutorial will cover the DCPS layer of DDS
- ▶ It will give you enough details and examples to make sure that you can get started writing DDS applications

## Software

- ▶ OpenSplice DDS
  - ▶ <http://www.opensplice.org>
- ▶ SIMple Dds (SIMD)
  - ▶ <http://code.google.com/p/simd-cxx>



## Prerequisite

- ▶ Basic C++ understanding

# What You've Learned on Part I

- ▶ Defining Topics and Topic Types
- ▶ Scoping Information with Partitions
- ▶ Writing Data
- ▶ Reading (Taking) data with Waitsets and Listeners
- ▶ Writing an example that demonstrate all of the above

# What we'll Cover Today

- ▶ Content Filtered Topics and Queries
- ▶ QoS and the Request vs. Offered Model
- ▶ Setting QoS on DDS Entities
- ▶ Tuning OpenSplice DDS Configuration

# OpenSplice|DDS

Delivering Performance, Openness, and Freedom

## Your will learn:

- What Filters and Queries are
- The available SQL92 subset
- Programming Filters and Queries

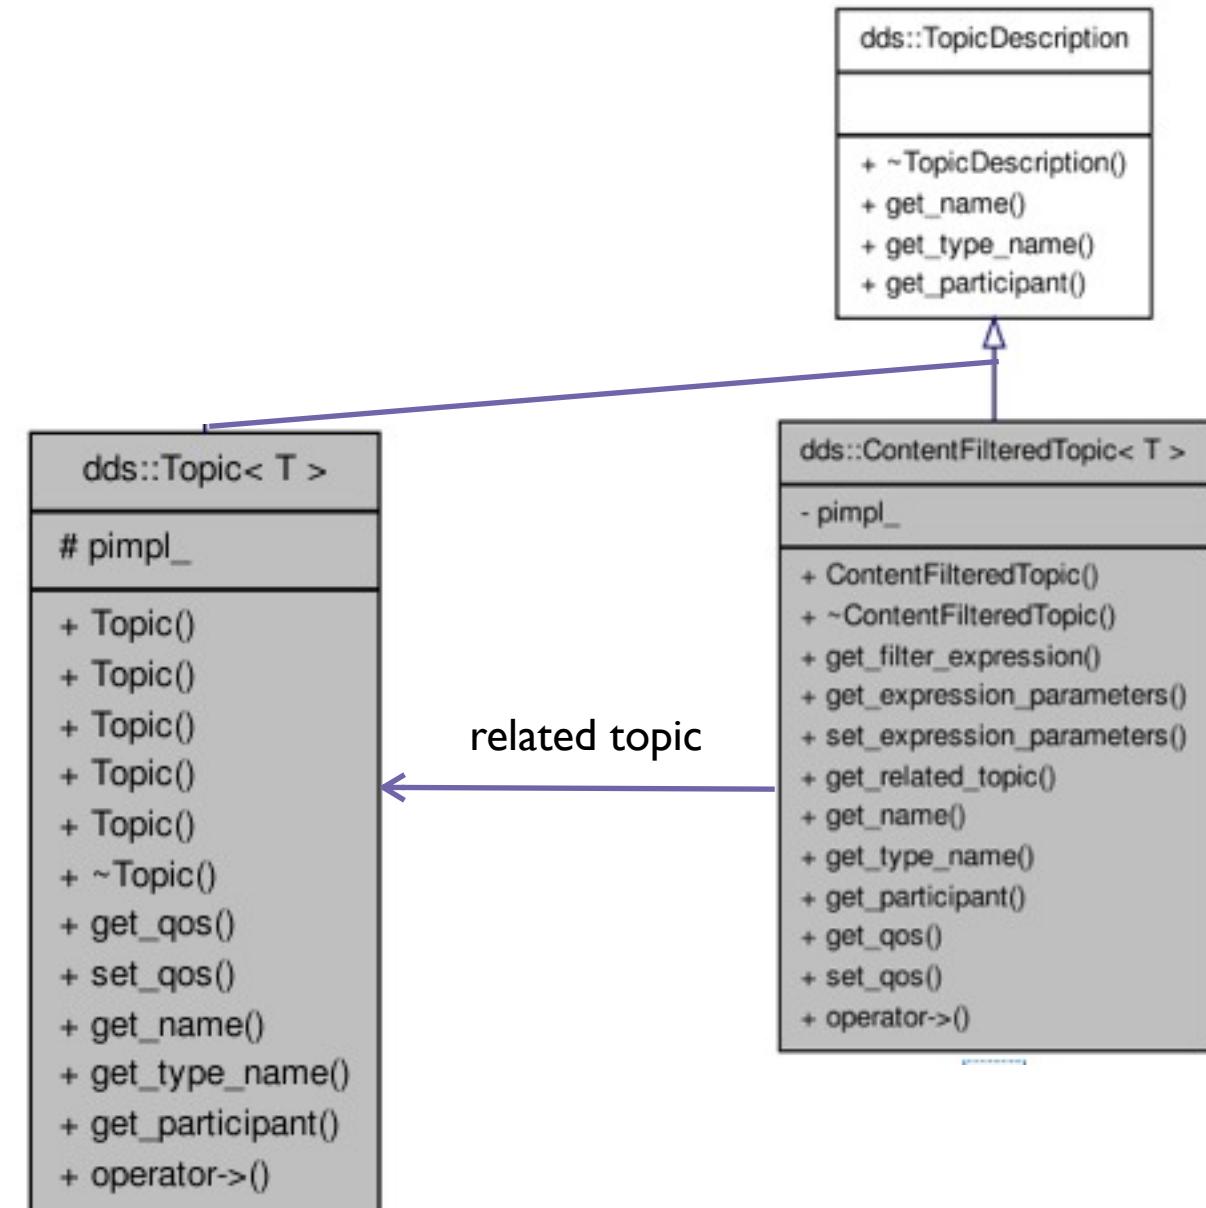
## Filters and Queries

# DDS Filters & Queries

- ▶ DDS provides means for **filtering data** using an application specified **condition expression**
- ▶ The two mechanism provided in order to **filter data** are:
  - ▶ `ContentFilteredTopics`
  - ▶ `QueryCondition`
- ▶ Both working in conjunction with a `DataReader`

# ContentFilteredTopic

- ▶ A ContentFilteredTopic can be seen as a decorator of a user defined topic defining a specific filter over the data published for the given topic
- ▶ As a result, a ContentFilteredTopic always exist in conjunction with its related Topic
- ▶ A DataReader created with a specific ContentFilteredTopic will only receive the data that matches the filter condition
- ▶ ContentFilteredTopic can be thought of as Continuous Queries



# QueryCondition

- ▶ A **QueryCondition** can be created over an existing **DataReader** in order to select, among received data, only the subset matching the condition associated with the **QueryCondition**
- ▶ **QueryCondition** are created for a specific **DataReader** and used as an argument of the **DataReader::read\_w\_condition** method

# Filters & Queries Grammar

*Condition ::= Predicate*

- | *Condition* 'AND' *Condition*
- | *Condition* 'OR' *Condition*
- | 'NOT' *Condition* | '(' *Condition* ')'

*Predicate ::= ComparisonPredicate | BetweenPredicate*

*ComparisonPredicate ::= FIELDNAME RelOp Parameter | Parameter RelOp FIELDNAME*

*BetweenPredicate ::= FIELDNAME 'BETWEEN' Range | FIELDNAME 'NOT BETWEEN' Range*

*RelOp ::= '=' | '>' | '>=' | '<' | '<=' | '<>' | like*

*Range ::= Parameter 'AND' Parameter*

*Parameter ::= INTEGERVALUE | FLOATVALUE  
| STRING | ENUMERATEDVALUE | PARAMETER*

where PARAMETER has the form %n (with n in 0..100)

# ContentFilteredTopics in SIMD

```
template <typename T>
class dds::ContentFilteredTopic : public dds::TopicDescription {
public:
    ContentFilteredTopic(const std::string& name,
                         const dds::Topic<T>& t,
                         const std::string& filter,
                         const std::vector<std::string>& params);

    virtual ~ContentFilteredTopic();
public:
    std::string get_filter_expression() const;

    std::vector<std::string> get_expression_parameters() const;
    void set_expression_parameters(const std::vector<std::string>& params);

    dds::Topic<T> get_related_topic() const;

    virtual std::string get_name() const;

    virtual std::string get_type_name() const;

    virtual dds::DomainParticipant get_participant() const;

    TopicQos get_qos() const;
    void set_qos(const TopicQos& qos);
};
```

# Using ContentFilteredTopics

```
// Create the "TempSensor" Topic
dds::Topic<TempSensorType> tsTopic("TempSensor");

// Create the filter parameters
std::vector<std::string> params(2);
params[0] = "30";
params[1] = "0.6";

// Create the ContentFilteredTopic
dds::ContentFilteredTopic<TempSensorType>
cfTsTopic("TempSensor-1",
          tsTopic,
          "(temp < %0) AND (hum < %1)",
          params);

// Create the DataReader with for the ContentFilteredTopic
dds::DataReader<TempSensorType> dr(cfTsTopic);
```

```
enum TemperatureScale {
    CELSIUS,
    FAHRENHEIT,
    KELVIN
};

struct TempSensorType {
    short id;
    float temp;
    float hum;
    TemperatureScale scale;
};
#pragma keylist TempSensorType id
```

# QueryCondition in SIMD

```
template <typename T>
class DataReader {
public:
    // [...] Other DataReader methods
    QueryCondition create_querycondition(const std::string& expression, const std::vector& params);

    QueryCondition create_querycondition(const SampleStateMask& samples_state,
                                         const ViewStateMask& views_state,
                                         const InstanceStateMask& instances_state,
                                         const std::string& expression, const std::vector& params);

    ReturnCode_t read_w_condition(TSeq& samples, SampleInfoSeq& infos, const QueryCondition& cond);

    ReturnCode_t read_w_condition(TSeq& samples, SampleInfoSeq& infos,
                                 const QueryCondition& cond, long max_samples);

    ReturnCode_t take_w_condition(TSeq& samples, SampleInfoSeq& infos, const QueryCondition& cond);

    ReturnCode_t take_w_condition(TSeq& samples, SampleInfoSeq& infos, const QueryCondition& cond,
                                 long max_samples);
}
```

# Using QueryConditions

```
dds::Topic<TempSensorType> tsTopic("TempSensor");
dds::DataReader<TempSensorType> dr(tsTopic);

std::vector<std::string> params(2);
params[0] = "30";
params[1] = "0.6";

// Create Query Condition
dds::QueryCondition cond =
    dr.create_querycondition("(temp < %0) AND (hum < %1)", params);

TempSensorTypeSeq data;
SampleInfoSeq info;
// Read with Condition
dr.read_w_condition(data, info, cond);
```

# Filters vs Queries

DataReader



(temp < 30) AND (hum < 0.6)

**ContentFilteredTopic**

1	26	0.4	C
1	28	0.5	C
1	31	0.5	C
1	30	0.7	C
1	29	0.5	C

**QueryCondition**

(temp < 30) AND (hum < 0.6)

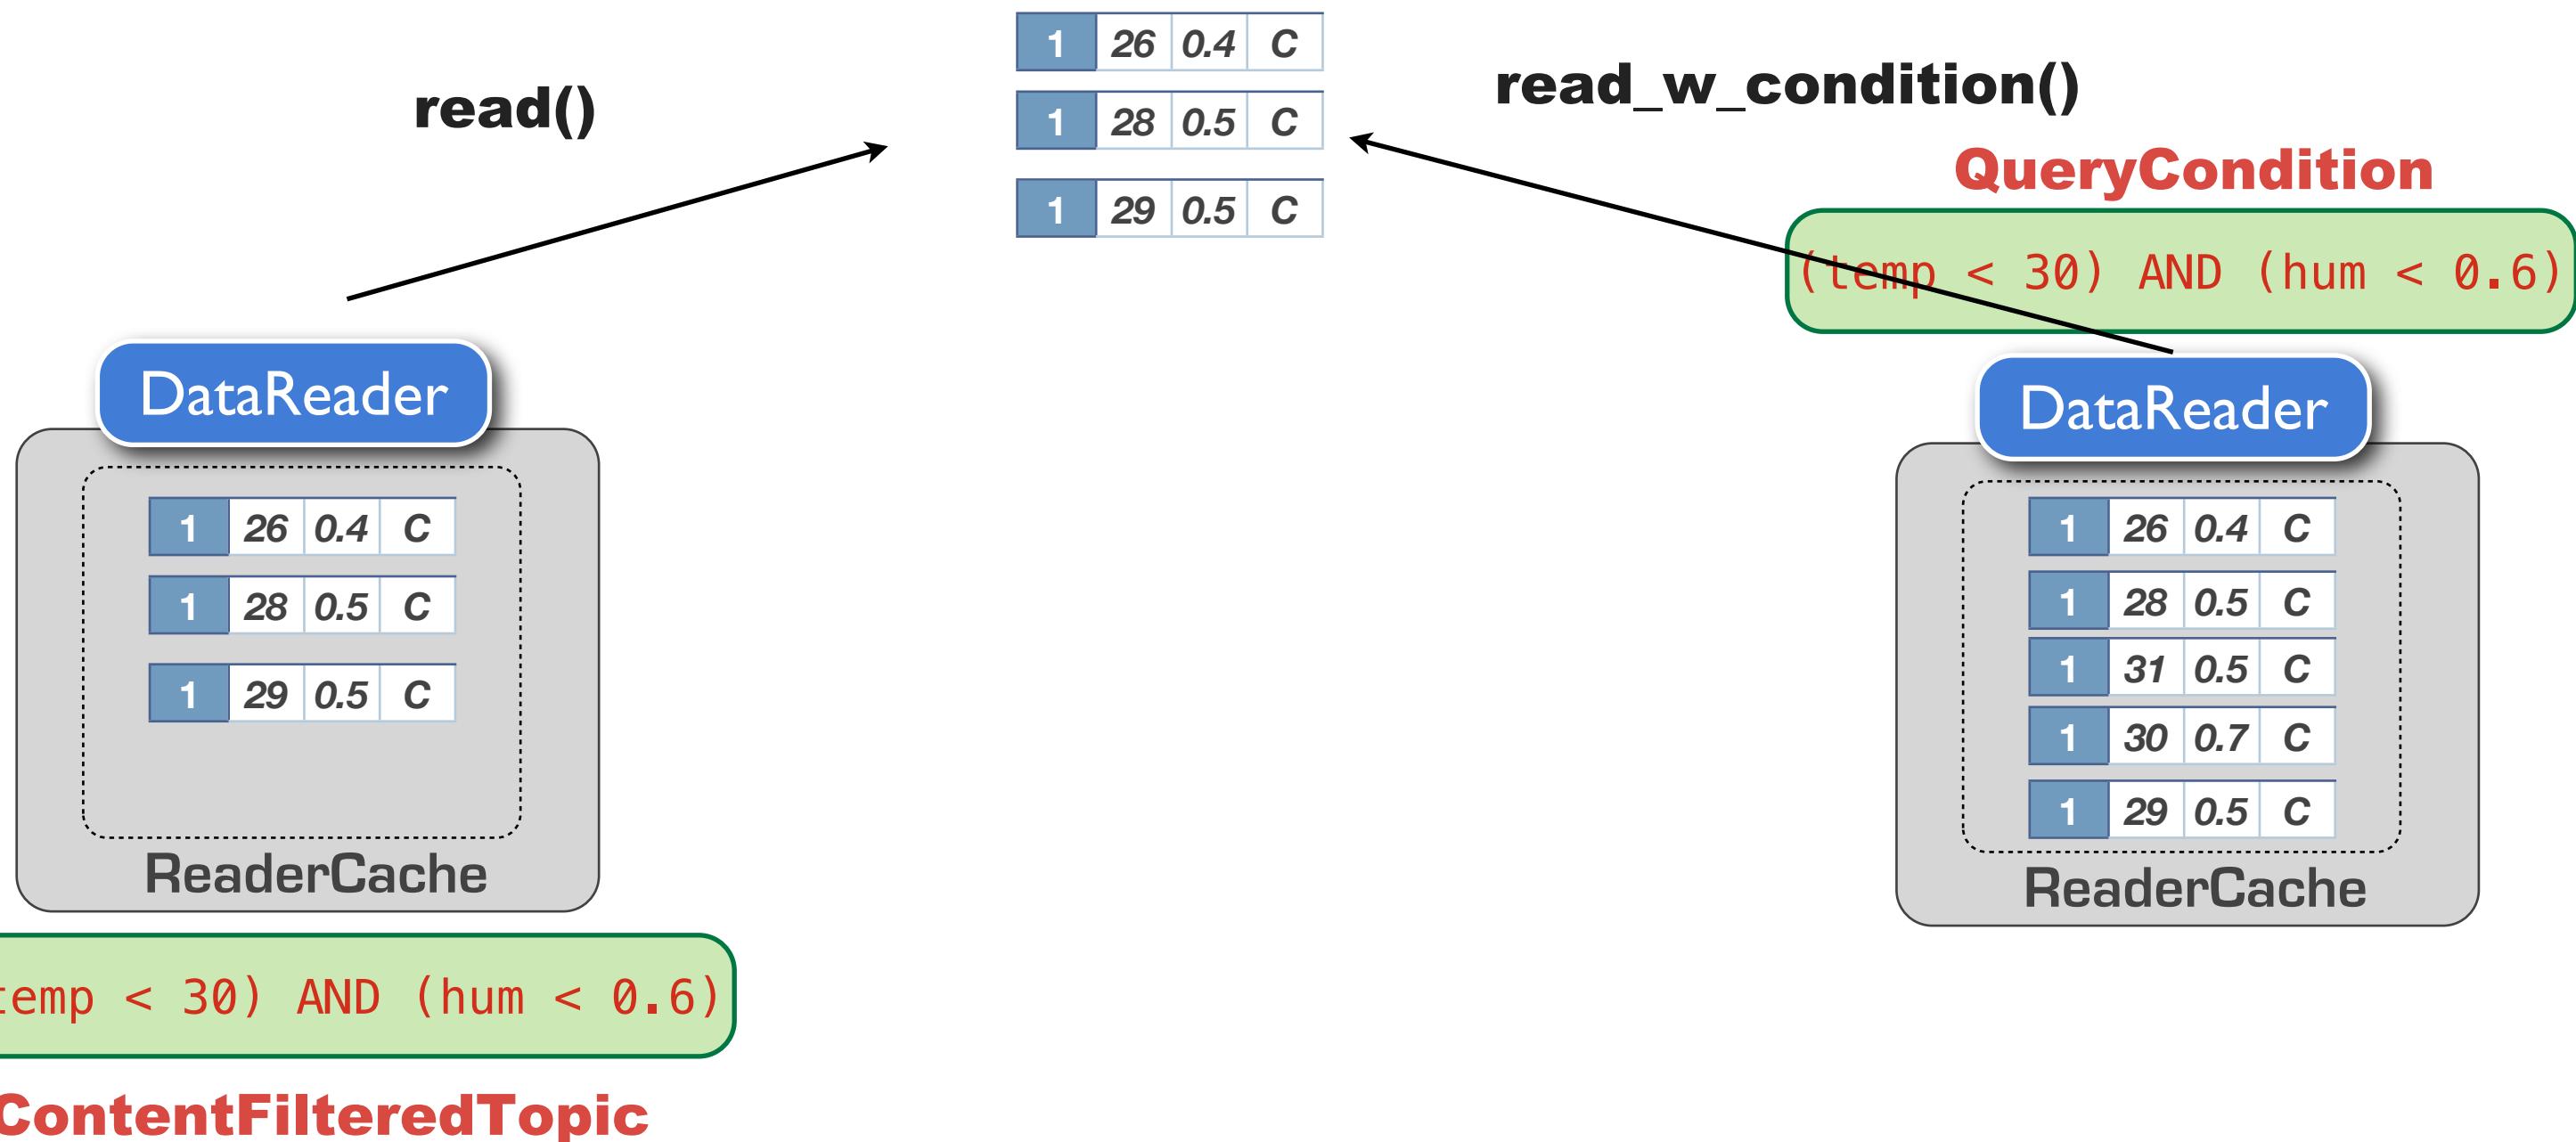
DataReader



# Filters vs Queries



# Filters vs Queries



# OpenSplice|DDS

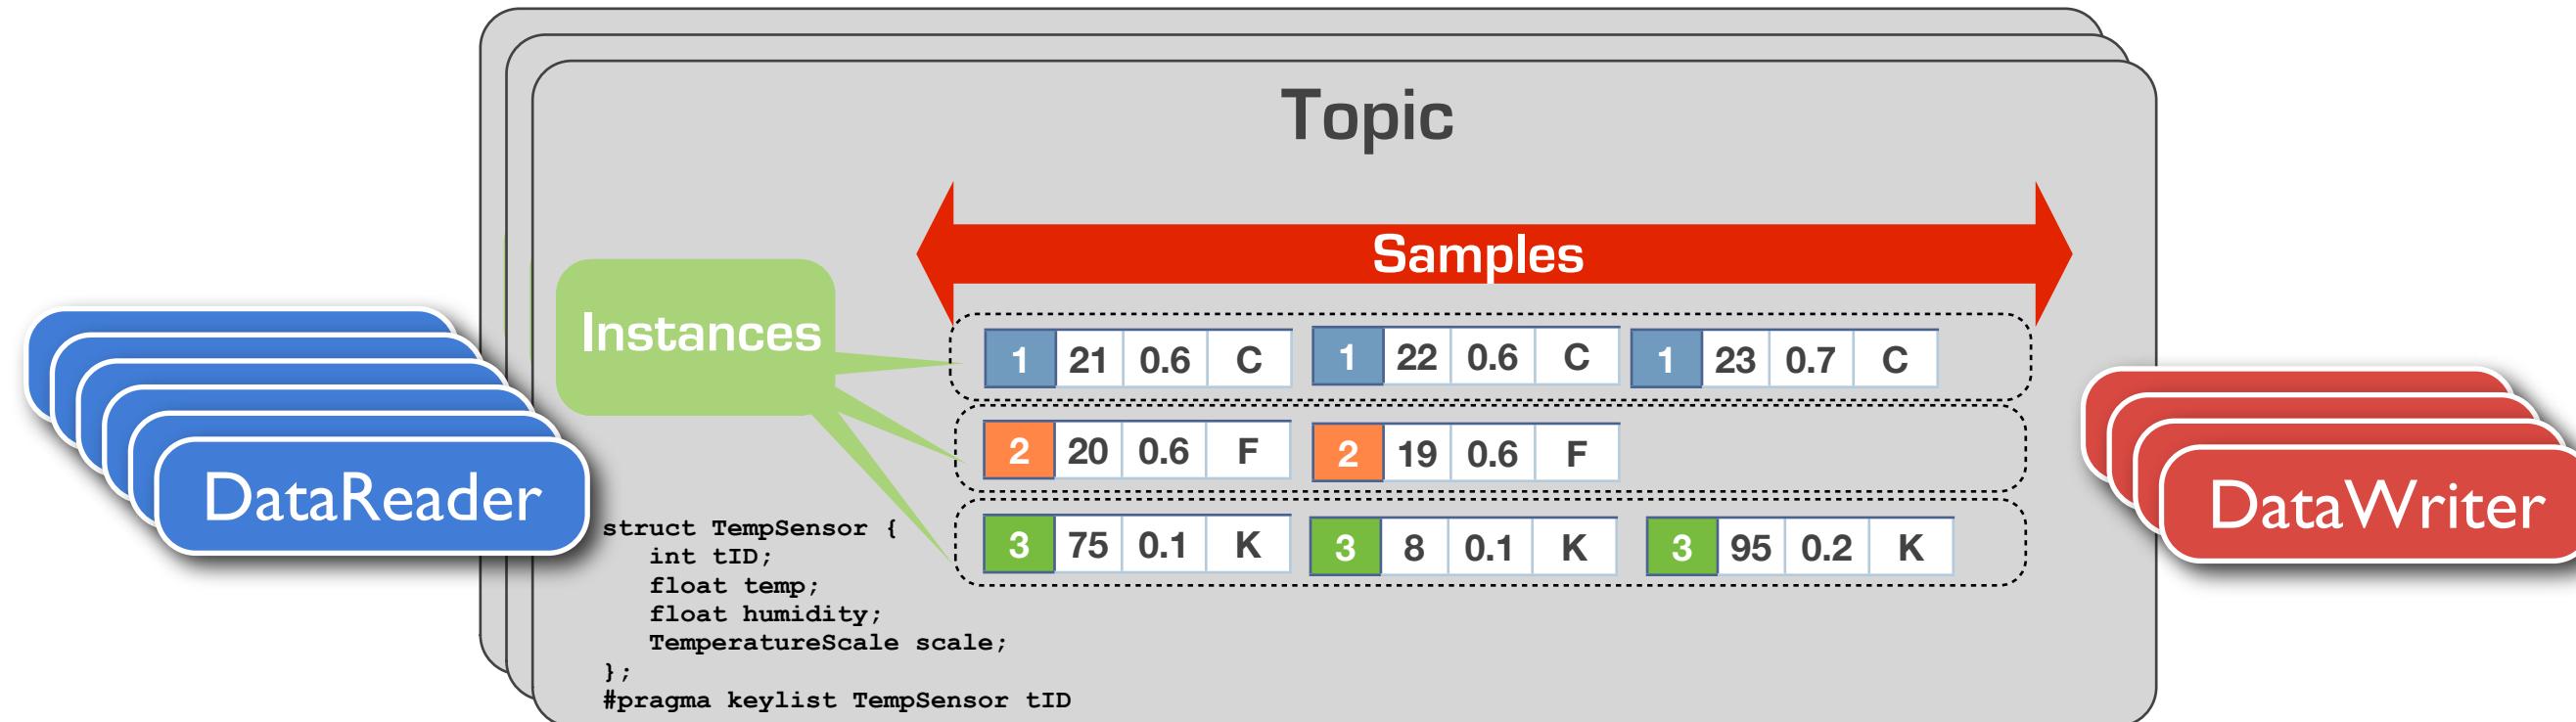
Delivering Performance, Openness, and Freedom

## Your will learn:

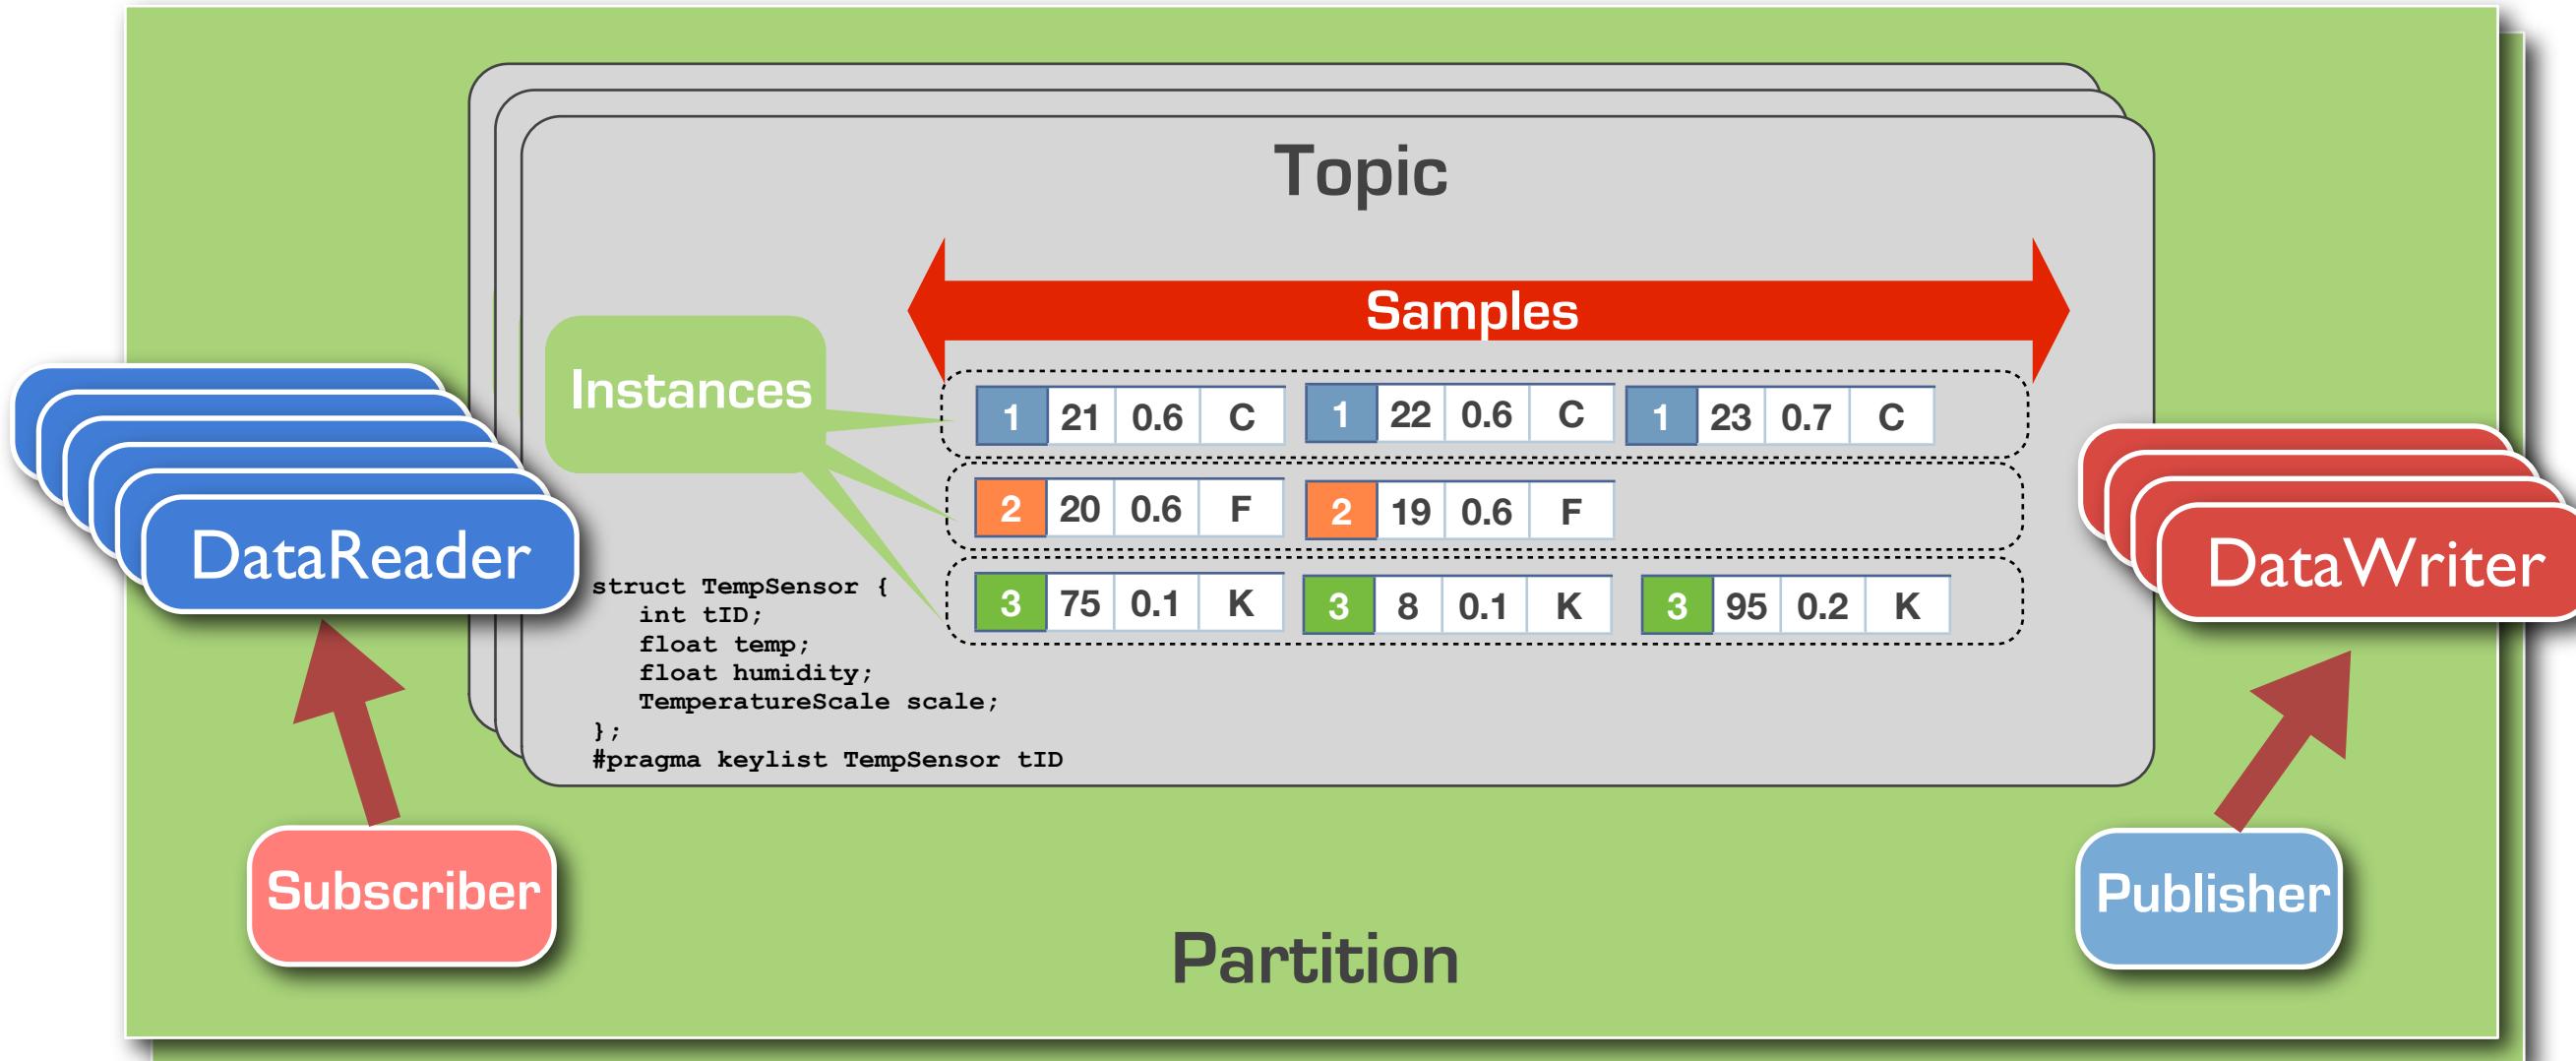
- What can you control with QoS
- Request vs. Offered in DDS
- QoS Patterns

## Quality of Service (QoS)

# Anatomy of a DDS Application

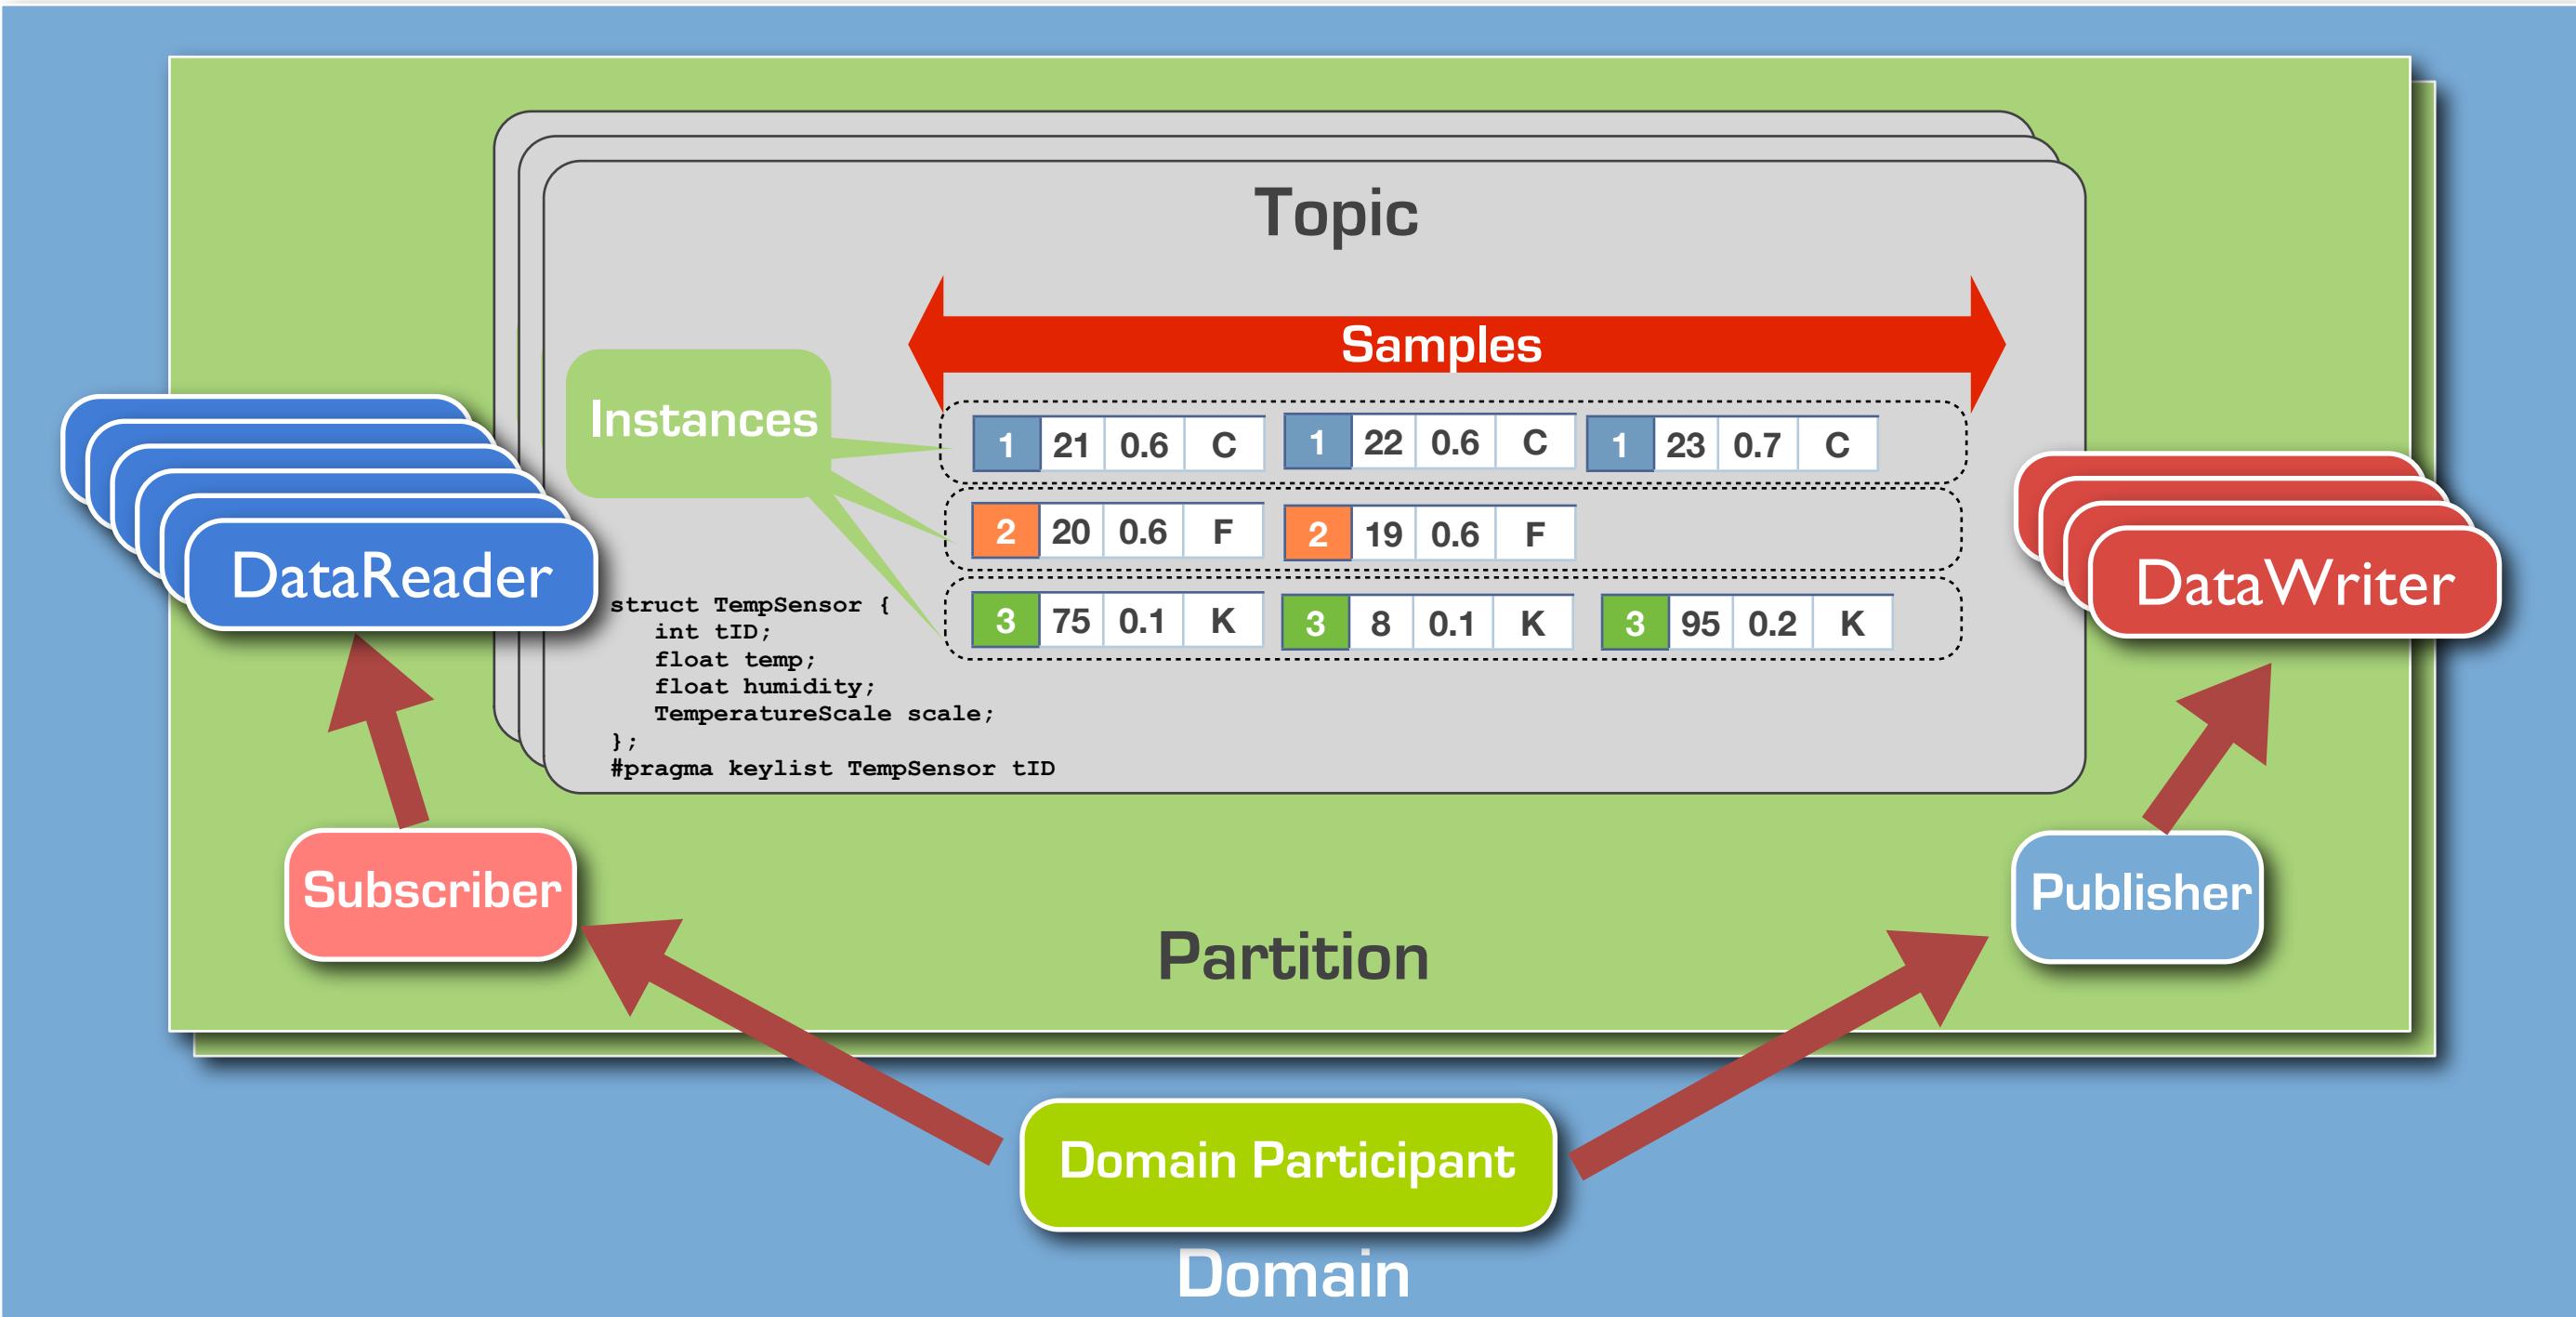


# Anatomy of a DDS Application



# Anatomy of a DDS Application

→  
Arrows  
show  
structural  
relationship  
s, not data-  
flows



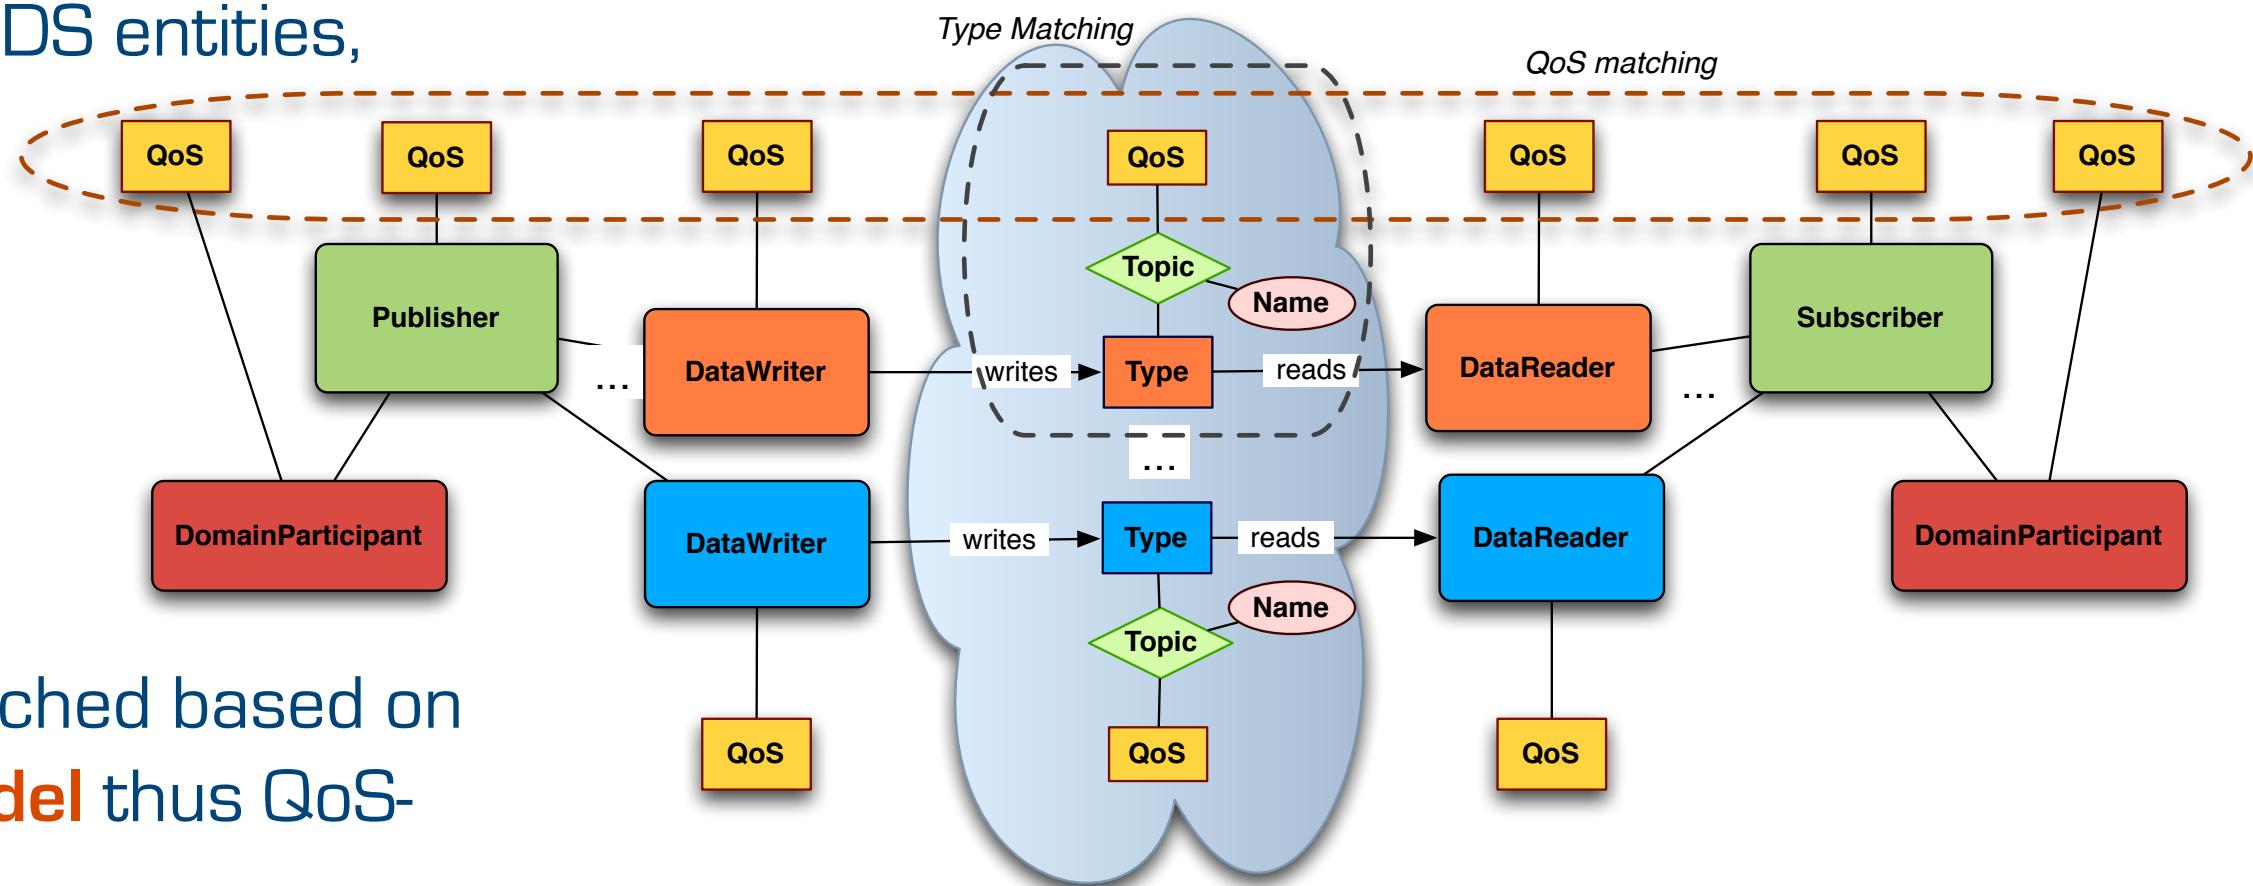
# QoS Model

- ▶ QoS-Policies are used to control relevant properties of OpenSplice DDS entities, such as:

- ▶ Temporal Properties
- ▶ Priority
- ▶ Durability
- ▶ Availability
- ▶ ...

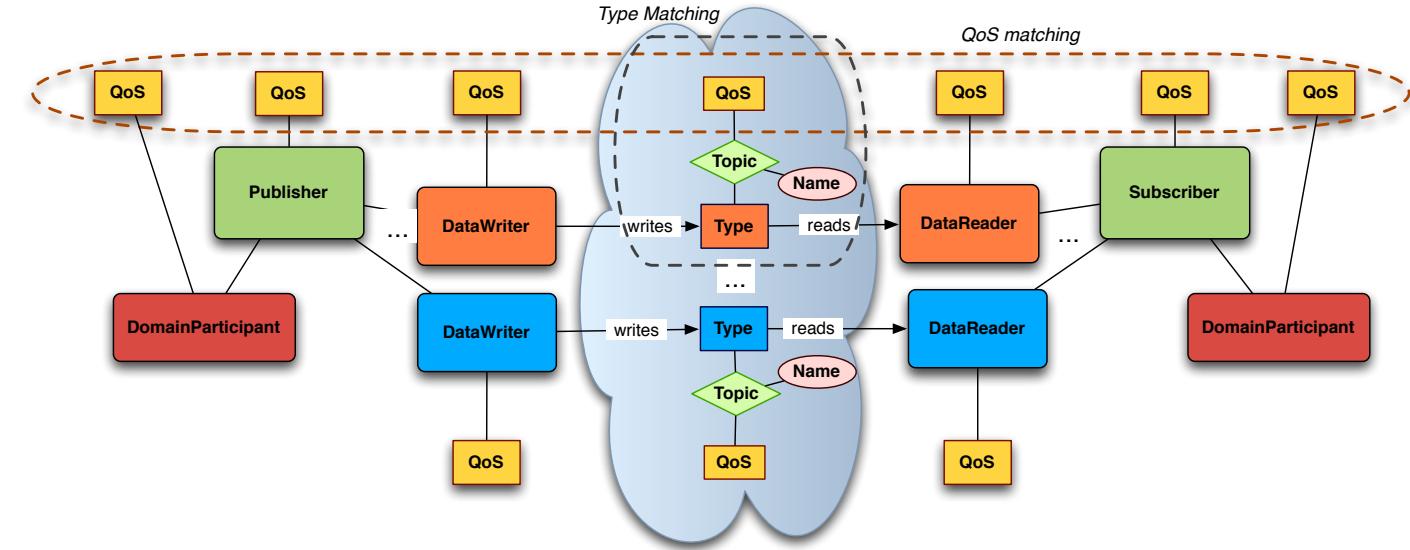
- ▶ Some QoS-Policies are matched based on a **Request vs. Offered Model** thus QoS-enforcement

- ▶ Publications and Subscriptions match only if the declared vs. requested QoS are compatible
  - ▶ e.g., it is not possible to match a publisher which delivers data unreliable with a subscriber which requires reliability



# QoS Policies

QoS Policy	Applicability	RxO	Modifiable	
DURABILITY	T, DR, DW	Y	N	Data Availability
DURABILITY SERVICE	T, DW	N	N	
LIFESPAN	T, DW	N/A	Y	
HISTORY	T, DR, DW	N	N	
PRESENTATION	P, S	Y	N	
RELIABILITY	T, DR, DW	Y	N	
PARTITION	P, S	N	Y	
DESTINATION ORDER	T, DR, DW	Y	N	
OWNERSHIP	T, DR, DW	Y	N	
OWNERSHIP STRENGTH	DW	N/A	Y	
DEADLINE	T, DR, DW	Y	Y	Data Timeliness
LATENCY BUDGET	T, DR, DW	Y	Y	
TRANSPORT PRIORITY	T, DW	N/A	Y	
TIME BASED FILTER	DR	N/A	Y	Resources
RESOURCE LIMITS	T, DR, DW	N	N	
USER_DATA	DP, DR, DW	N	Y	Configuration
TOPIC_DATA	T	N	Y	
GROUP_DATA	P, S	N	Y	



- Rich set of QoS allow to configure several different aspects of data availability, delivery and timeliness
- QoS can be used to control and optimize network as well as computing resource

# OpenSplice|DDS

Delivering Performance, Openness, and Freedom

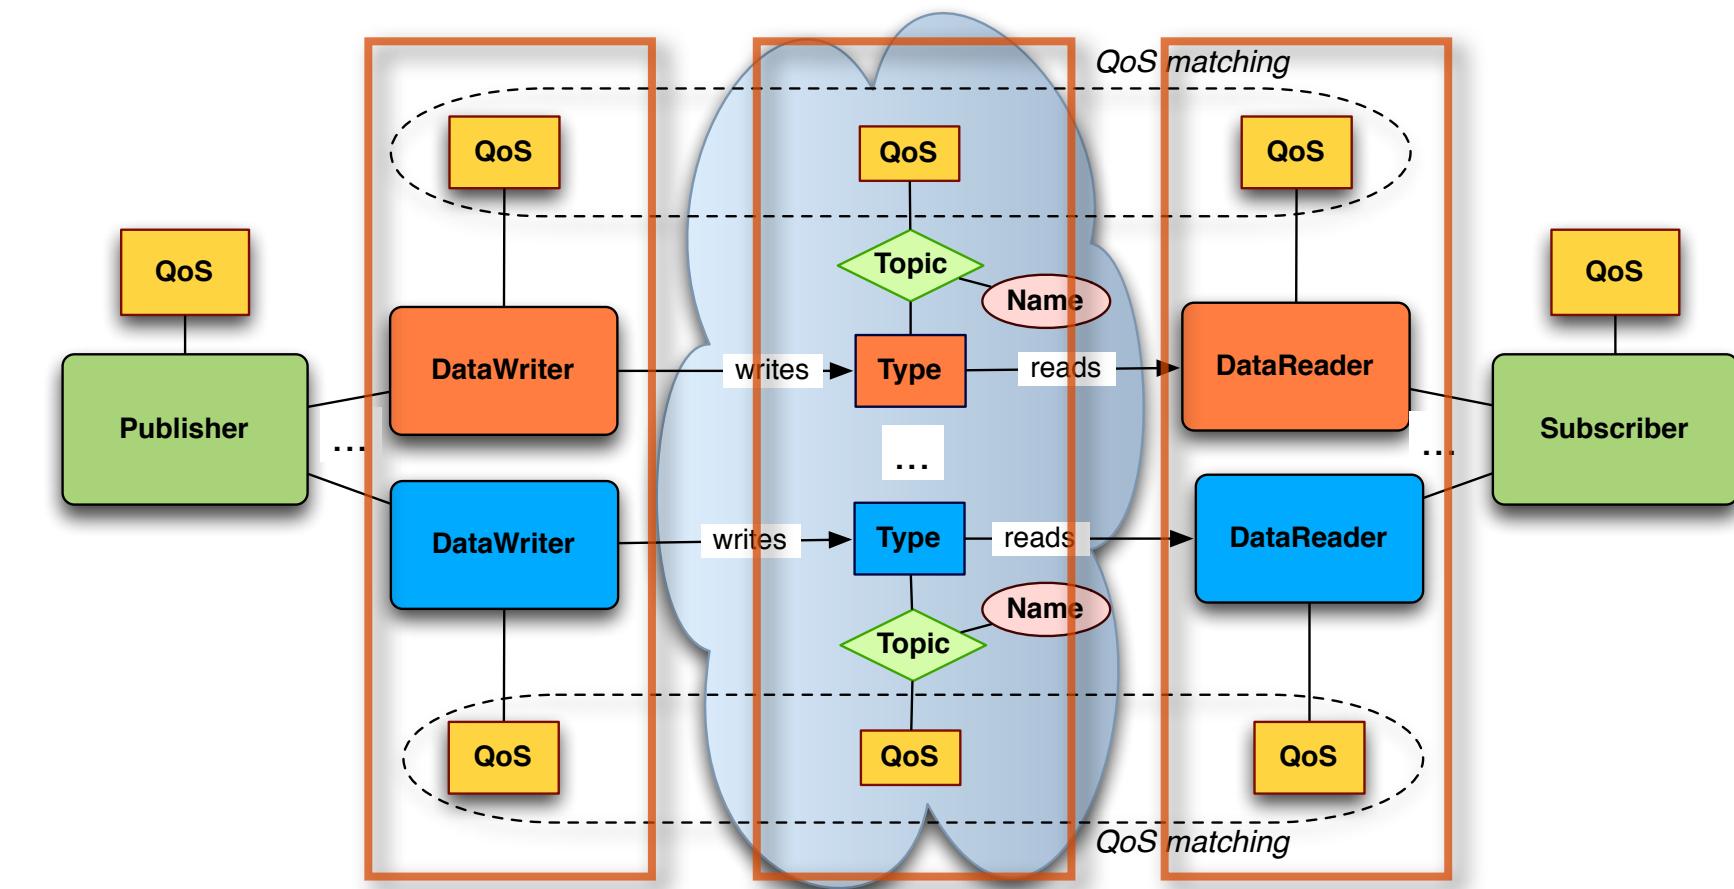
Controlling Reliability

# Reliability

The **RELIABILITY** QoS indicate the level of guarantee offered by the DDS in delivering data to subscribers. Possible variants are:

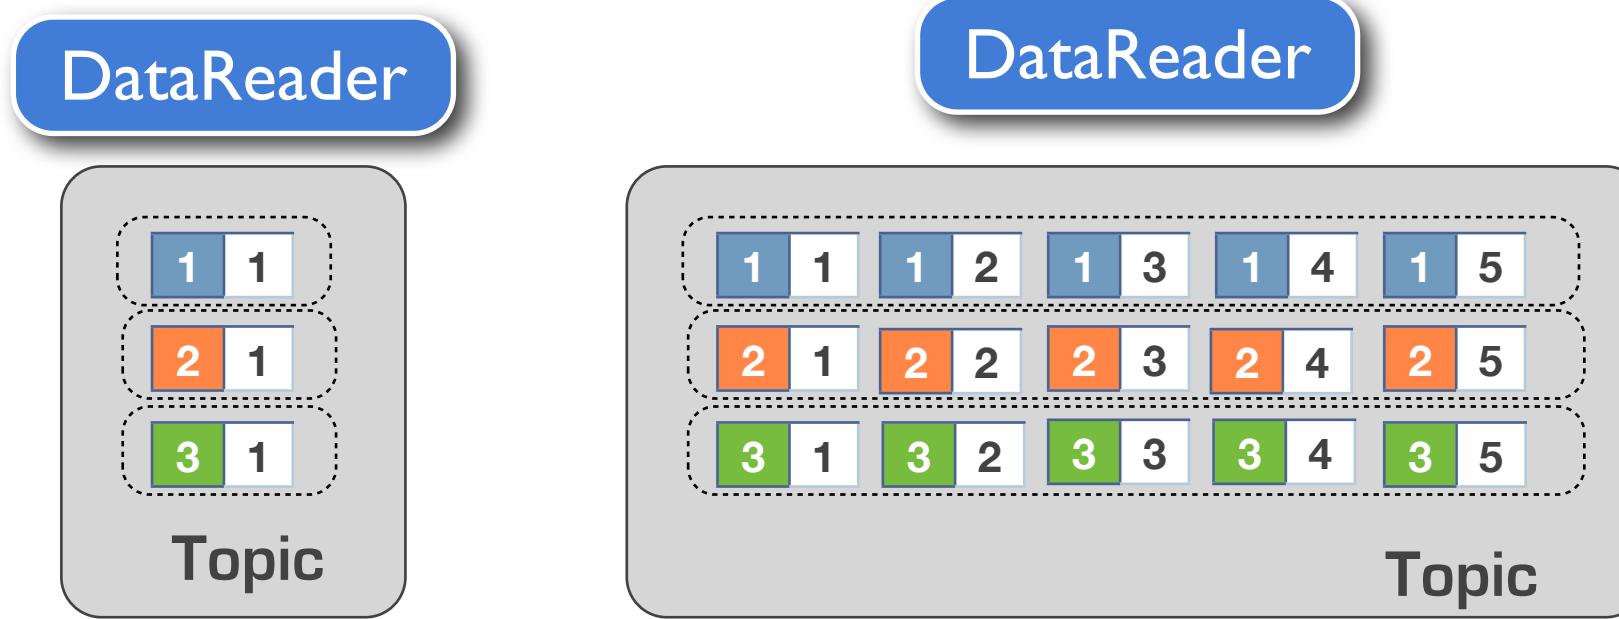
- ▶ **Reliable.** In steady-state the middleware guarantees that all samples in the DataWriter history will eventually be delivered to all the DataReader
- ▶ **Best Effort.** Indicates that it is acceptable to not retry propagation of any samples

QoS Policy	Applicability	RxO	Modifiable
RELIABILITY	T, DR, DW	Y	N



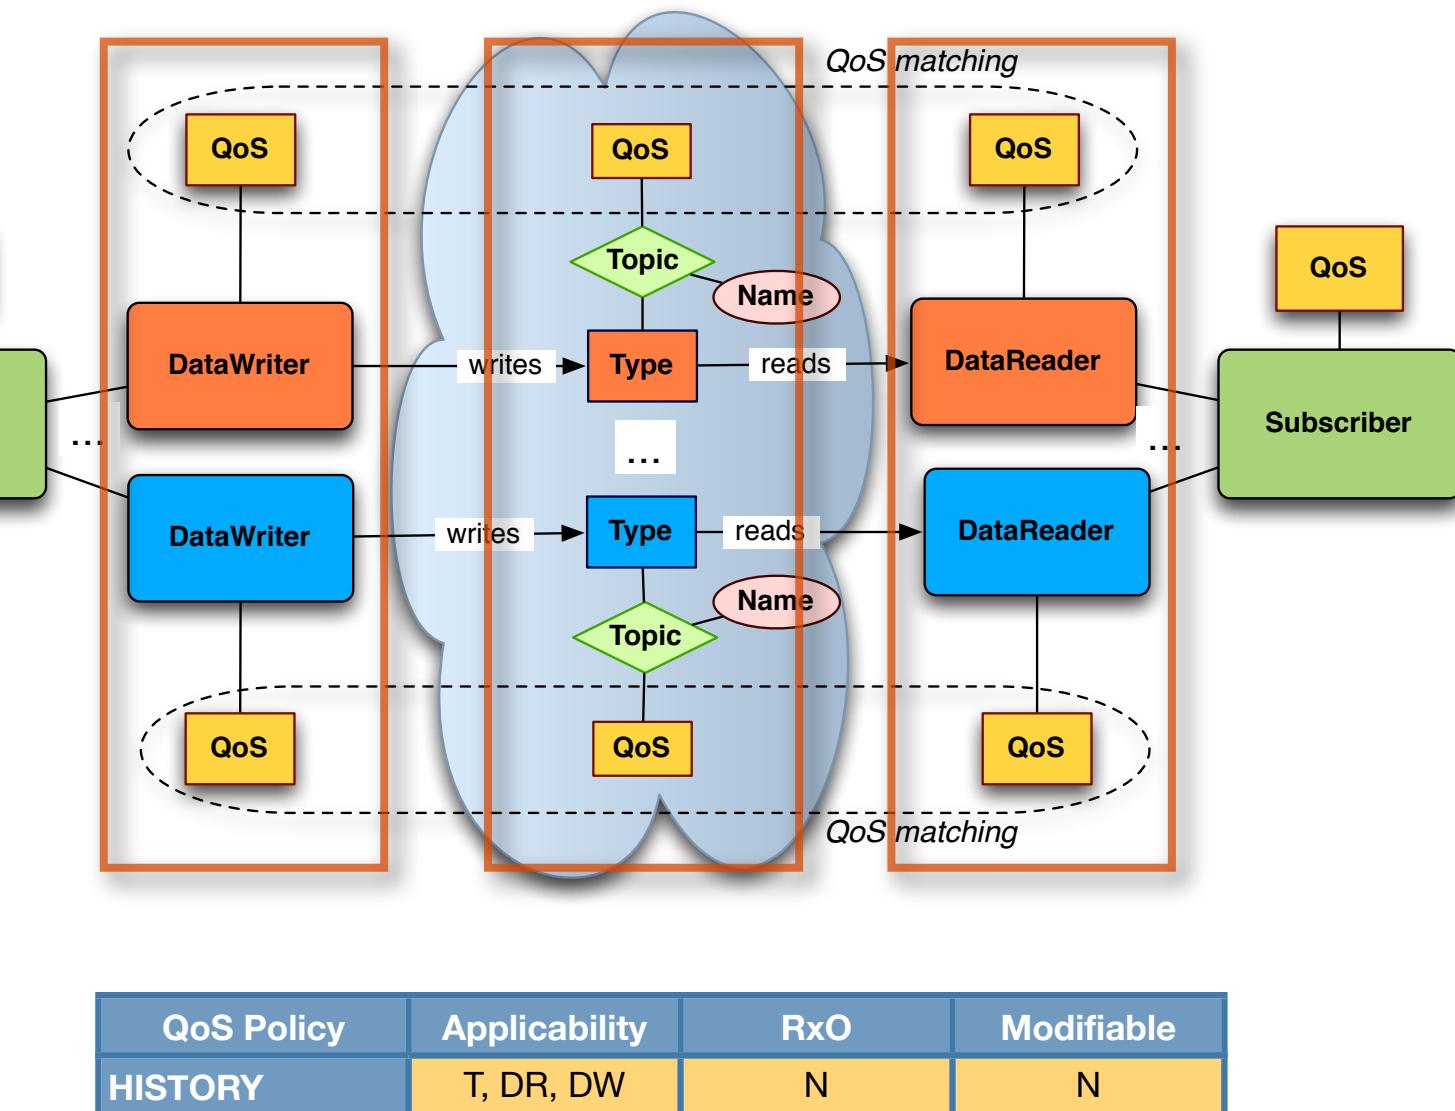
# History

- ▶ The History QoS Controls the number of samples-per-instance that will be stored by the middleware on behalf of a Reader
- ▶ **Keep Last K.** The History QoS can be set so to always have the latest **K** samples
- ▶ **Keep All.** The History QoS can be set so keep all samples produced by the writer and not yet taken, until resource limits are not reached



History Depth = 1 (DDS Default)

History Depth = 5



# OpenSplice|DDS

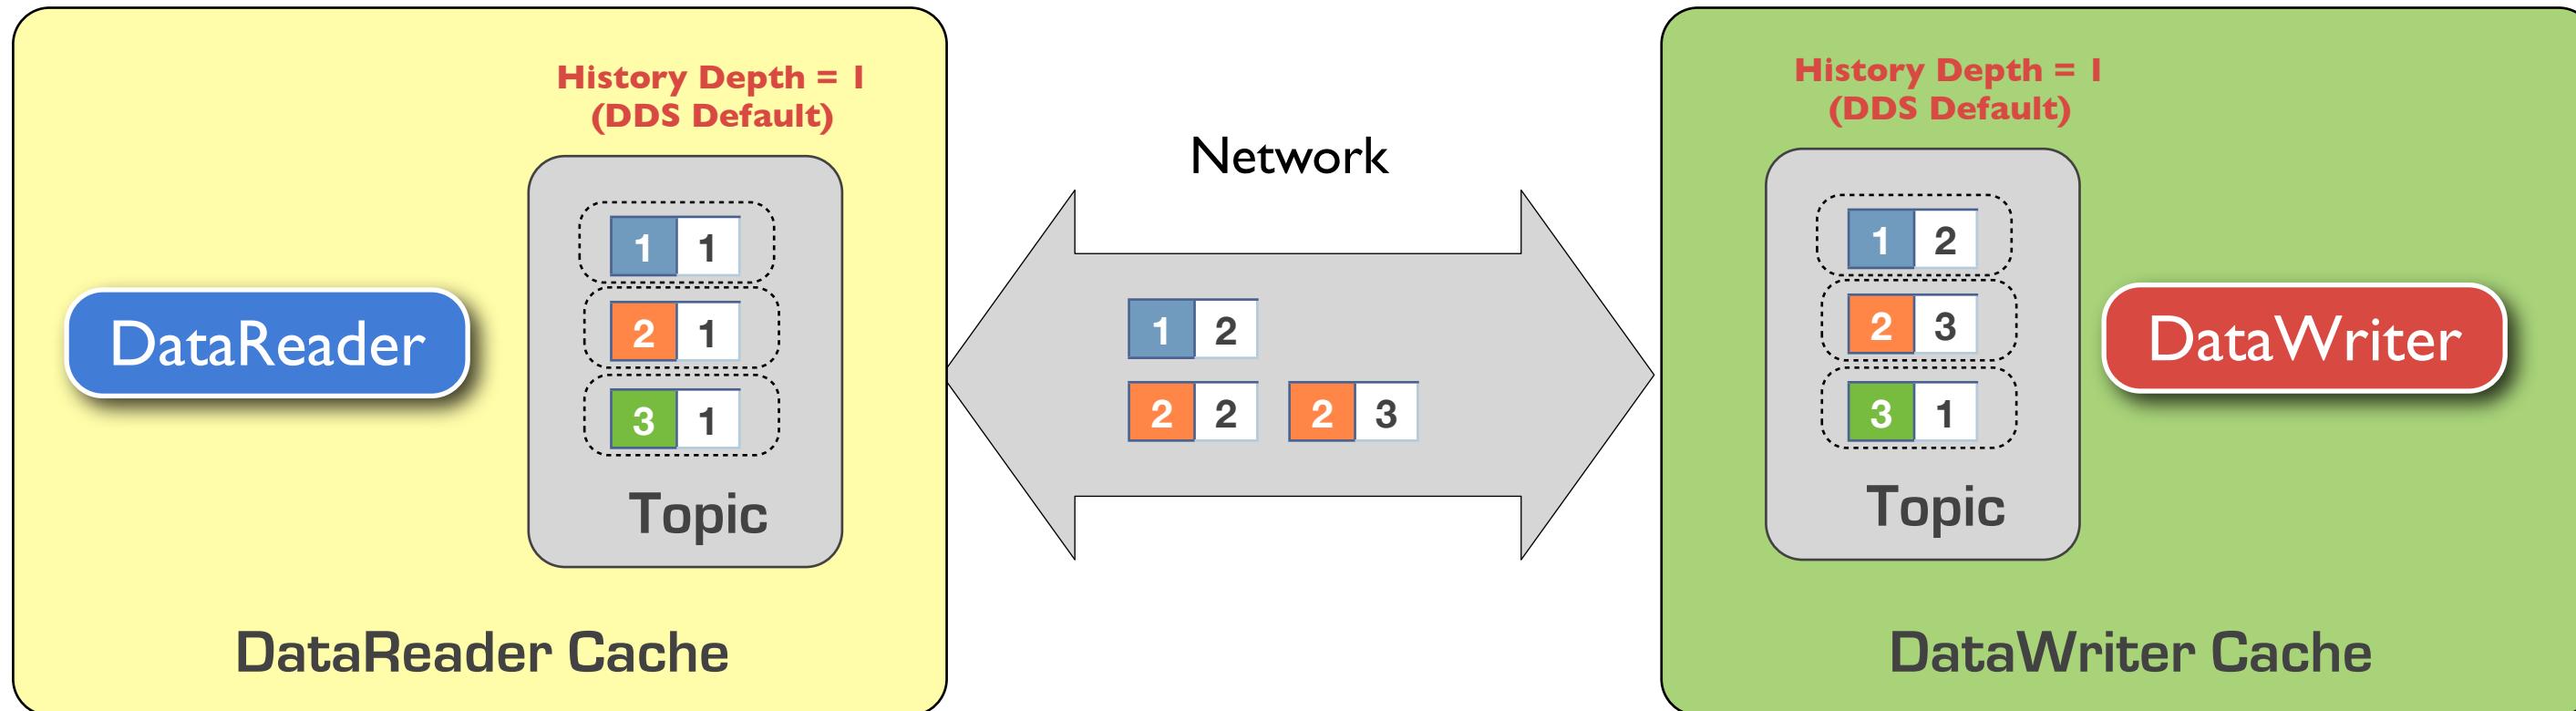
Delivering Performance, Openness, and Freedom

**Topic Used in next section.**

```
struct Counter {  
    int cID;  
    int count;  
};  
#pragma keylist Counter cID
```

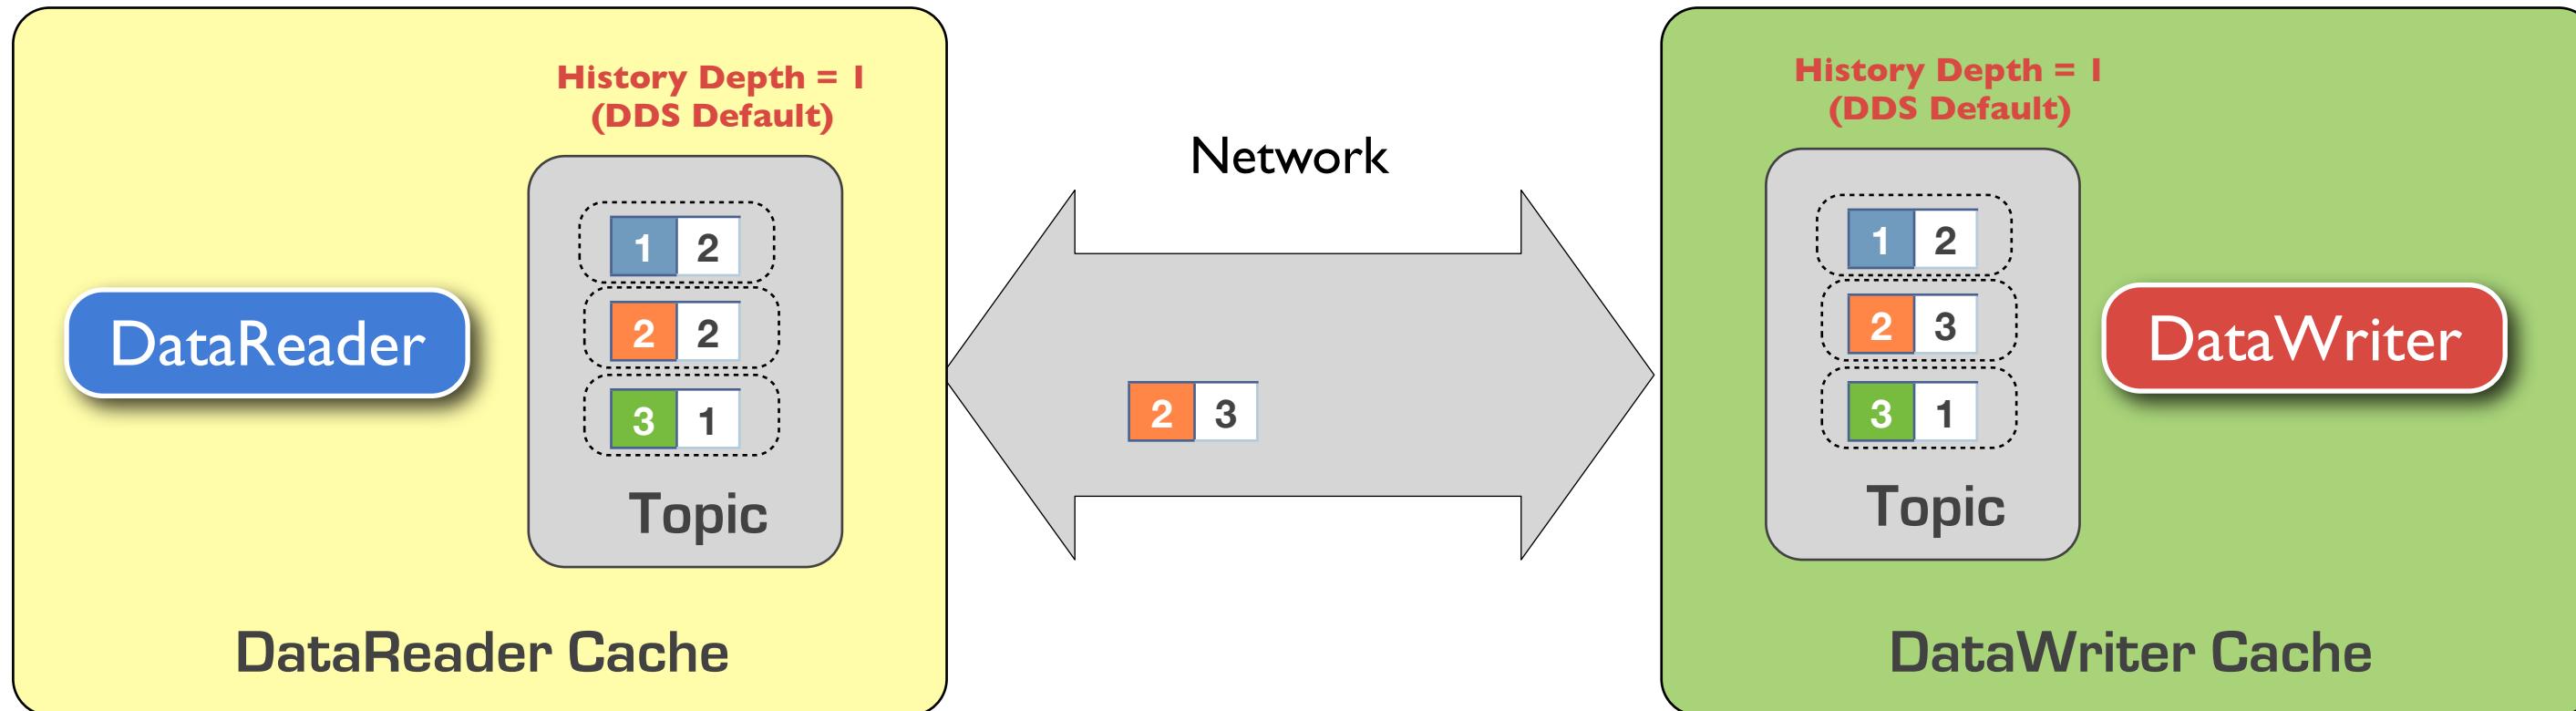
## History in Action

# History in Action



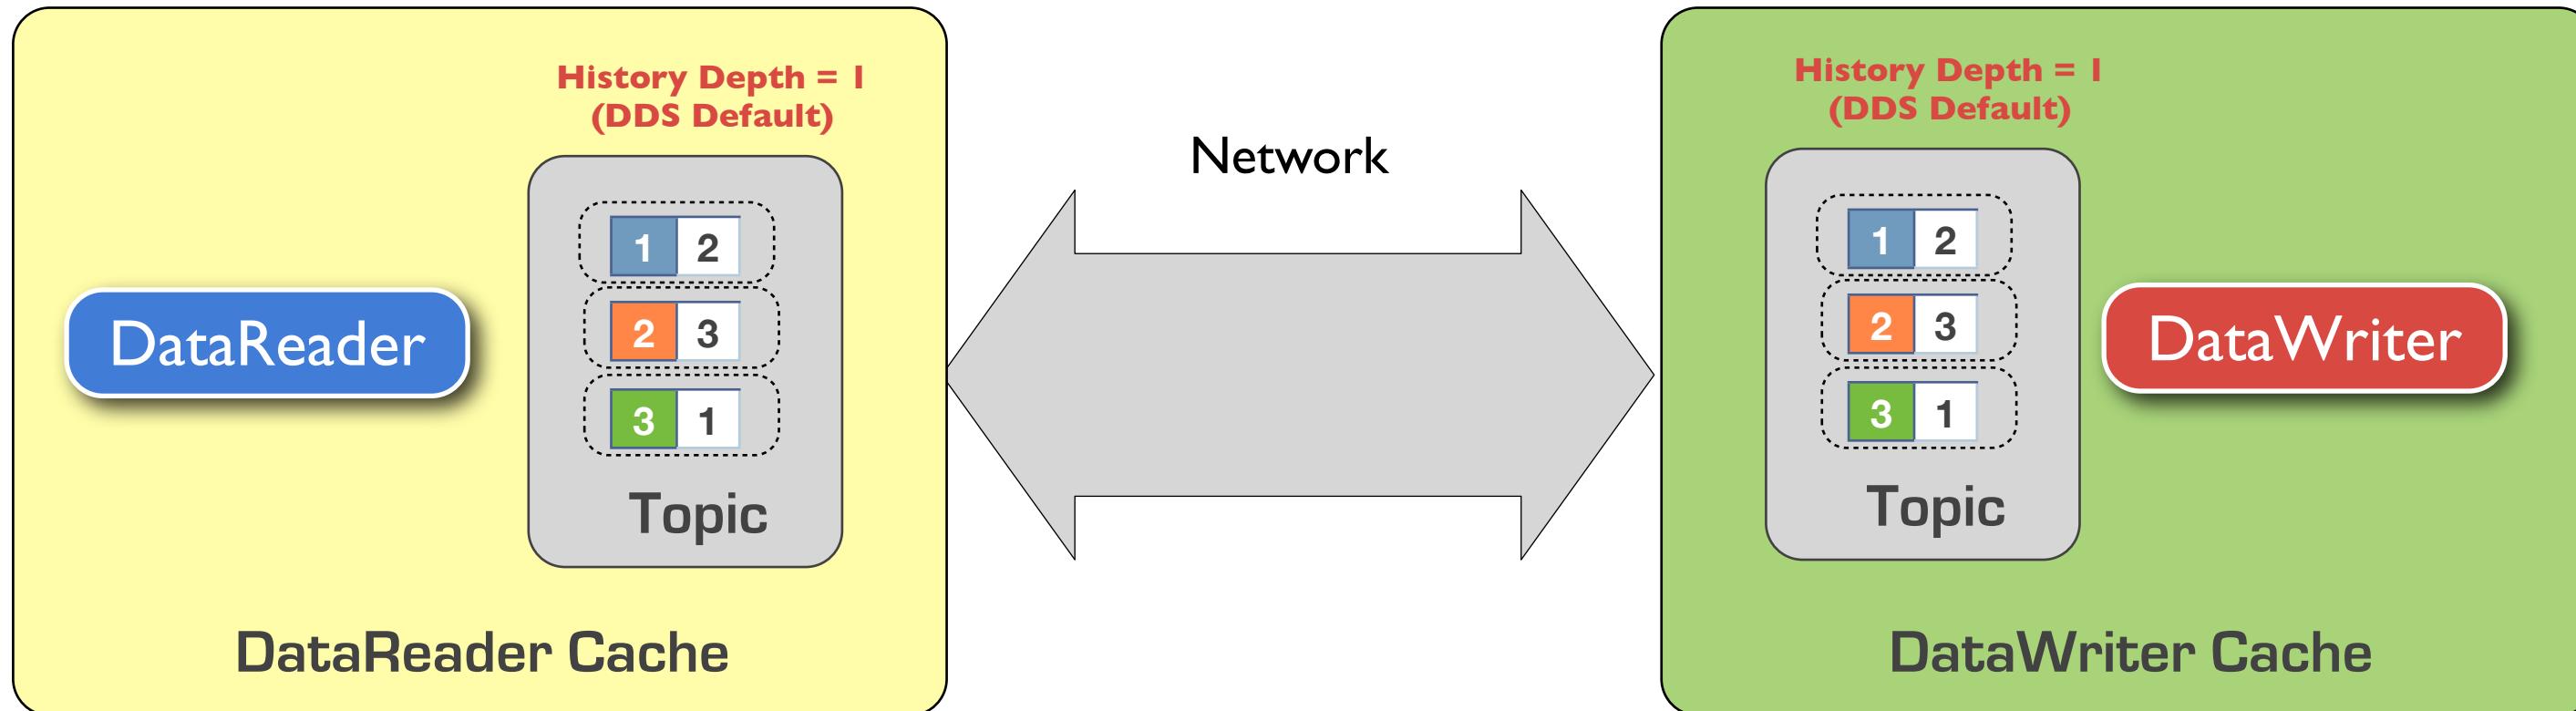
**Note:** The Reliability QoS controls whether data is sent reliably, or best-effort, from the DataWriter to matched DataReaders

# History in Action



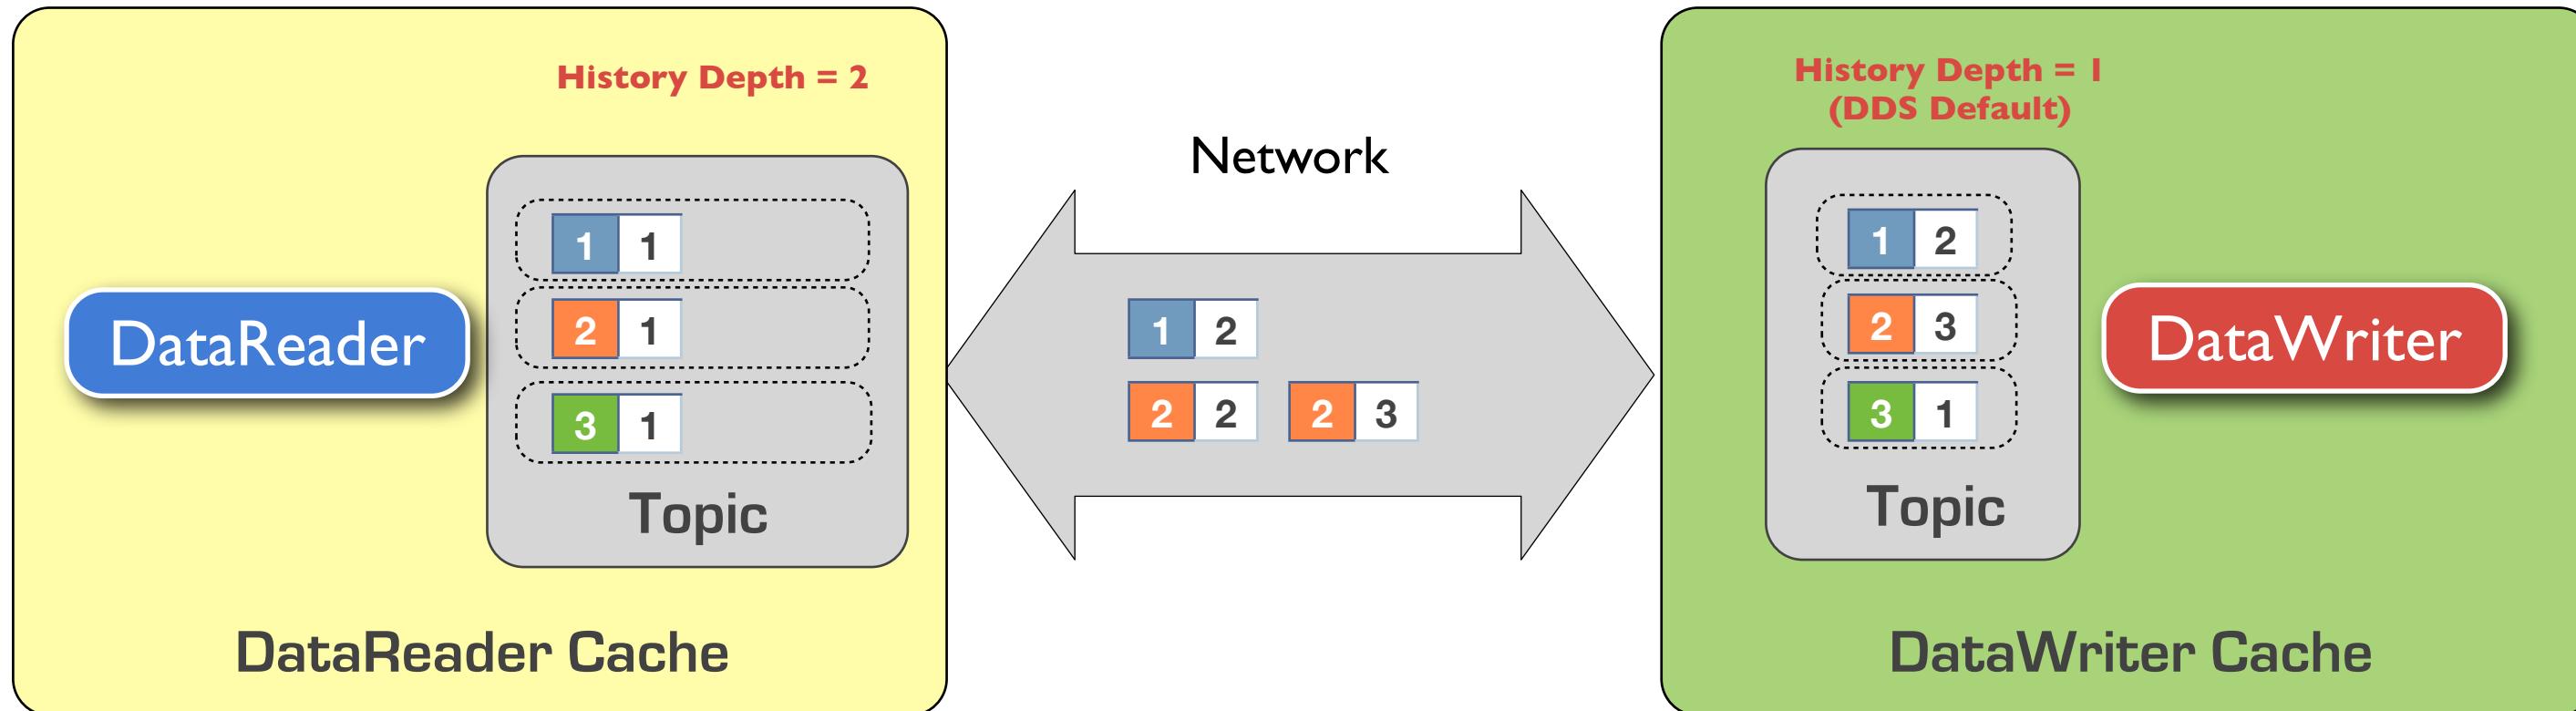
**Note:** The Reliability QoS controls whether data is sent reliably, or best-effort, from the DataWriter to matched DataReaders

# History in Action



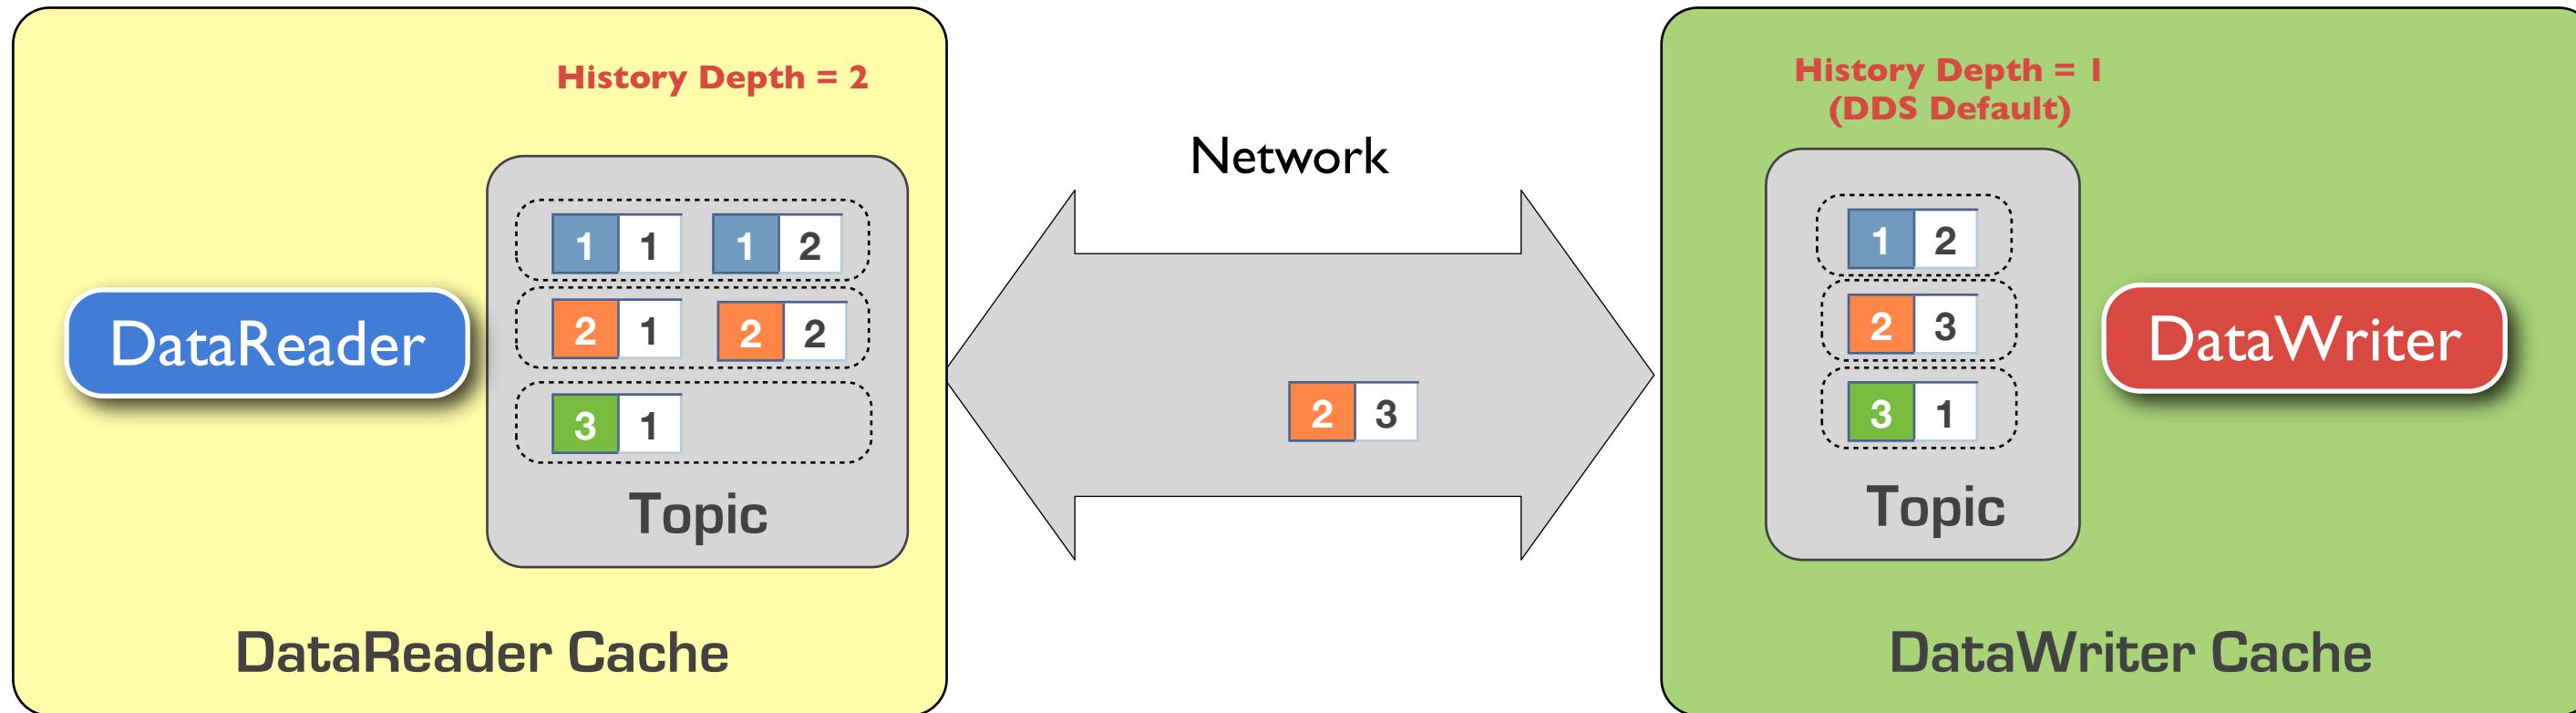
**Note:** The Reliability QoS controls whether data is sent reliably, or best-effort, from the DataWriter to matched DataReaders

# History in Action



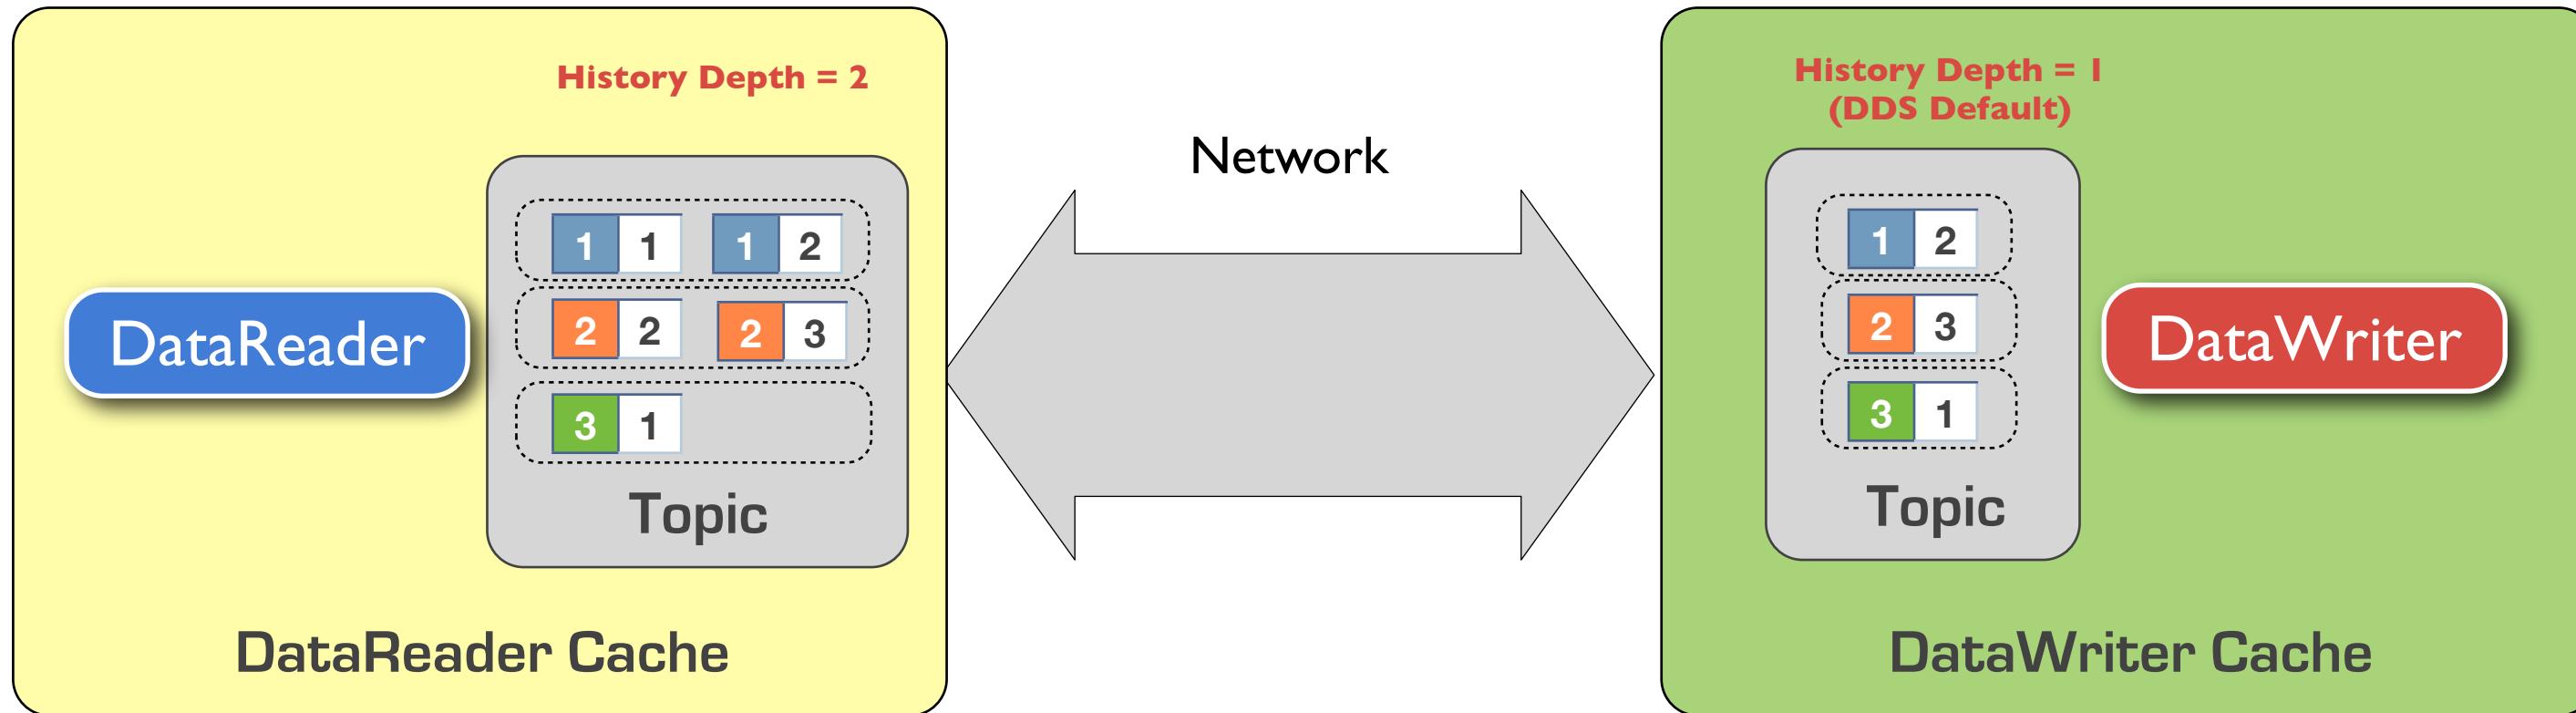
**Note:** The Reliability QoS controls whether data is sent reliably, or best-effort, from the DataWriter to matched DataReaders

# History in Action



**Note:** The Reliability QoS controls whether data is sent reliably, or best-effort, from the DataWriter to matched DataReaders

# History in Action

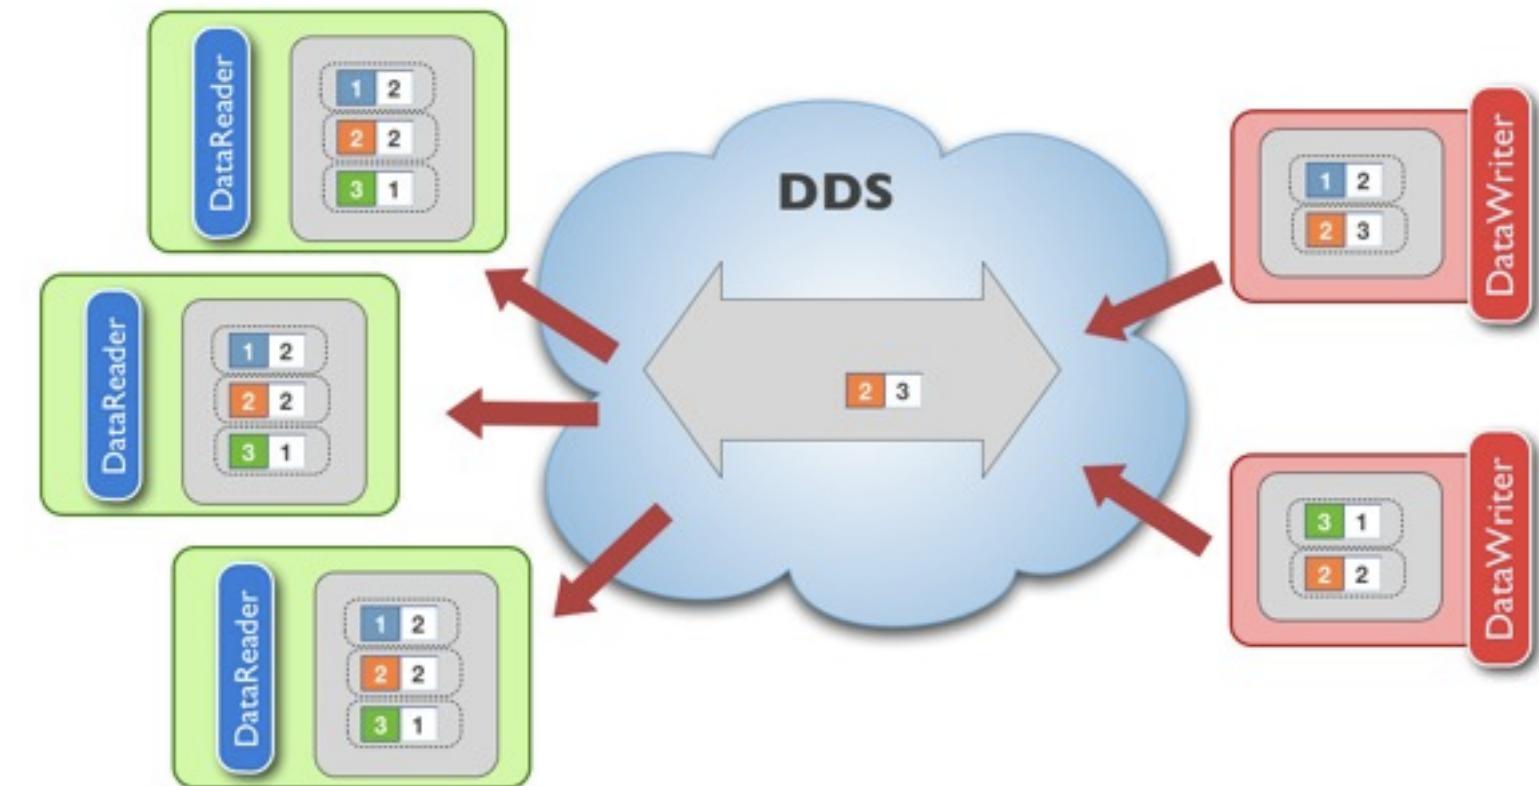


**Note:** The Reliability QoS controls whether data is sent reliably, or best-effort, from the DataWriter to matched DataReaders

# Putting it All Together

The reliability with which data is delivered to applications is impacted in DDS by the following qualities of service

- ▶ RELIABILITY
  - ▶ BEST\_EFORT
  - ▶ RELIABLE
- ▶ HISTORY
  - ▶ KEEP\_LAST (K)
  - ▶ KEEP\_ALL
- ▶ Theoretically, the only way to assure that **an application will see all the samples** produced by a writer is to use **RELIABLE+KEEP\_ALL**. Any other combination could induce to samples being discarded on the receiving side because of the HISTORY depth



# OpenSplice|DDS

Delivering Performance, Openness, and Freedom

Controlling Real-Time  
Properties

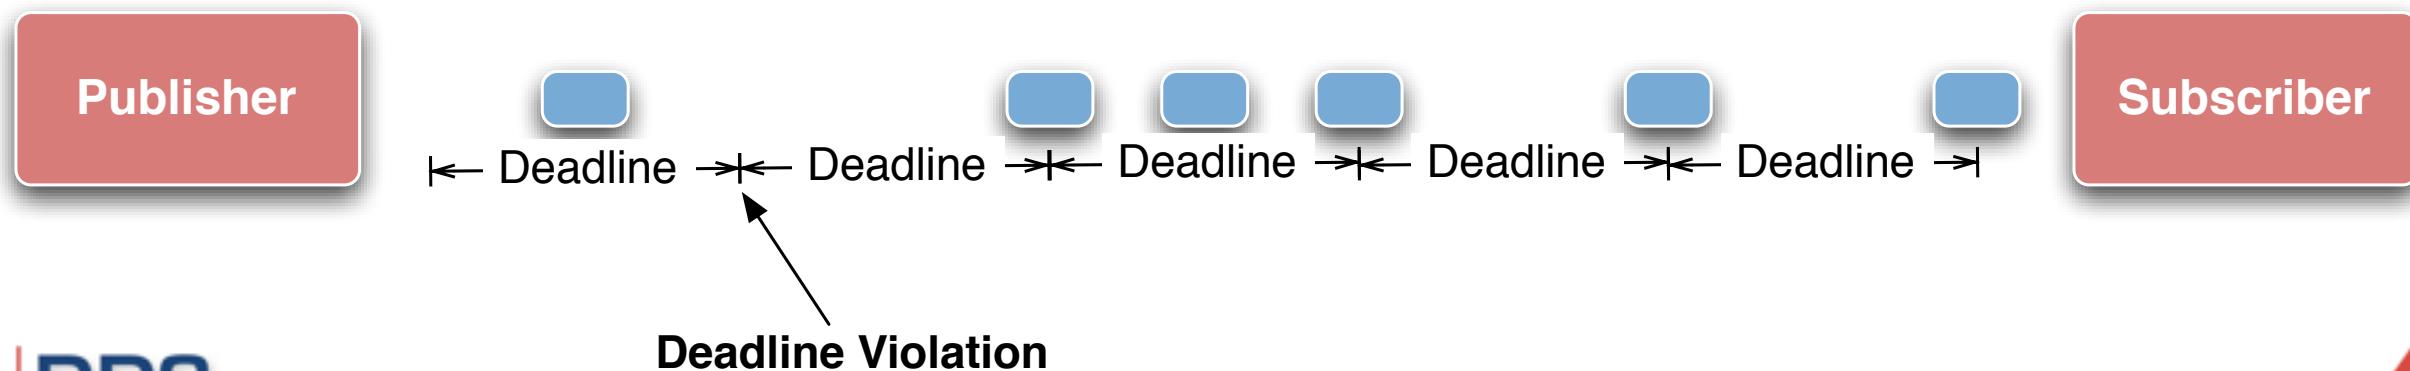
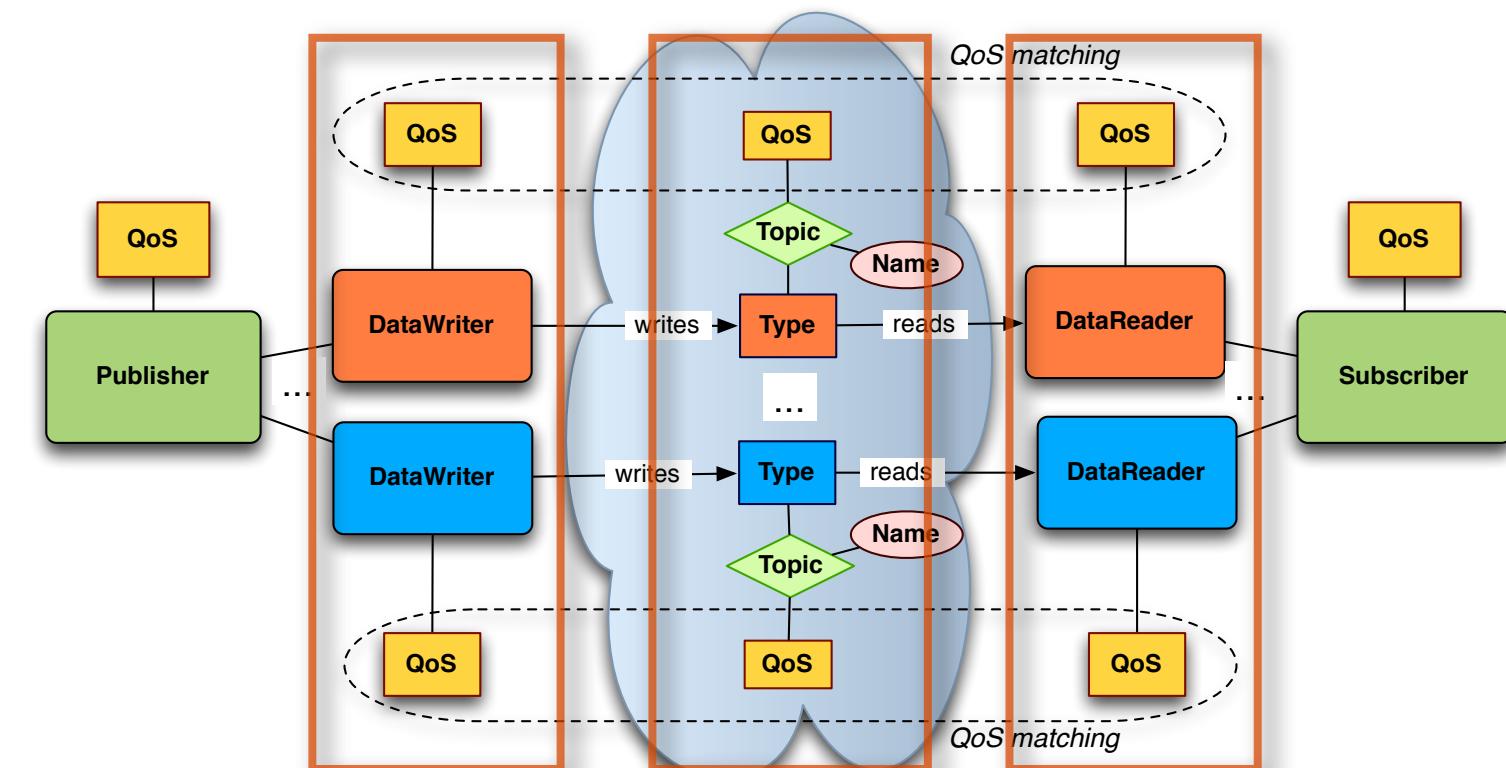
# Deadline

You can't be later than...

The **DEADLINE** QoS policy allows to define the **maximum inter-arrival time** between data samples

QoS Policy	Applicability	RxO	Modifiable
<b>DEADLINE</b>	T, DR, DW	Y	Y

- ▶ DataWriter indicates that the application commits to write a new value at least once every deadline period
- ▶ DataReaders are notified by the DDS when the DEADLINE QoS contract is violated

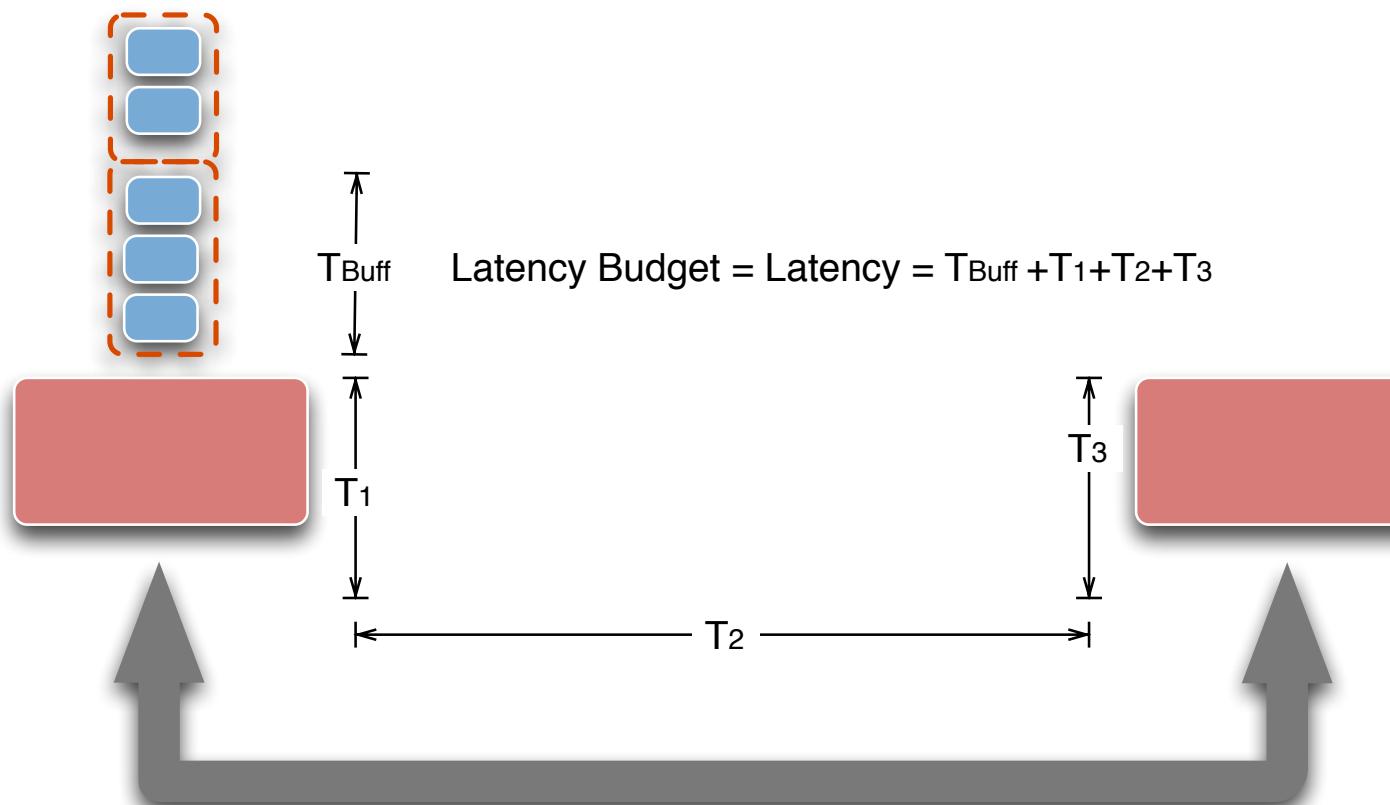


# Latency Budget

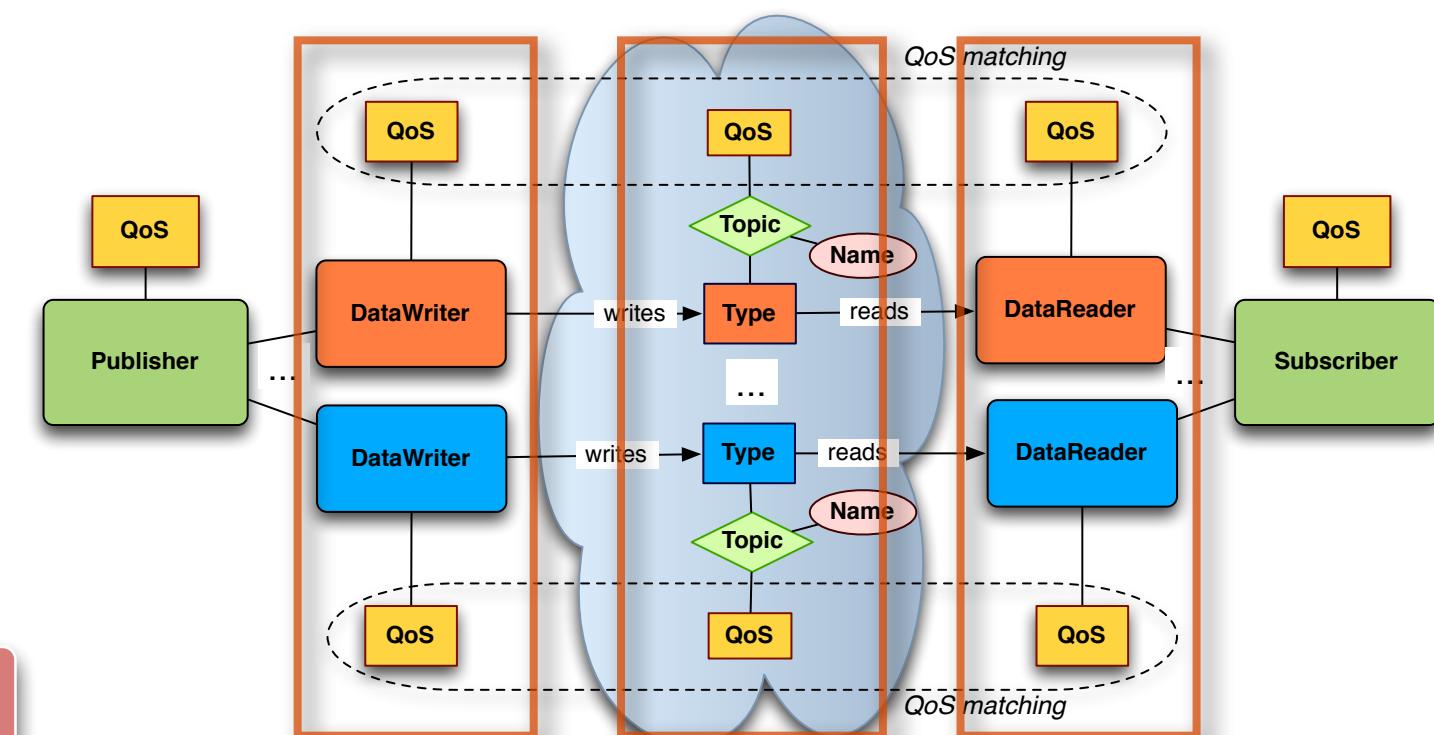
I need to get there in at most...

The **LATENCY\_BUDGET QoS** policy specifies the maximum acceptable delay from the time the data is written until the data is inserted in the receiver's application-cache

- ▶ The default value of the duration is zero indicating that the delay should be minimized
- ▶ This policy is a hint to the DDS, not something that must be monitored or enforced.



QoS Policy	Applicability	RxO	Modifiable
LATENCY_BUDGET	T, DR, DW	Y	Y

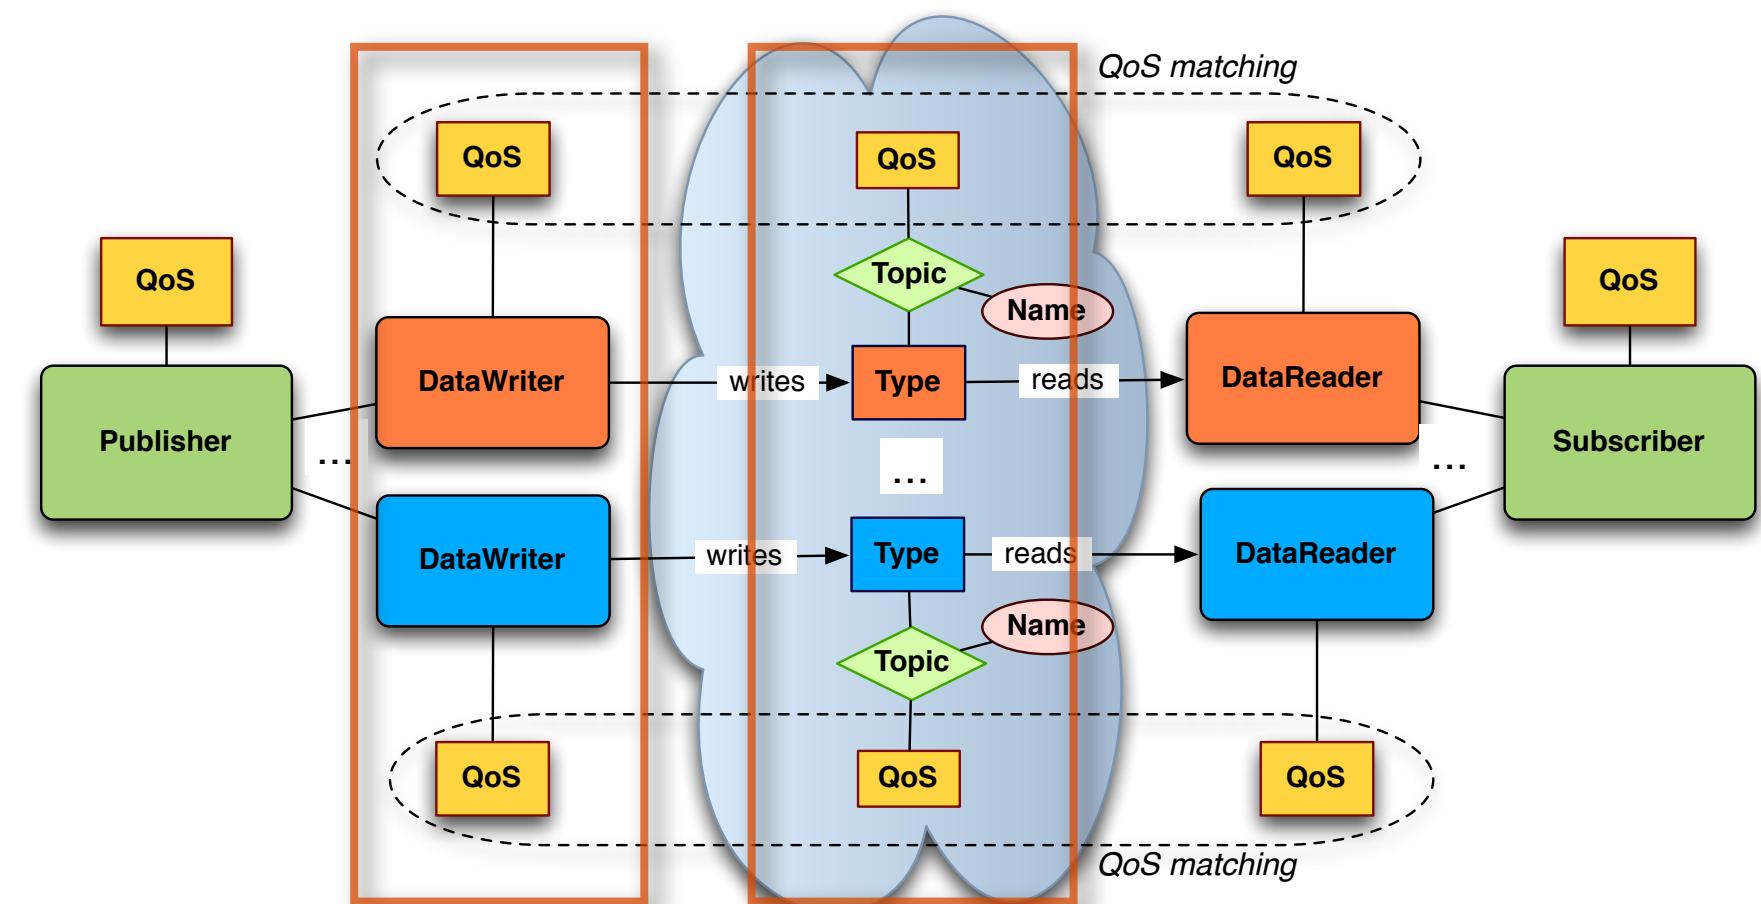


# Transport Priority

VIP Data, stay clear!

QoS Policy	Applicability	RxO	Modifiable
TRANSPORT PRIORITY	T, DW	-	Y

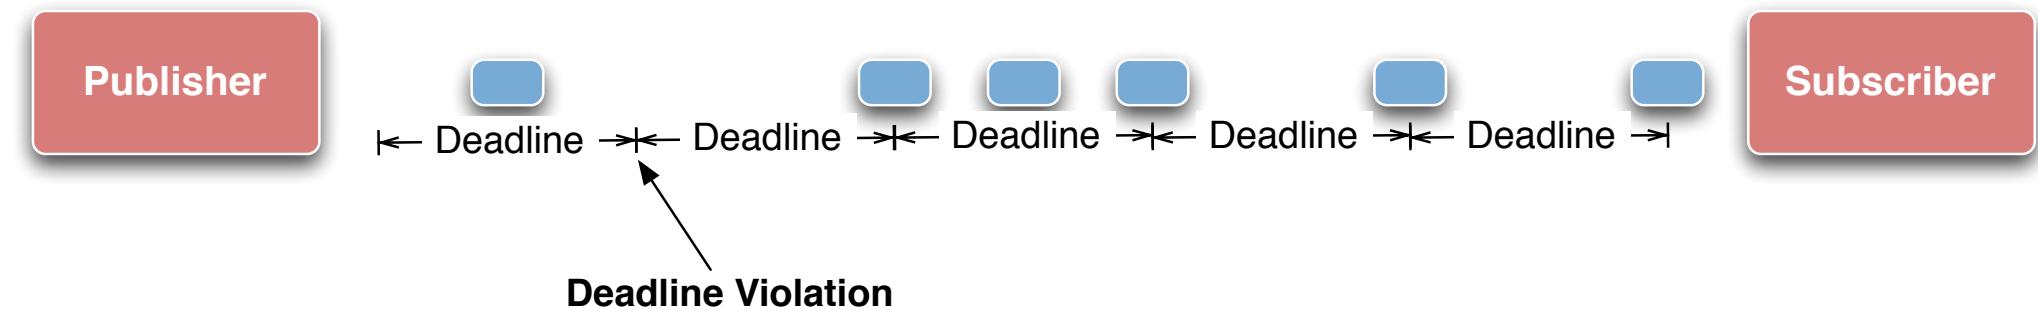
The **TRANSPORT\_PRIORITY** QoS policy is a hint to the infrastructure as to how to set the priority of the underlying transport used to send the data.



# Putting it all Together

The real-time properties with which data is delivered to applications is impacted in DDS by the following qualities of service:

- ▶ **TRANSPORT\_PRIORITY**
- ▶ **LATENCY\_BUDGET**
- ▶ In addition, DDS provides means for detecting performance failure, e.g., Deadline miss, by means of the **DEADLINE** QoS
- ▶ Given a periodic task-set  $\{T\}$  with periods  $D_i$  (with  $D_i < D_{i+1}$ ) and deadline equal to the period, than QoS should be set as follows:
  - ▶ Assign to each task  $T_i$  a TRANSPORT\_PRIORITY  $P_i$  such that  $P_i > P_{i+1}$
  - ▶ Set for each task  $T_i$  a DEADLINE QoS of  $D_i$
  - ▶ For maximizing throughput and minimizing resource usage set for each  $T_i$  a LATENCY\_BUDGET QoS between  $D_i/2$  and  $D_i/3$  (this is a rule of thumb, the upper bound is  $D_i - (RTT/2)$ )



# OpenSplice|DDS

Delivering Performance, Openness, and Freedom

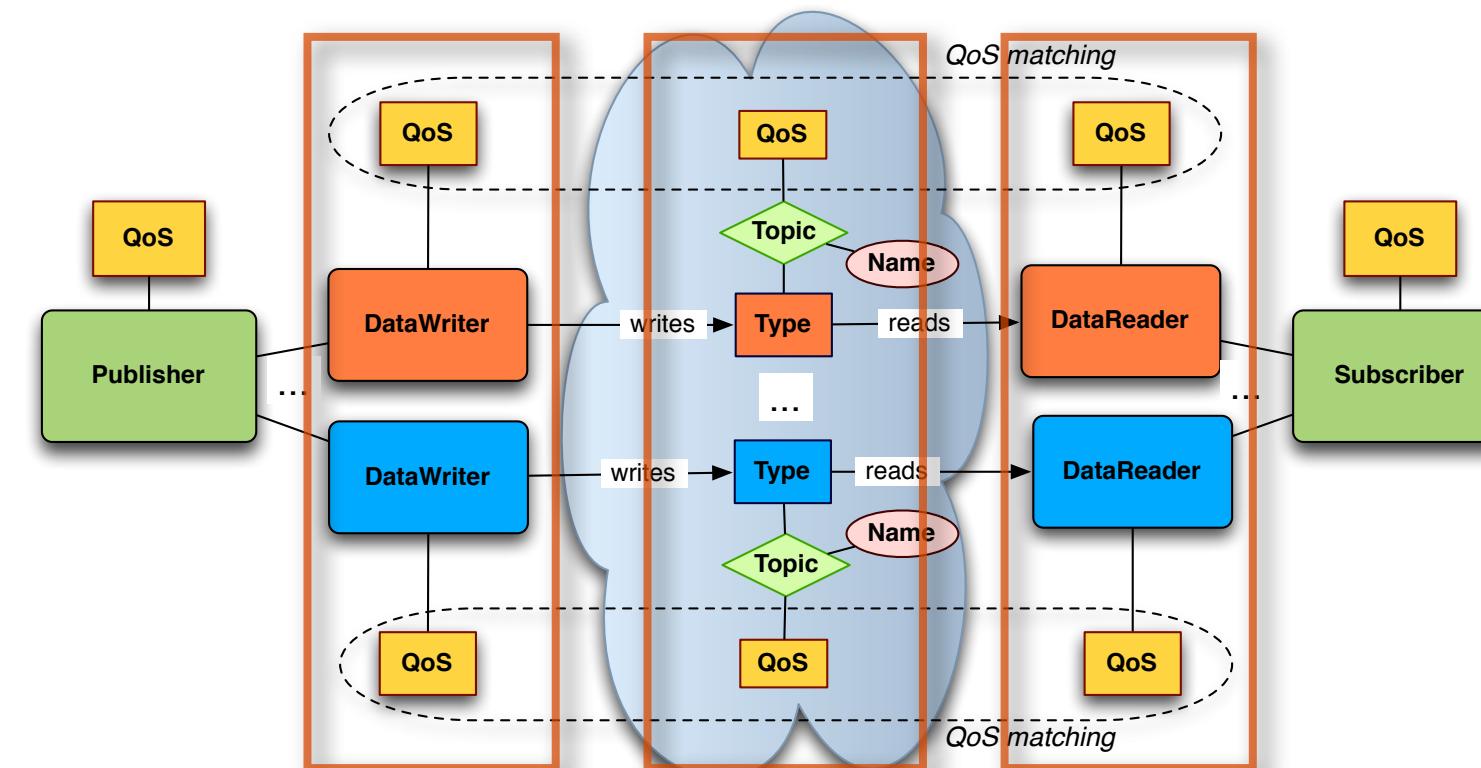
Controlling the  
Consistency Model

# Durability

The **DURABILITY** QoS controls the data availability w.r.t. late joiners, specifically the DDS provides the following variants:

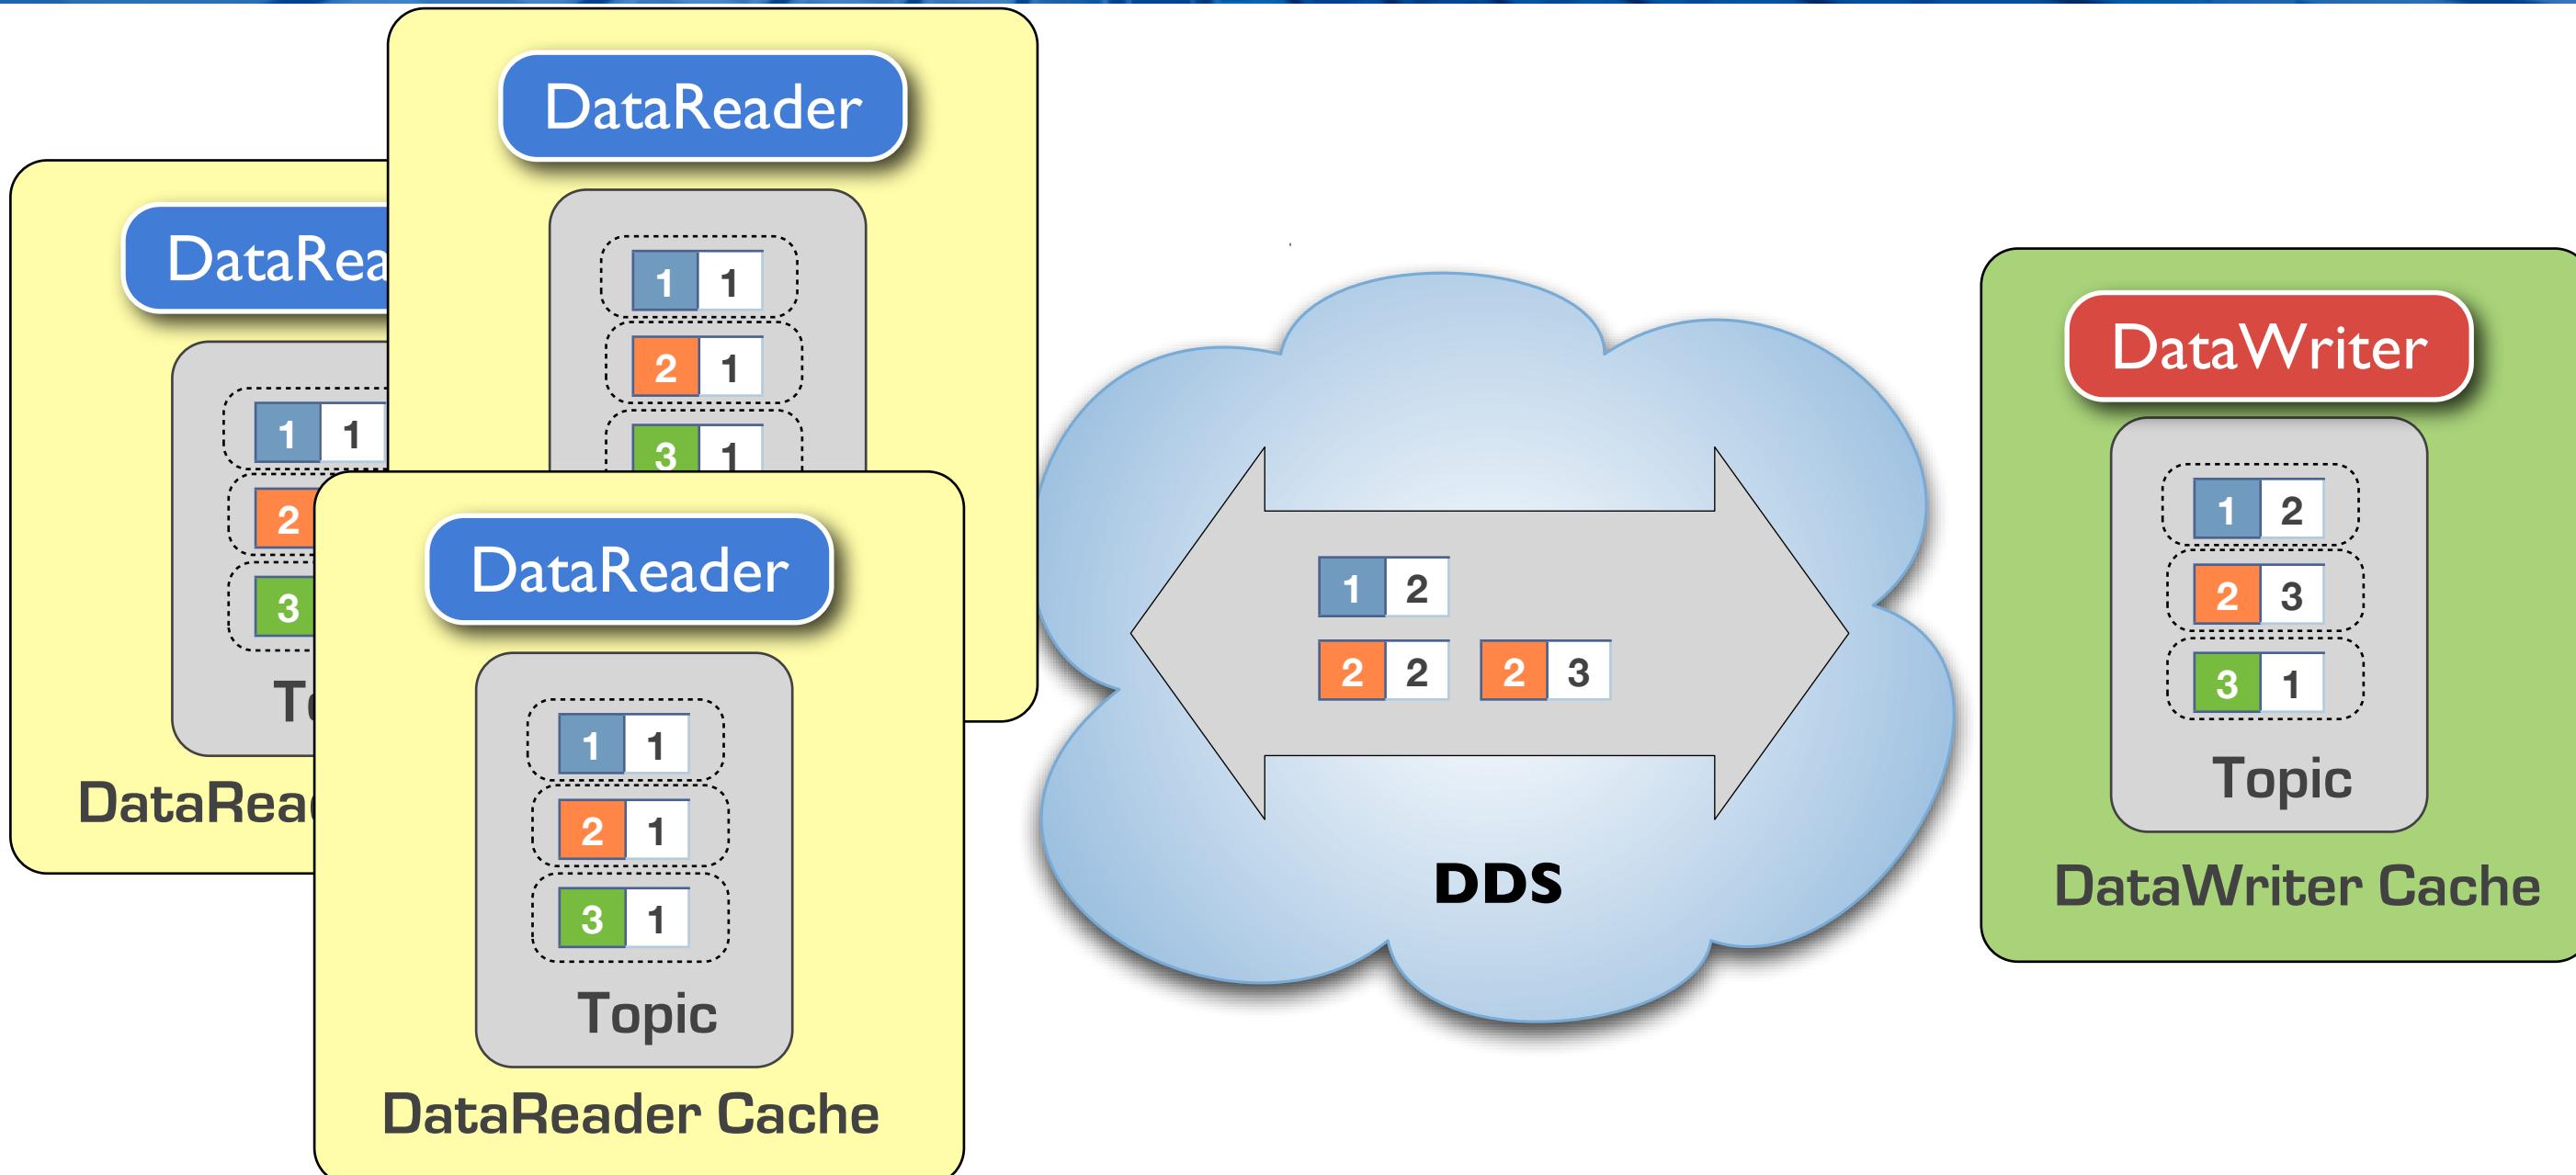
- ▶ **Volatile.** No need to keep data instances for late joining data readers
- ▶ **Transient Local.** Data instance availability for late joining data reader is tied to the data writer availability
- ▶ **Transient.** Data instance availability outlives the data writer
- ▶ **Persistent.** Data instance availability outlives system restarts

QoS Policy	Applicability	RxO	Modifiable
DURABILITY	T, DR, DW	Y	N
DURABILITY SERVICE	T, DW	N	N



The **DURABILITY\_SERVICE** QoS provide control over configuration of the service that implements the transient and persistent durability features

# Eventual Consistency & R/W Caches



Under an Eventual Consistency Model, DDS guarantees that all matched Reader Caches will eventually be identical of the respective Writer Cache

# QoS Impacting the Consistency Model

The DDS Consistency Model is a property that can be associated to Topics or further refined by Reader/Writers. The property is controlled by the following QoS Policies:

- ▶ **DURABILITY**
  - ▶ VOLATILE | TRANSIENT\_LOCAL | TRANSIENT | PERSISTENT
- ▶ **LIFESPAN**
- ▶ **RELIABILITY**
  - ▶ RELIABLE | BEST\_EFFORT
- ▶ **DESTINATION ORDER**
  - ▶ SOURCE\_TIMESTAMP | DESTINATION\_TIMESTAMP

QoS Policy	Applicability	RxO	Modifiable
DURABILITY	T, DR, DW	Y	N
LIFESPAN	T, DW	-	Y
RELIABILITY	T, DR, DW	Y	N
DESTINATION ORDER	T, DR, DW	Y	N

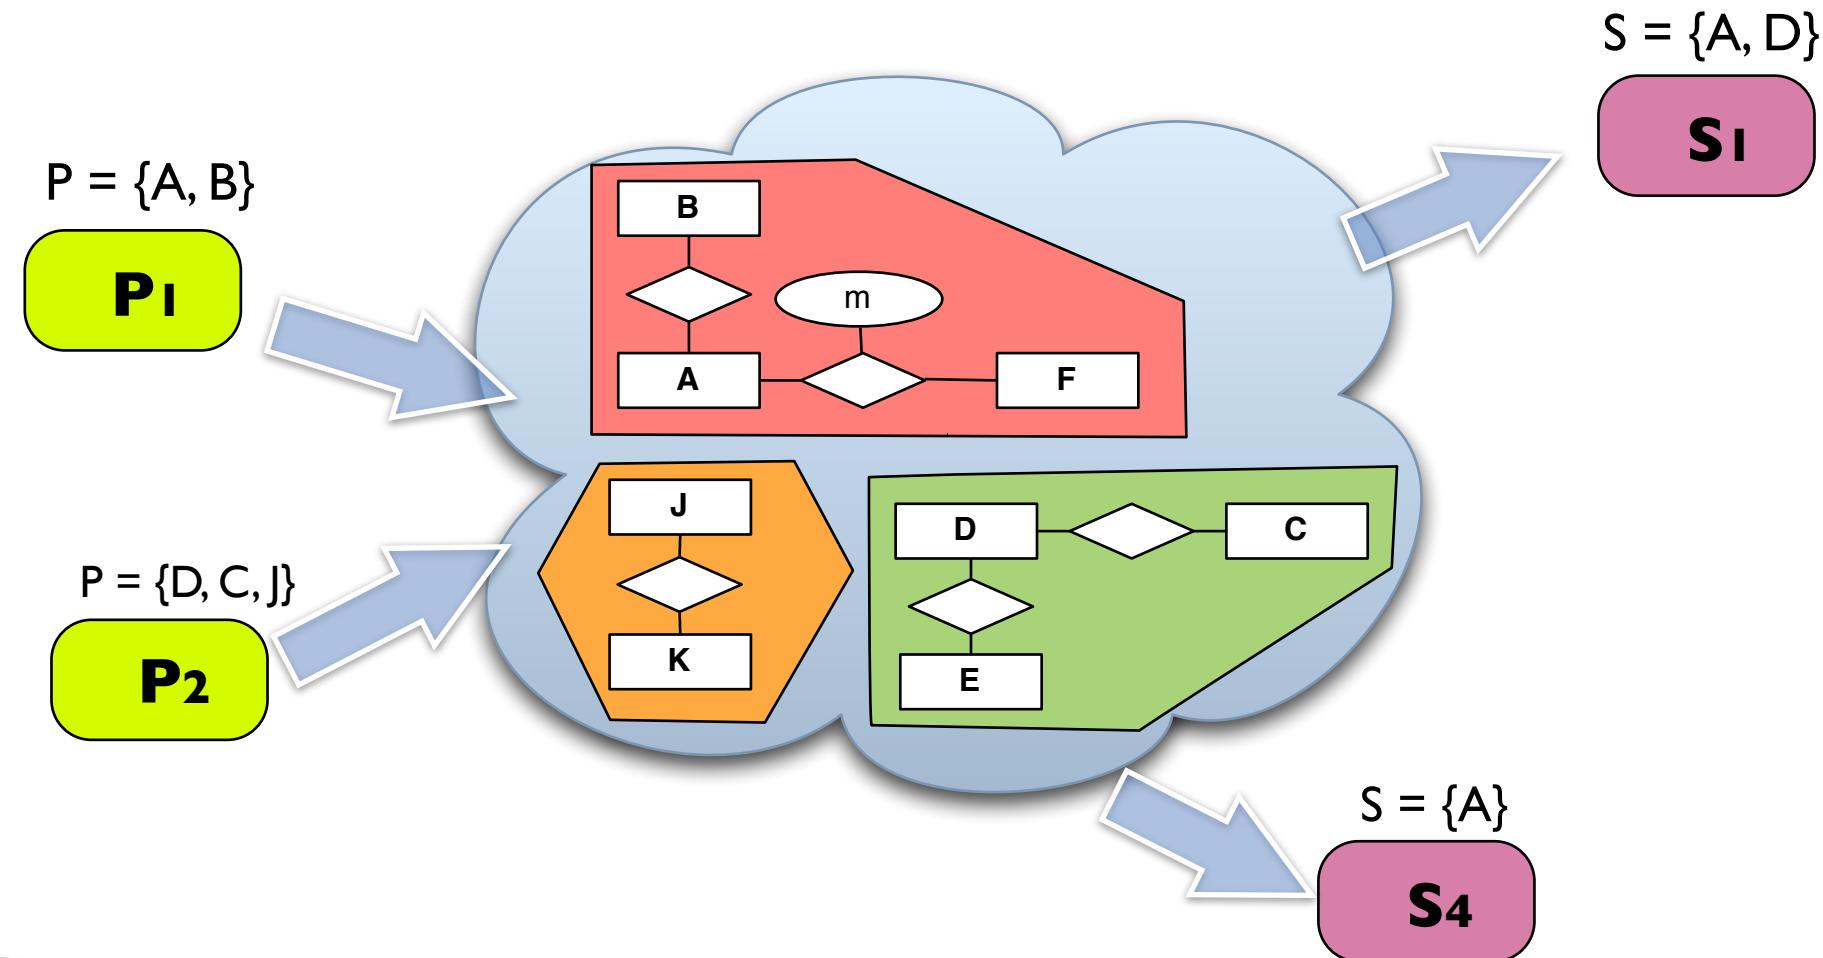
# QoS Impacting the Consistency Model

	DURABILITY	RELIABILITY	DESTINATION_ORDER	LIFESPAN
<b>Eventual Consistency (No Crash / Recovery)</b>	VOLATILE	RELIABLE	SOURCE_TIMESTAMP	INF.
<b>Eventual Consistency (Reader Crash / Recovery)</b>	TRANSIENT_LOCAL	RELIABLE	SOURCE_TIMESTAMP	INF.
<b>Eventual Consistency (Crash/Recovery)</b>	TRANSIENT	RELIABLE	SOURCE_TIMESTAMP	INF.
<b>Eventual Consistency (Crash/Recovery)</b>	PERSISTENT	RELIABLE	SOURCE_TIMESTAMP	INF.
<b>Weak Consistency</b>	ANY	ANY	DESTINATION_TIMESTAMP	ANY
<b>Weak Consistency</b>	ANY	BEST EFFORT	ANY	ANY
<b>Weak Consistency</b>	ANY	ANY	ANY	N

# Eventual Consistency @ Work

	DURABILITY	RELIABILITY	DESTINATION_ORDER	LIFESPAN
Eventual Consistency (Reader Crash / Recovery)	TRANSIENT_LOCAL	RELIABLE	SOURCE_TIMESTAMP	INF.
Eventual Consistency (Crash/Recovery)	TRANSIENT	RELIABLE	SOURCE_TIMESTAMP	INF.
Weak Consistency	ANY	ANY	ANY	N

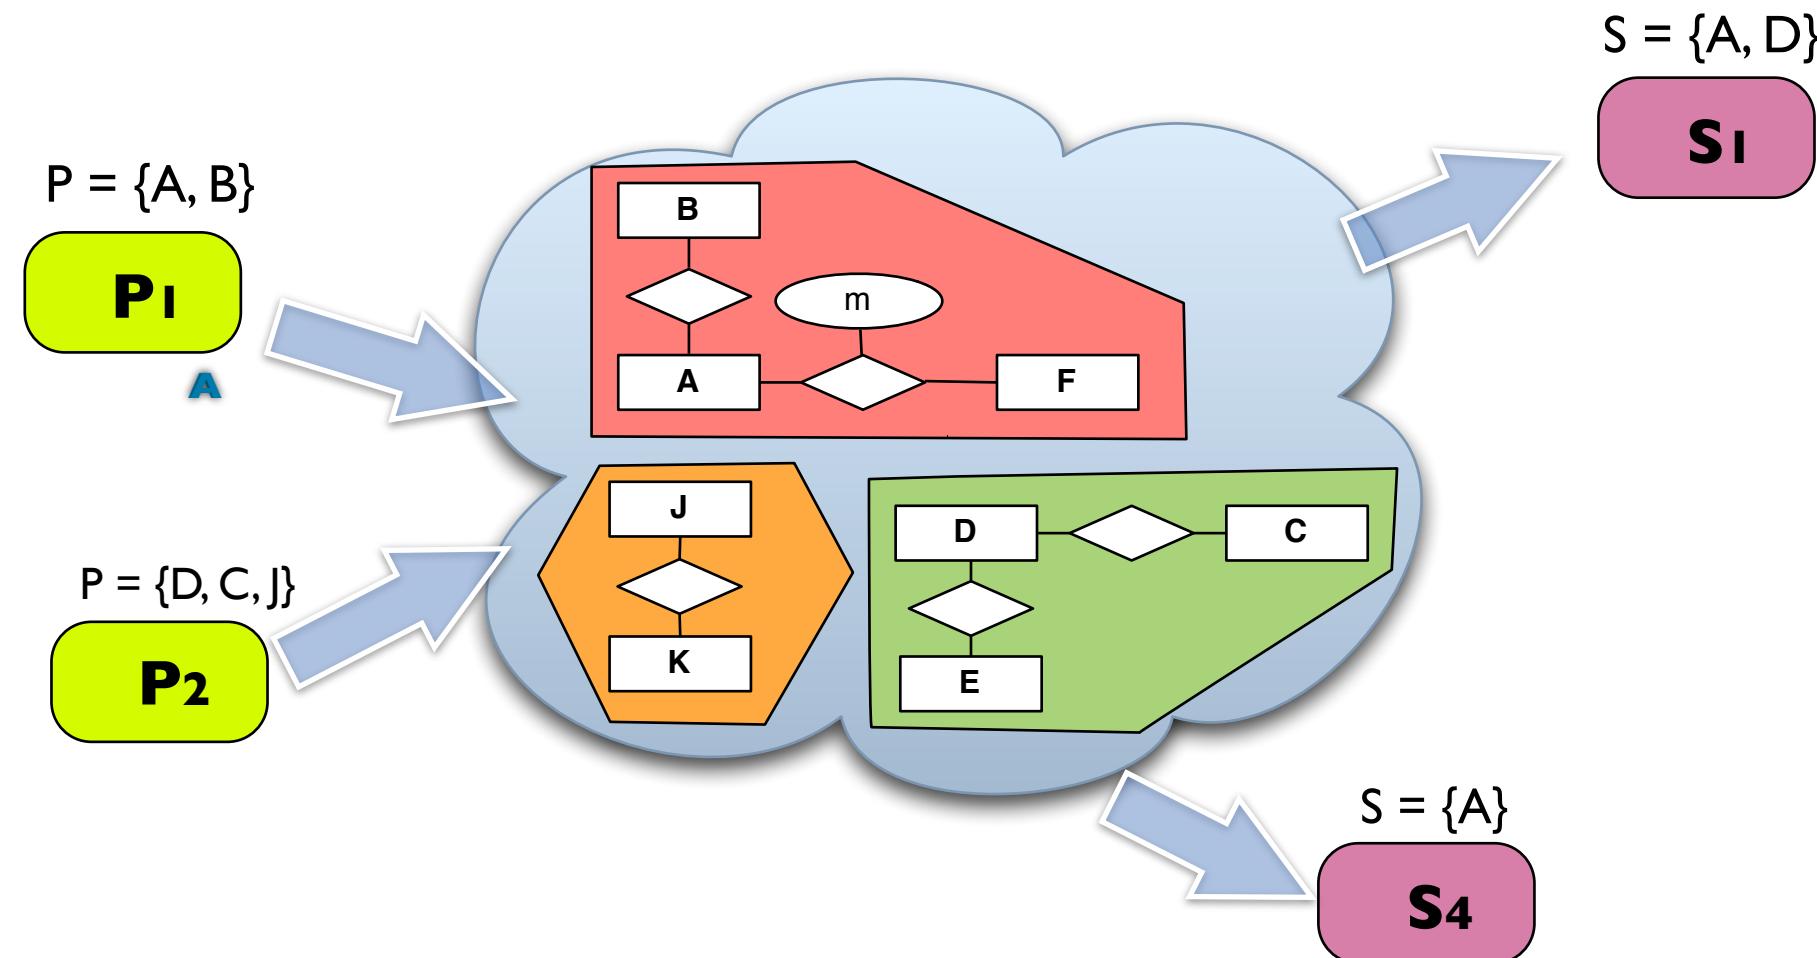
{A}  
{B}  
{J}



# Eventual Consistency @ Work

	DURABILITY	RELIABILITY	DESTINATION_ORDER	LIFESPAN
Eventual Consistency (Reader Crash / Recovery)	TRANSIENT_LOCAL	RELIABLE	SOURCE_TIMESTAMP	INF.
Eventual Consistency (Crash/Recovery)	TRANSIENT	RELIABLE	SOURCE_TIMESTAMP	INF.
Weak Consistency	ANY	ANY	ANY	N

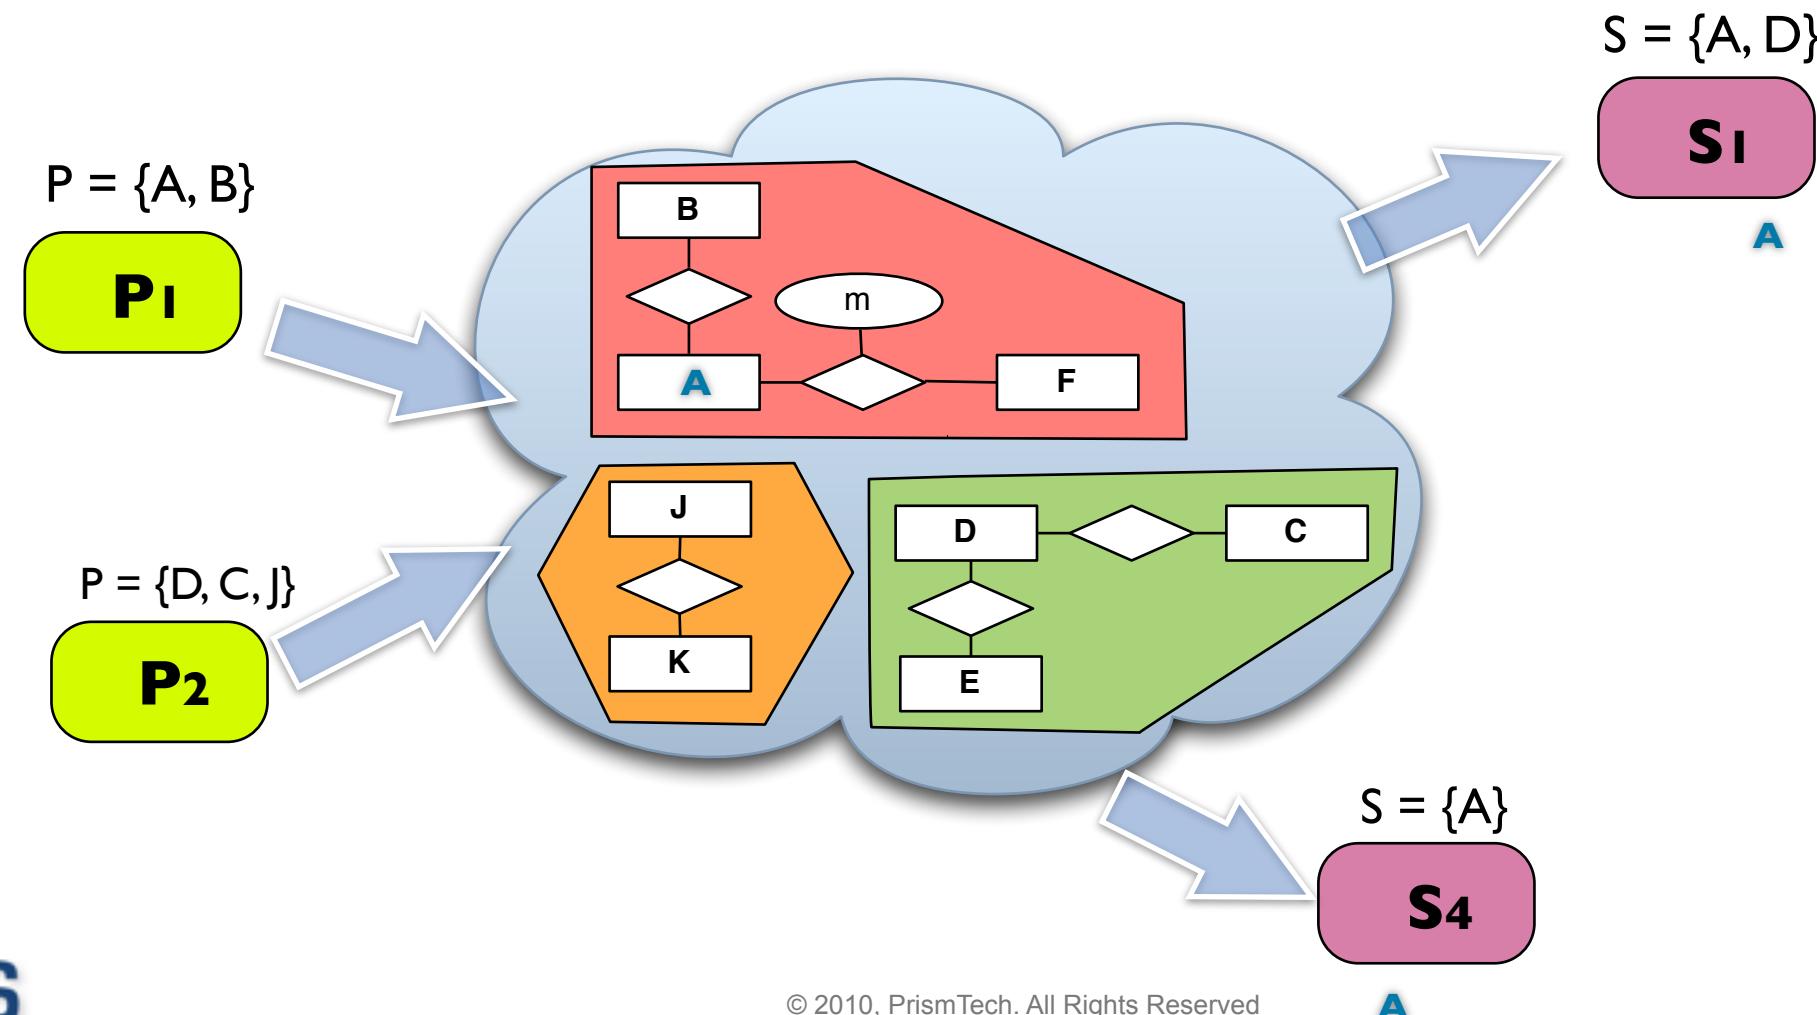
{A}  
{B}  
{J}



# Eventual Consistency @ Work

	DURABILITY	RELIABILITY	DESTINATION_ORDER	LIFESPAN
Eventual Consistency (Reader Crash / Recovery)	TRANSIENT_LOCAL	RELIABLE	SOURCE_TIMESTAMP	INF.
Eventual Consistency (Crash/Recovery)	TRANSIENT	RELIABLE	SOURCE_TIMESTAMP	INF.
Weak Consistency	ANY	ANY	ANY	N

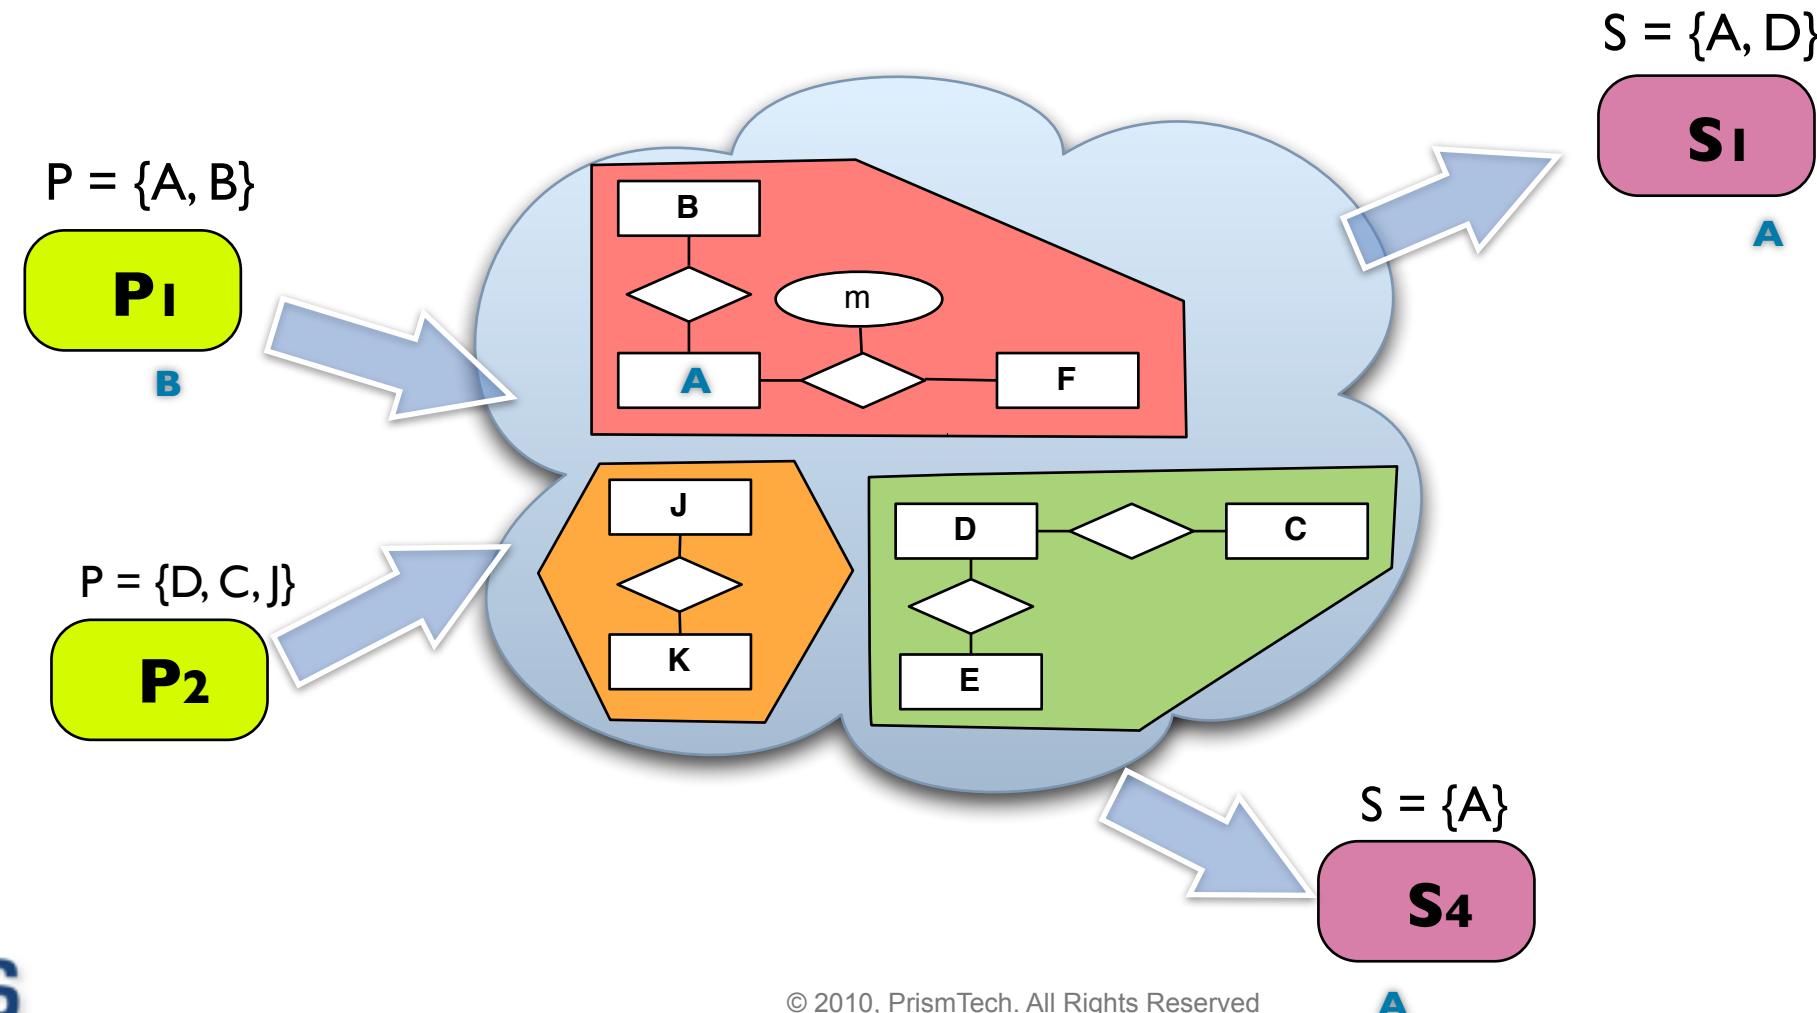
{A}  
{B}  
{J}



# Eventual Consistency @ Work

	DURABILITY	RELIABILITY	DESTINATION_ORDER	LIFESPAN
Eventual Consistency (Reader Crash / Recovery)	TRANSIENT_LOCAL	RELIABLE	SOURCE_TIMESTAMP	INF.
Eventual Consistency (Crash/Recovery)	TRANSIENT	RELIABLE	SOURCE_TIMESTAMP	INF.
Weak Consistency	ANY	ANY	ANY	N

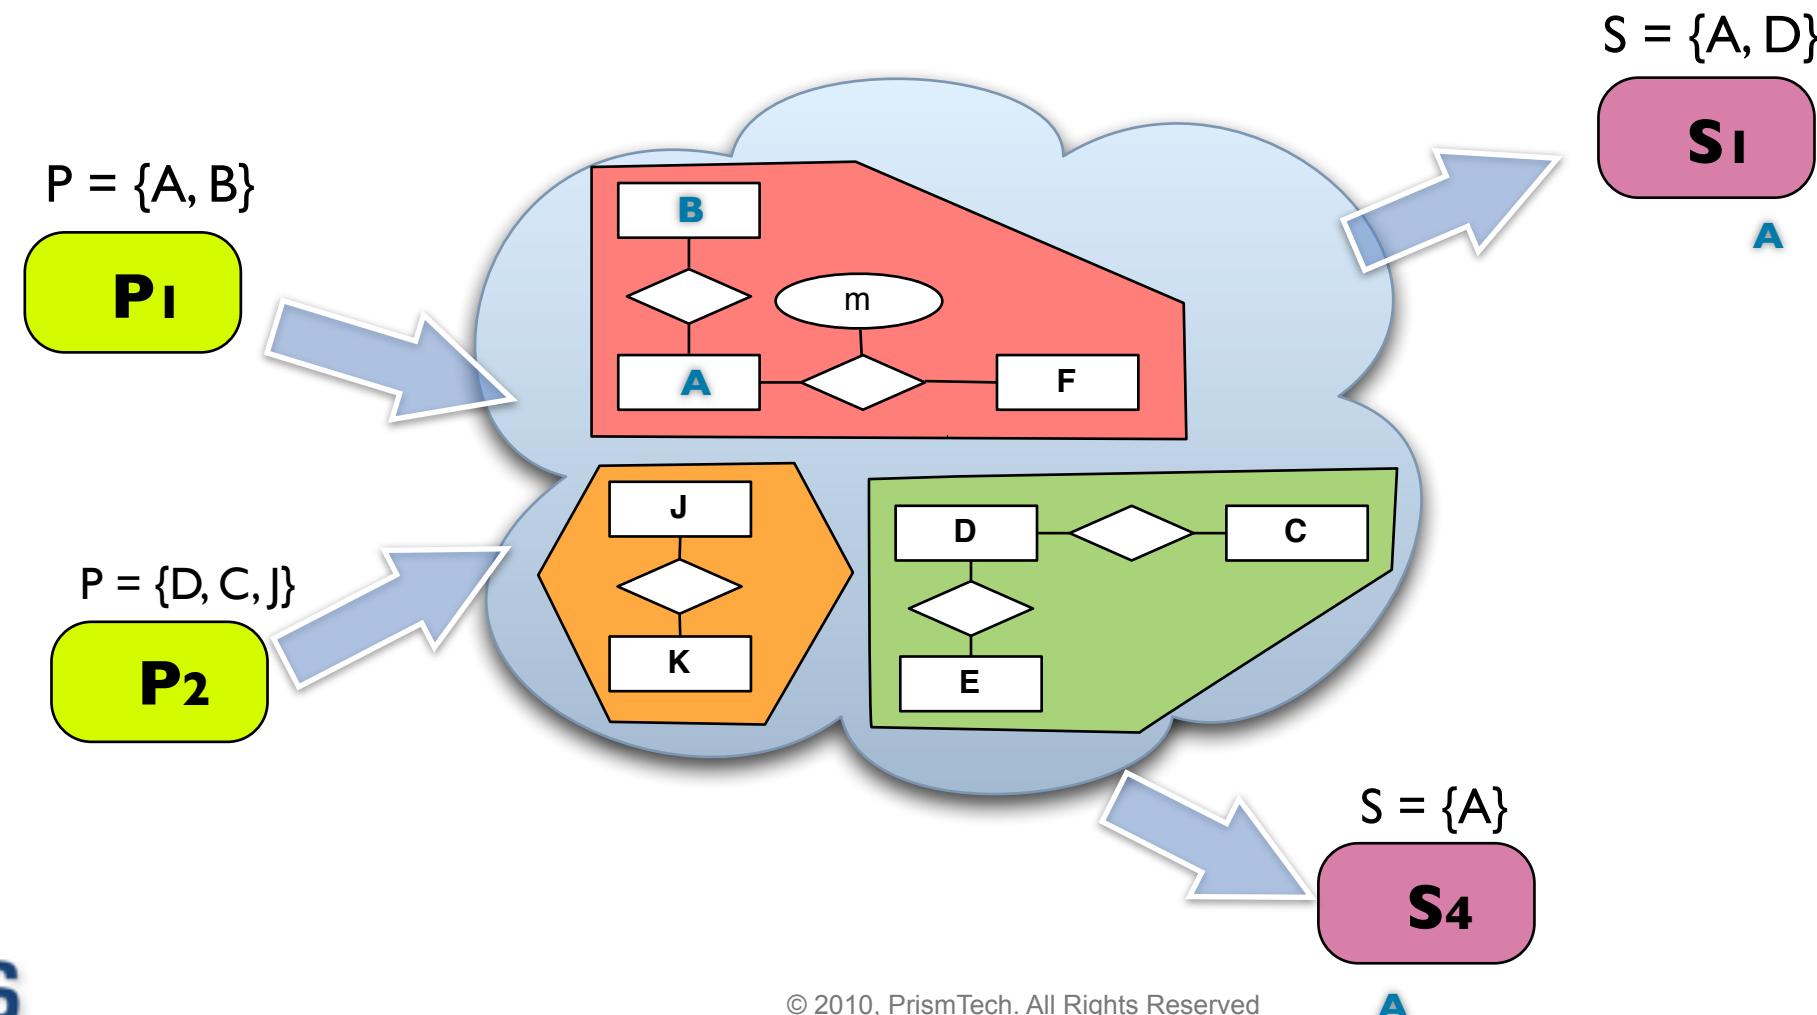
{A}  
{B}  
{J}



# Eventual Consistency @ Work

	DURABILITY	RELIABILITY	DESTINATION_ORDER	LIFESPAN
Eventual Consistency (Reader Crash / Recovery)	TRANSIENT_LOCAL	RELIABLE	SOURCE_TIMESTAMP	INF.
Eventual Consistency (Crash/Recovery)	TRANSIENT	RELIABLE	SOURCE_TIMESTAMP	INF.
Weak Consistency	ANY	ANY	ANY	N

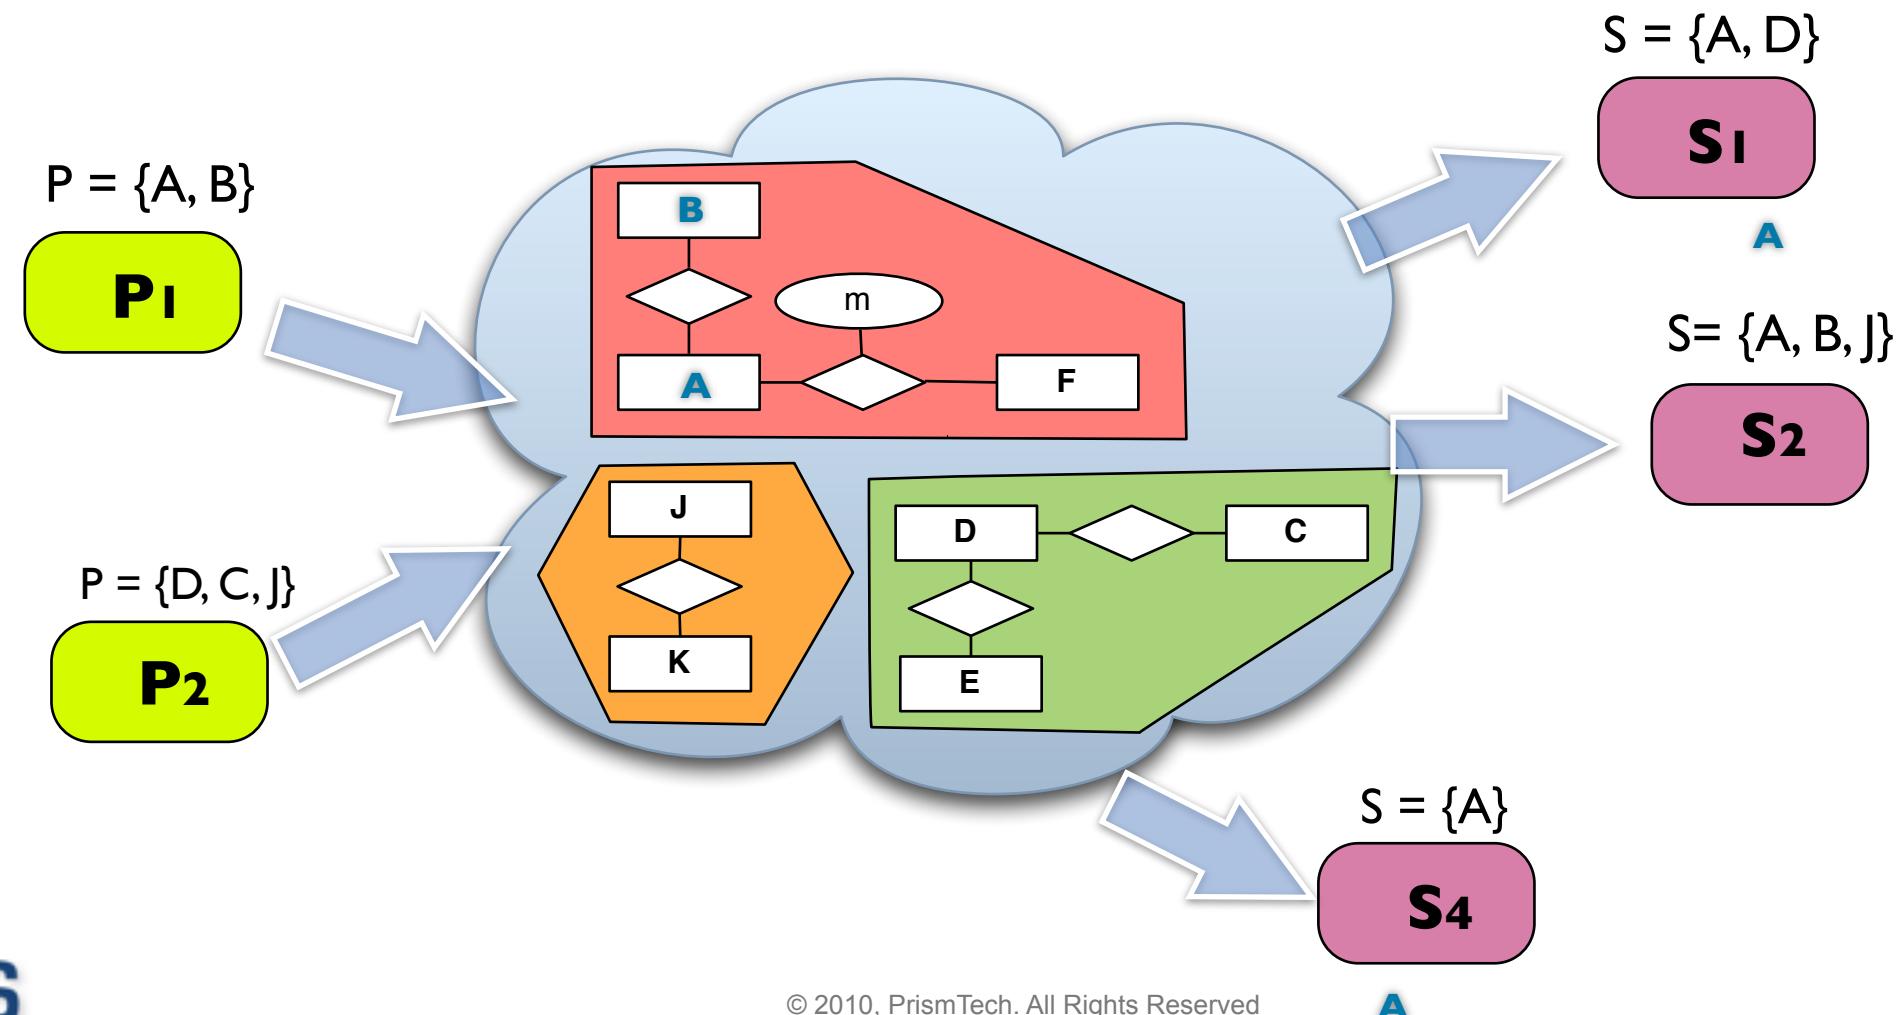
{A}  
{B}  
{J}



# Eventual Consistency @ Work

	DURABILITY	RELIABILITY	DESTINATION_ORDER	LIFESPAN
Eventual Consistency (Reader Crash / Recovery)	TRANSIENT_LOCAL	RELIABLE	SOURCE_TIMESTAMP	INF.
Eventual Consistency (Crash/Recovery)	TRANSIENT	RELIABLE	SOURCE_TIMESTAMP	INF.
Weak Consistency	ANY	ANY	ANY	N

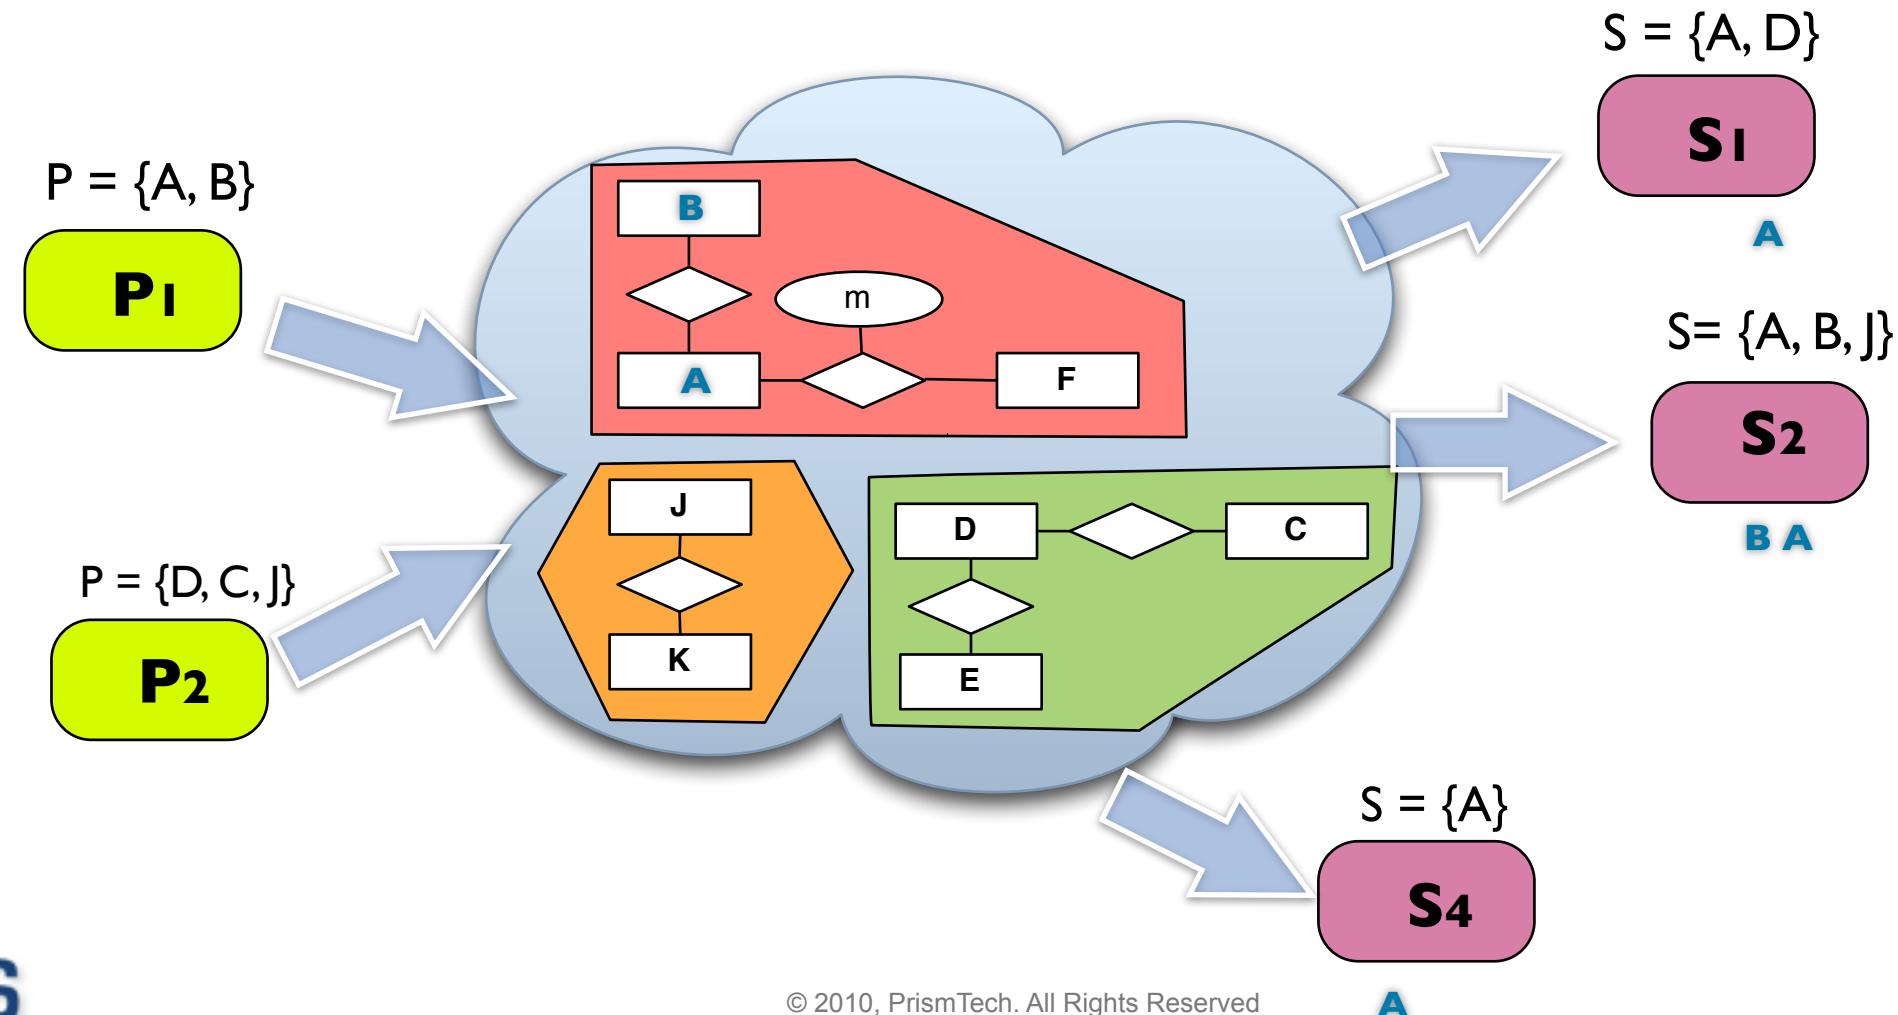
{A}  
{B}  
{J}



# Eventual Consistency @ Work

	DURABILITY	RELIABILITY	DESTINATION_ORDER	LIFESPAN
Eventual Consistency (Reader Crash / Recovery)	TRANSIENT_LOCAL	RELIABLE	SOURCE_TIMESTAMP	INF.
Eventual Consistency (Crash/Recovery)	TRANSIENT	RELIABLE	SOURCE_TIMESTAMP	INF.
Weak Consistency	ANY	ANY	ANY	N

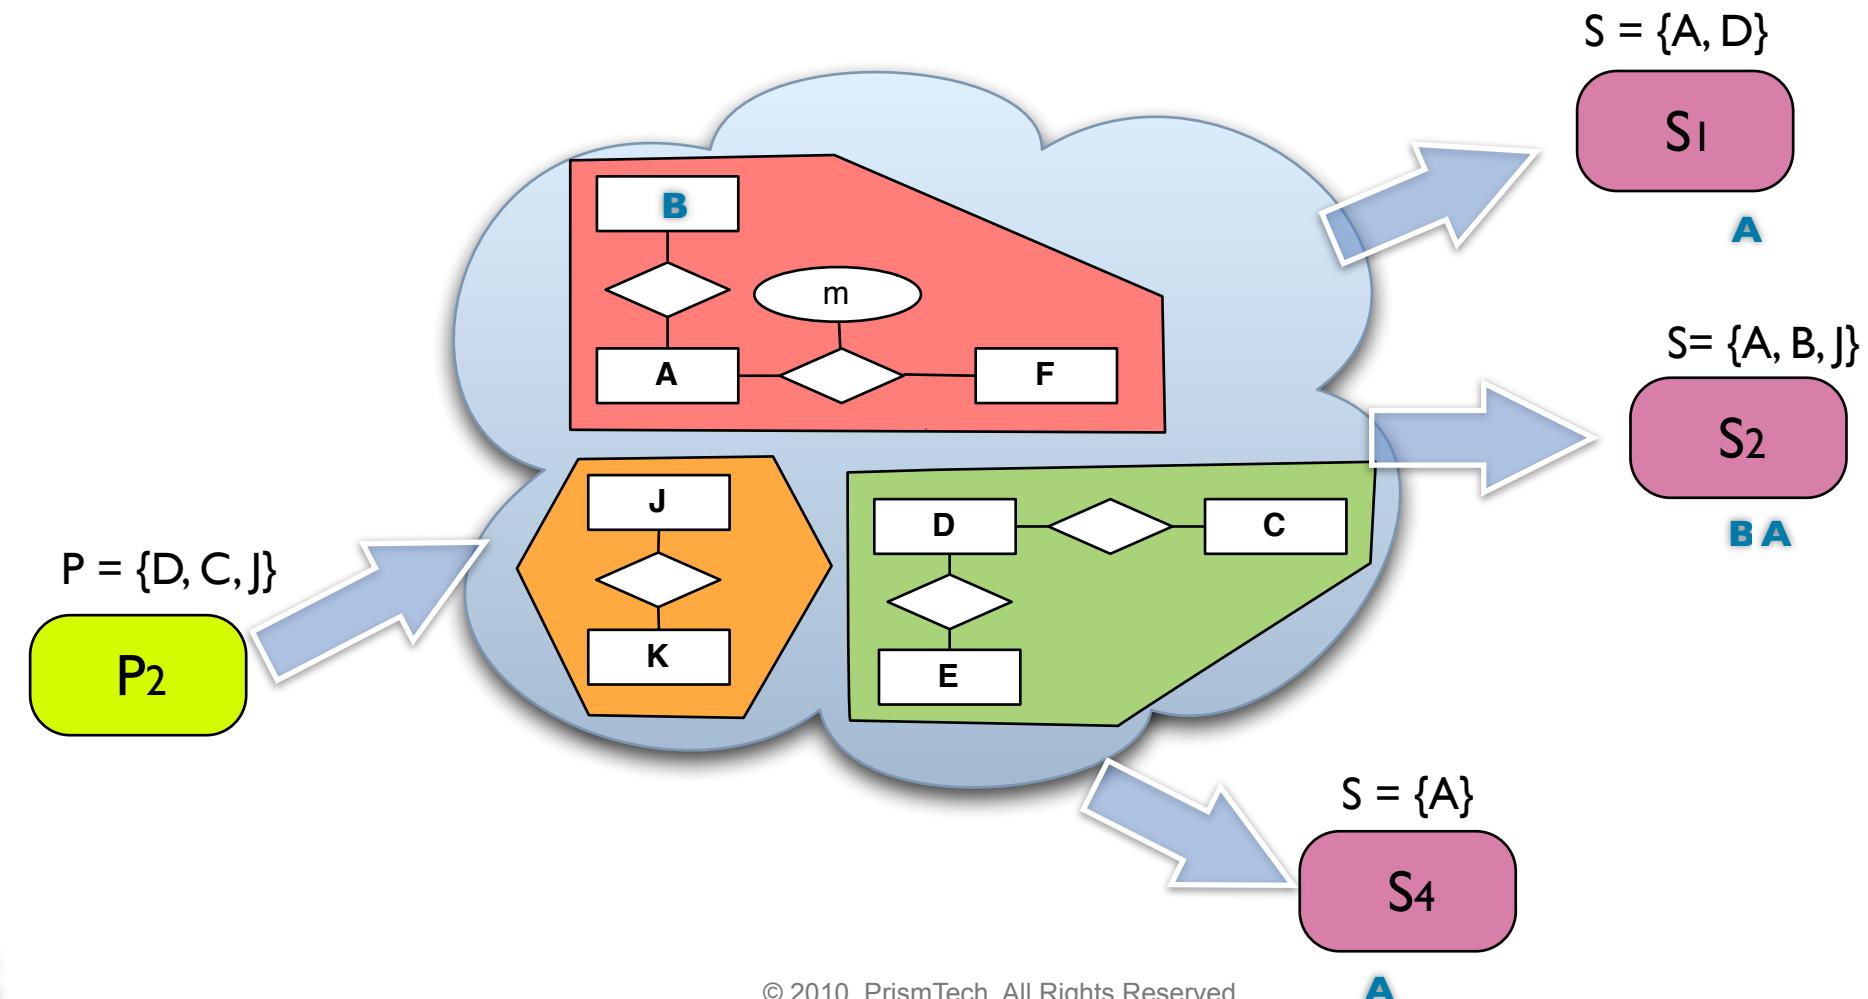
{A}  
{B}  
{J}



# Eventual Consistency @ Work

	DURABILITY	RELIABILITY	DESTINATION_ORDER	LIFESPAN
Eventual Consistency (Reader Crash / Recovery)	TRANSIENT_LOCAL	RELIABLE	SOURCE_TIMESTAMP	INF.
Eventual Consistency (Crash/Recovery)	TRANSIENT	RELIABLE	SOURCE_TIMESTAMP	INF.
Weak Consistency	ANY	ANY	ANY	N

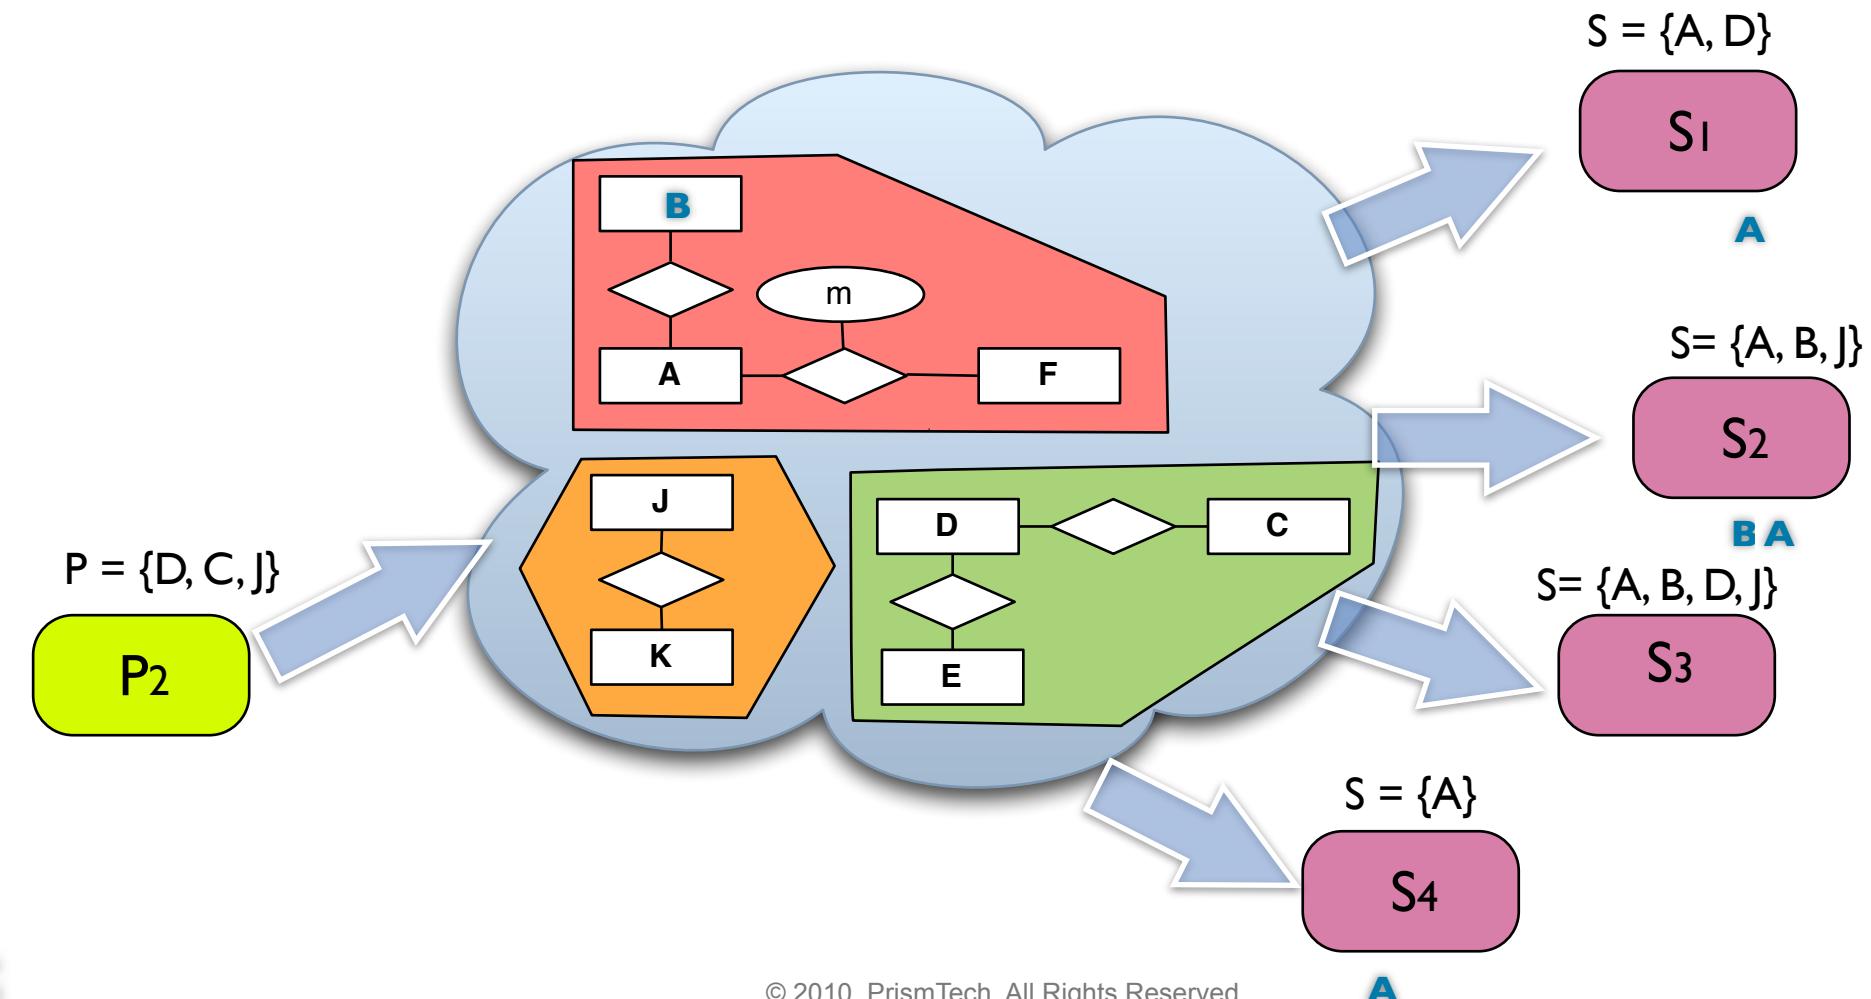
{A}  
{B}  
{J}



# Eventual Consistency @ Work

	DURABILITY	RELIABILITY	DESTINATION_ORDER	LIFESPAN
Eventual Consistency (Reader Crash / Recovery)	TRANSIENT_LOCAL	RELIABLE	SOURCE_TIMESTAMP	INF.
Eventual Consistency (Crash/Recovery)	TRANSIENT	RELIABLE	SOURCE_TIMESTAMP	INF.
Weak Consistency	ANY	ANY	ANY	N

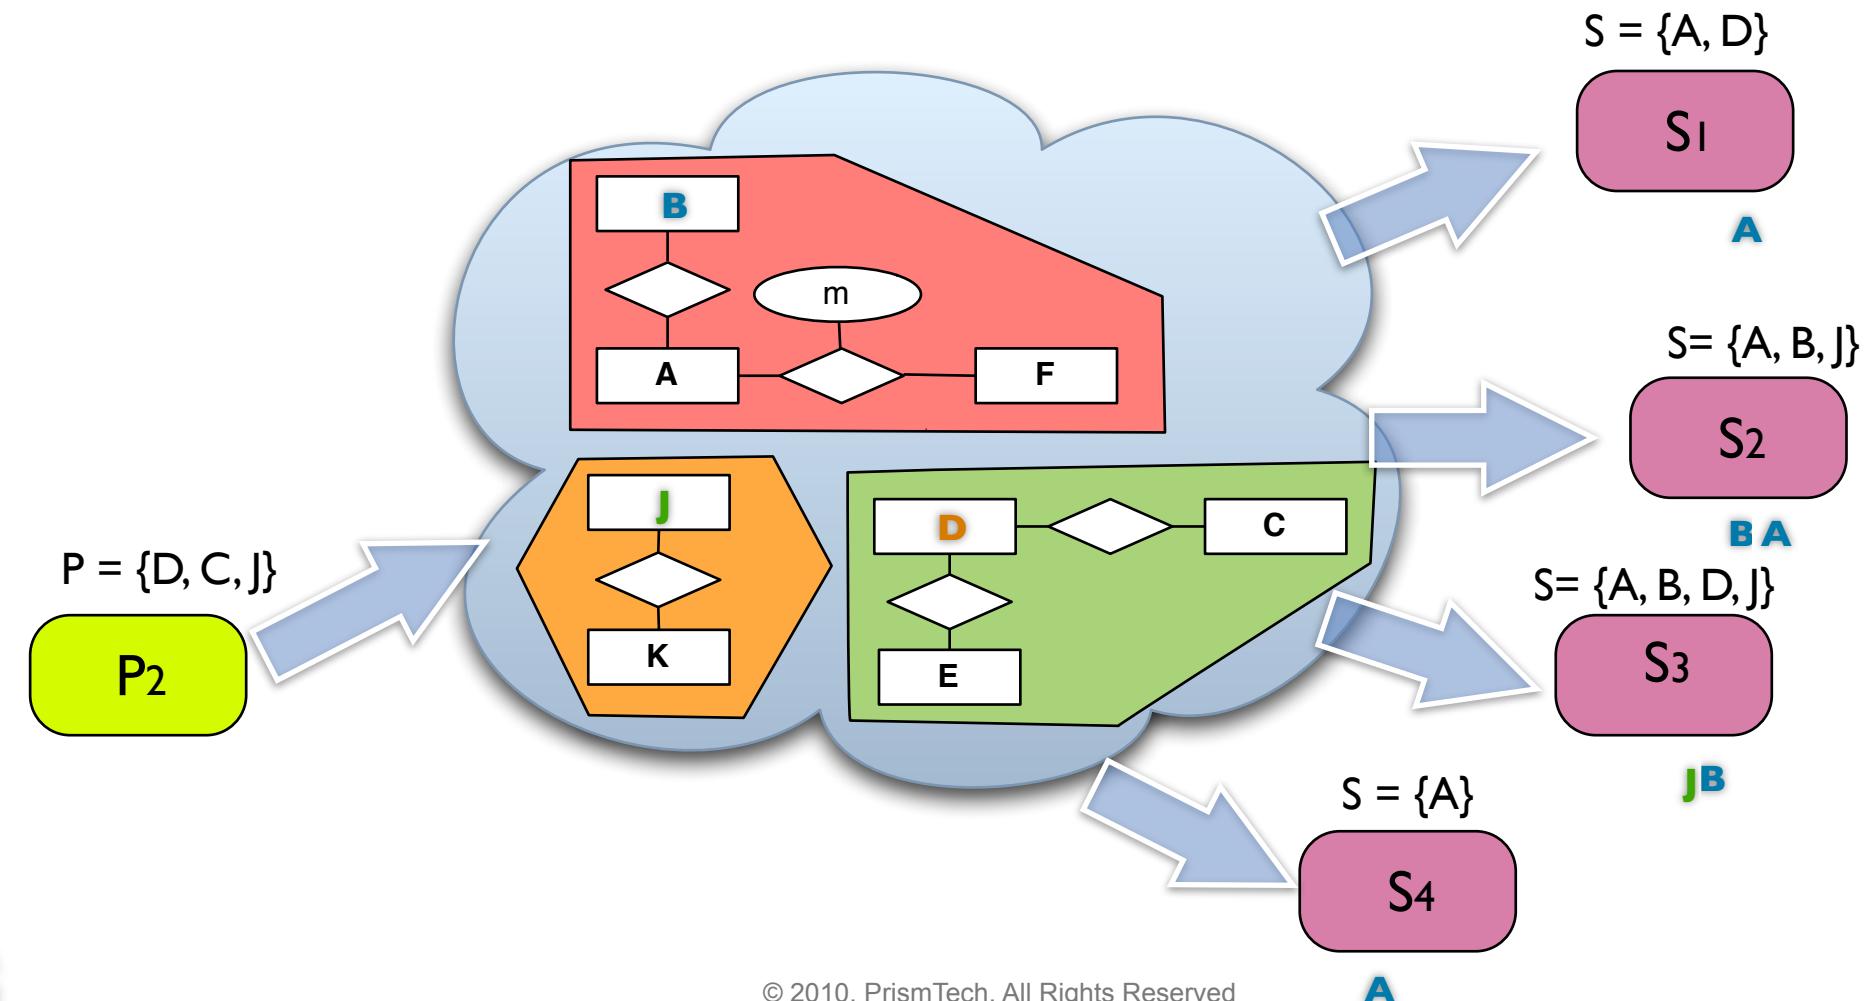
{A}  
{B}  
{J}



# Eventual Consistency @ Work

	DURABILITY	RELIABILITY	DESTINATION_ORDER	LIFESPAN
Eventual Consistency (Reader Crash / Recovery)	TRANSIENT_LOCAL	RELIABLE	SOURCE_TIMESTAMP	INF.
Eventual Consistency (Crash/Recovery)	TRANSIENT	RELIABLE	SOURCE_TIMESTAMP	INF.
Weak Consistency	ANY	ANY	ANY	N

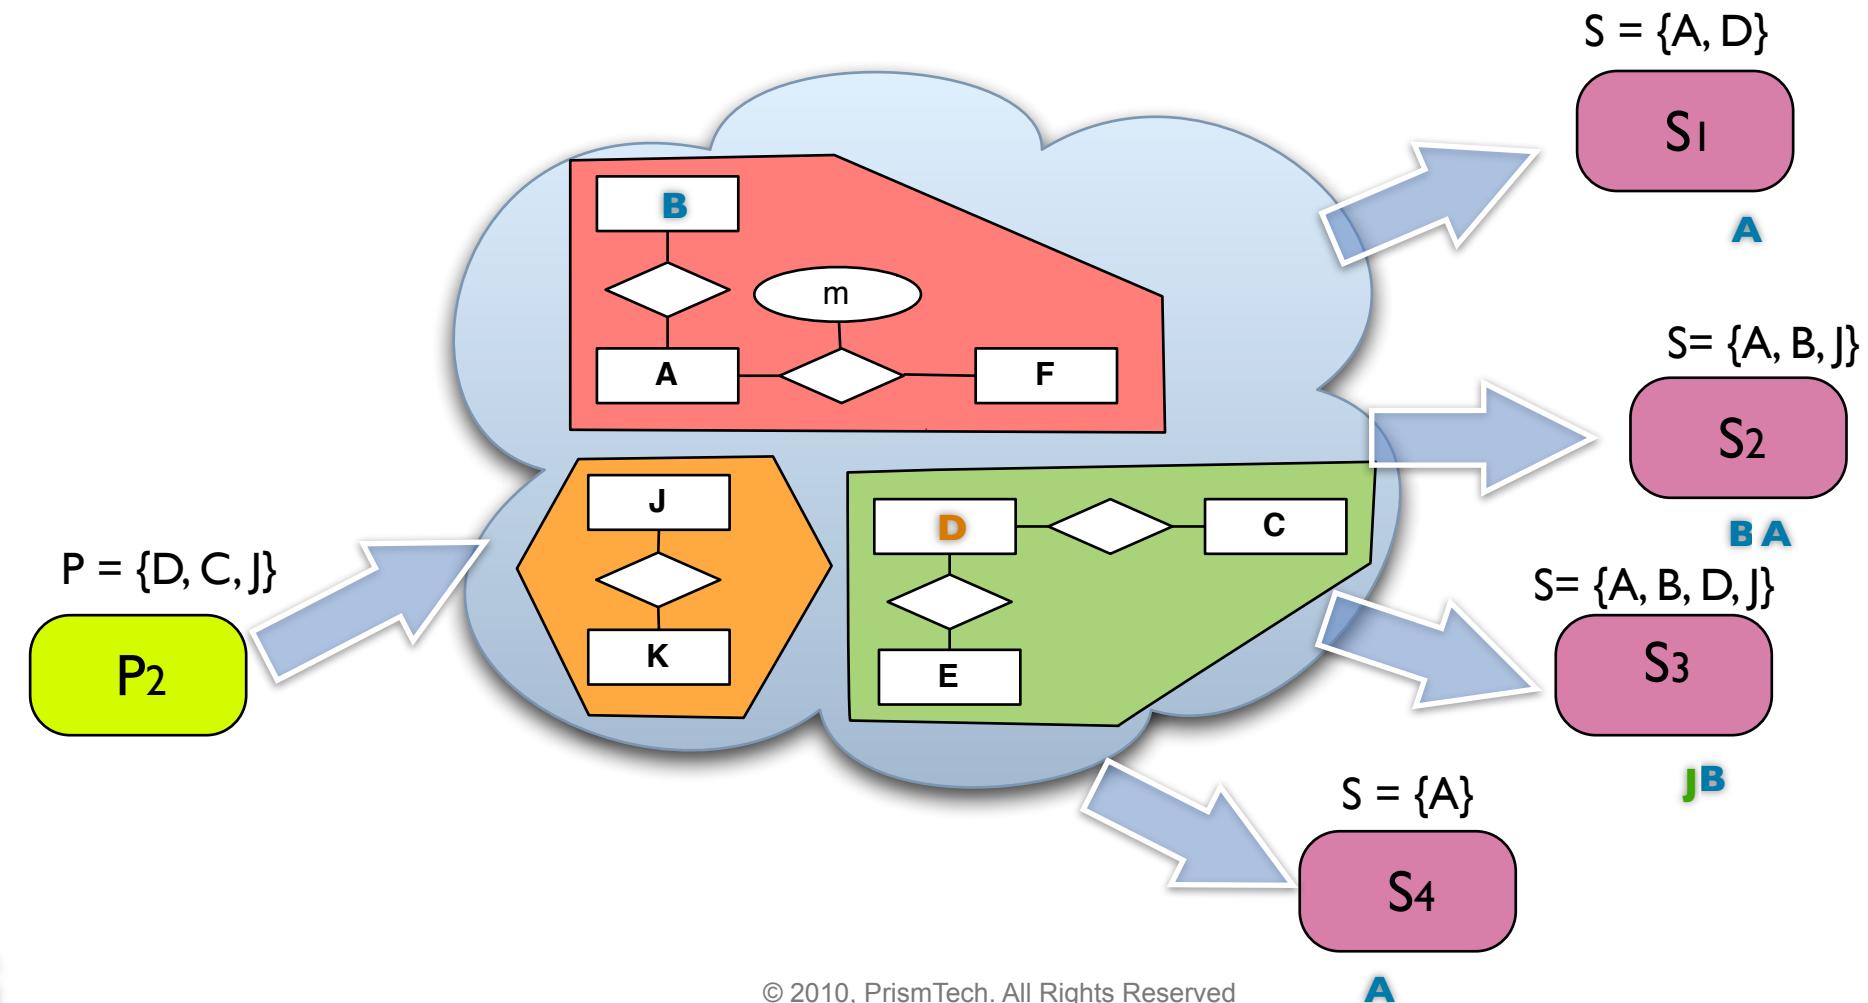
{A}  
{B}  
{J}



# Eventual Consistency @ Work

	DURABILITY	RELIABILITY	DESTINATION_ORDER	LIFESPAN
Eventual Consistency (Reader Crash / Recovery)	TRANSIENT_LOCAL	RELIABLE	SOURCE_TIMESTAMP	INF.
Eventual Consistency (Crash/Recovery)	TRANSIENT	RELIABLE	SOURCE_TIMESTAMP	INF.
Weak Consistency	ANY	ANY	ANY	N

{A}  
{B}  
{J}



# OpenSplice|DDS

Delivering Performance, Openness, and Freedom

## Controlling Replication

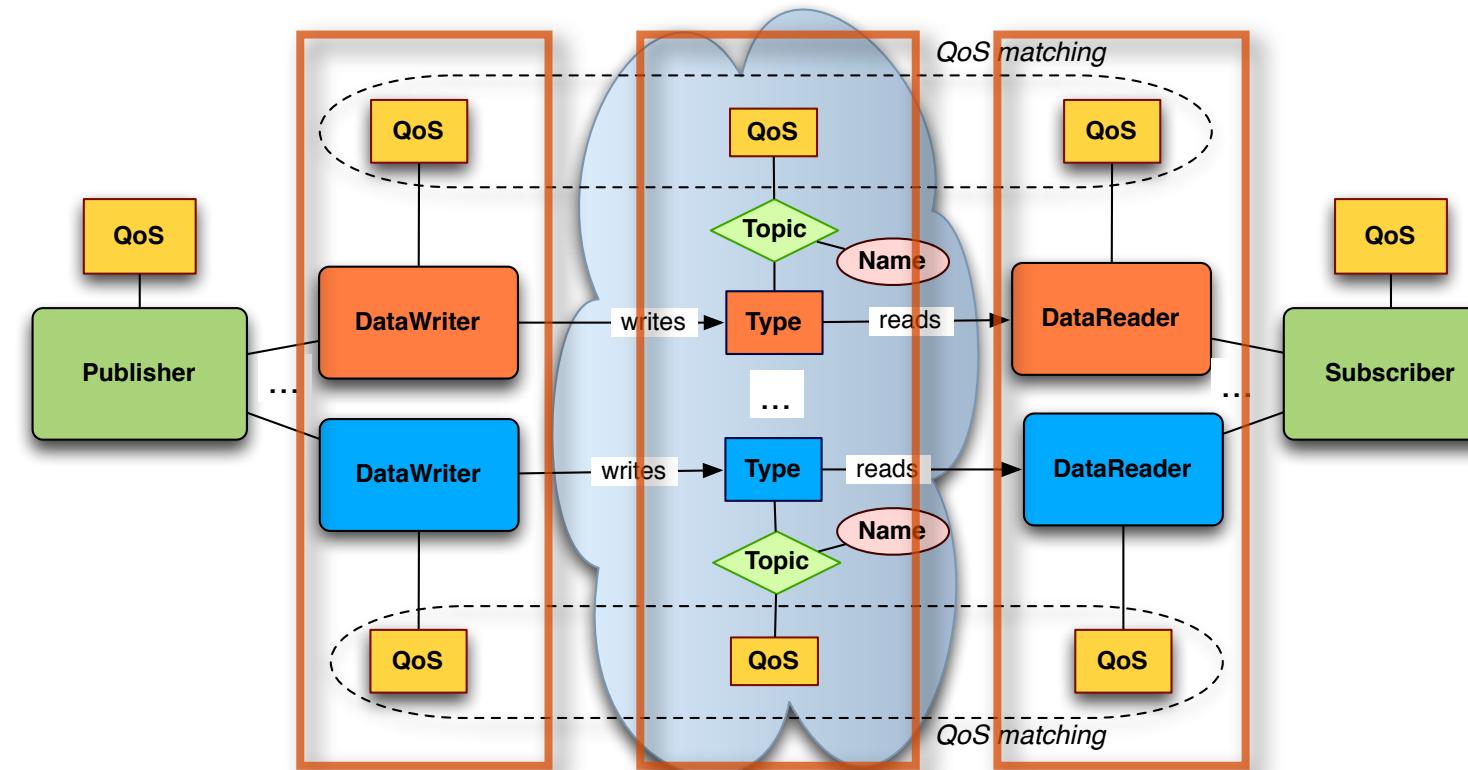
# Ownership

Who owns the data?

The **OWNERSHIP** QoS specifies whether it is allowed for multiple DataWriters to write the same instance of the data and if so, how these modifications should be arbitrated. Possible choices are:

- ▶ **Shared.** Multiple writers are allowed to update the same instance and all the updates are made available to the reader
- ▶ **Exclusive.** Indicates that each instance can only be owned by one DataWriter, but the owner of an instance can change dynamically – due to liveliness changes
- ▶ The selection of the owner is controlled by the setting of the **OWNERSHIP\_STRENGTH** **QoS policy**

QoS Policy	Applicability	RxO	Modifiable
OWNERSHIP	T, DR, DW	Y	N



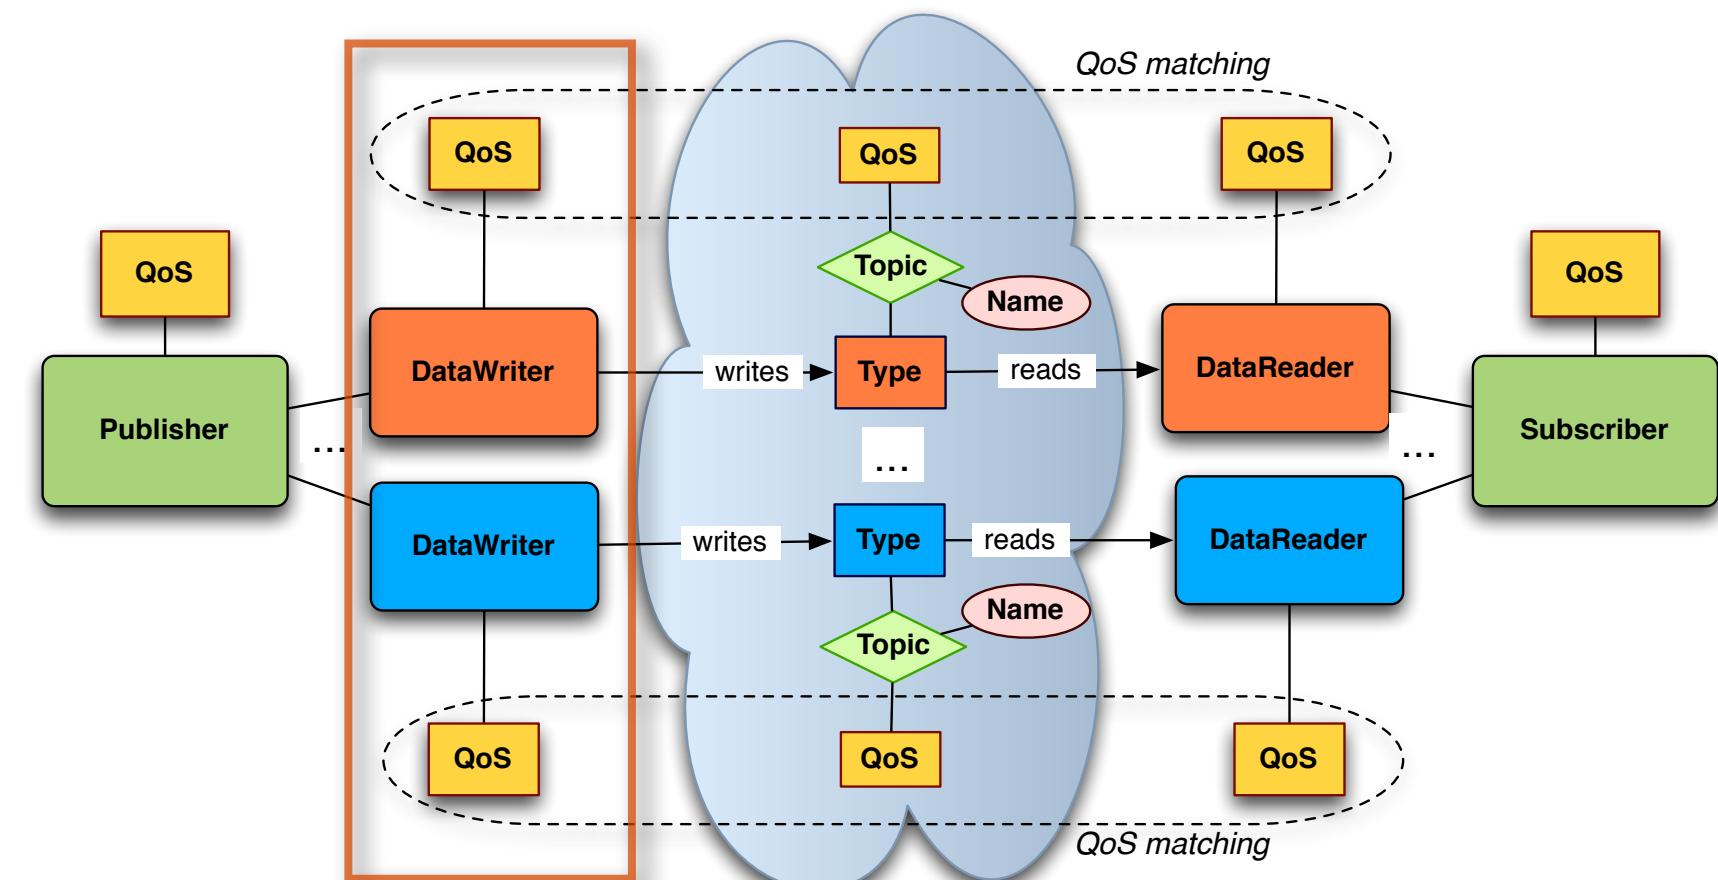
# Ownership Strength

How strong are you?

The **OWNERSHIP\_STRENGTH** Specifies the value of the “strength” used to arbitrate among DataWriters that attempt to modify the same data instance

- ▶ Data instances are identified by the couple (Topic, Key)
- ▶ The policy applies only if the OWNERSHIP is EXCLUSIVE

QoS Policy	Applicability	RxO	Modifiable
OWNERSHIP STRENGTH	DW	-	Y



# Next Steps

- ▶ By now you've learned most of what you need to write complex DDS applications
- ▶ However, as the wise Confucius used to say:
  - ▶ I hear and I forget. I see and I remember. I do I understand.
- ▶ The best way of really getting into DDS is to write some DDS applications, utilities or extensions



# Online Resources



Delivering Performance, Openness, and Freedom

- \* <http://www.opensplice.com/>
- \* <mailto:opensplicedd@prismtech.com>



- \* <http://bit.ly/1Sreg>



- \* <http://www.youtube.com/OpenSpliceTube>



- \* <http://www.slideshare.net/angelo.corsaro>



- \* <http://twitter.com/acorsaro/>



- \* <http://opensplice.blogspot.com>